



Kubernetes Arabic Course

The background of the slide features a dark blue gradient with a faint, abstract technical illustration. It includes several binary code numbers (e.g., 10010010, 11010110) scattered across the screen, some glowing blue. There are also stylized representations of mechanical components like gears and a steering wheel, all in a light blue color that blends with the background.

Kubernetes From Scratch - Hands-on

01 - Introduction



Course Contents

0 Introduction

- Course Introduction
- What is Kubernetes
- Kubernetes Features

1 Kubernetes Architecture

- Cluster Architecture
- Master Components
- Worker Components

0 Setup Lab Cluster

- Setup Single Node Cluster (Minikube)
- Setup Kind Cluster
- Setup Internal Docker Registry

2 Basic Concepts

- Namespaces
- Resource Quota
- Labels and Selectors
- Nodes

Course Contents

0 Pods

- Pods Explained
- Pods demo
- Multi-container pod

5 Deployments

- Deployment Explained
- Deployment Demo
- Rolling Update

0 Services

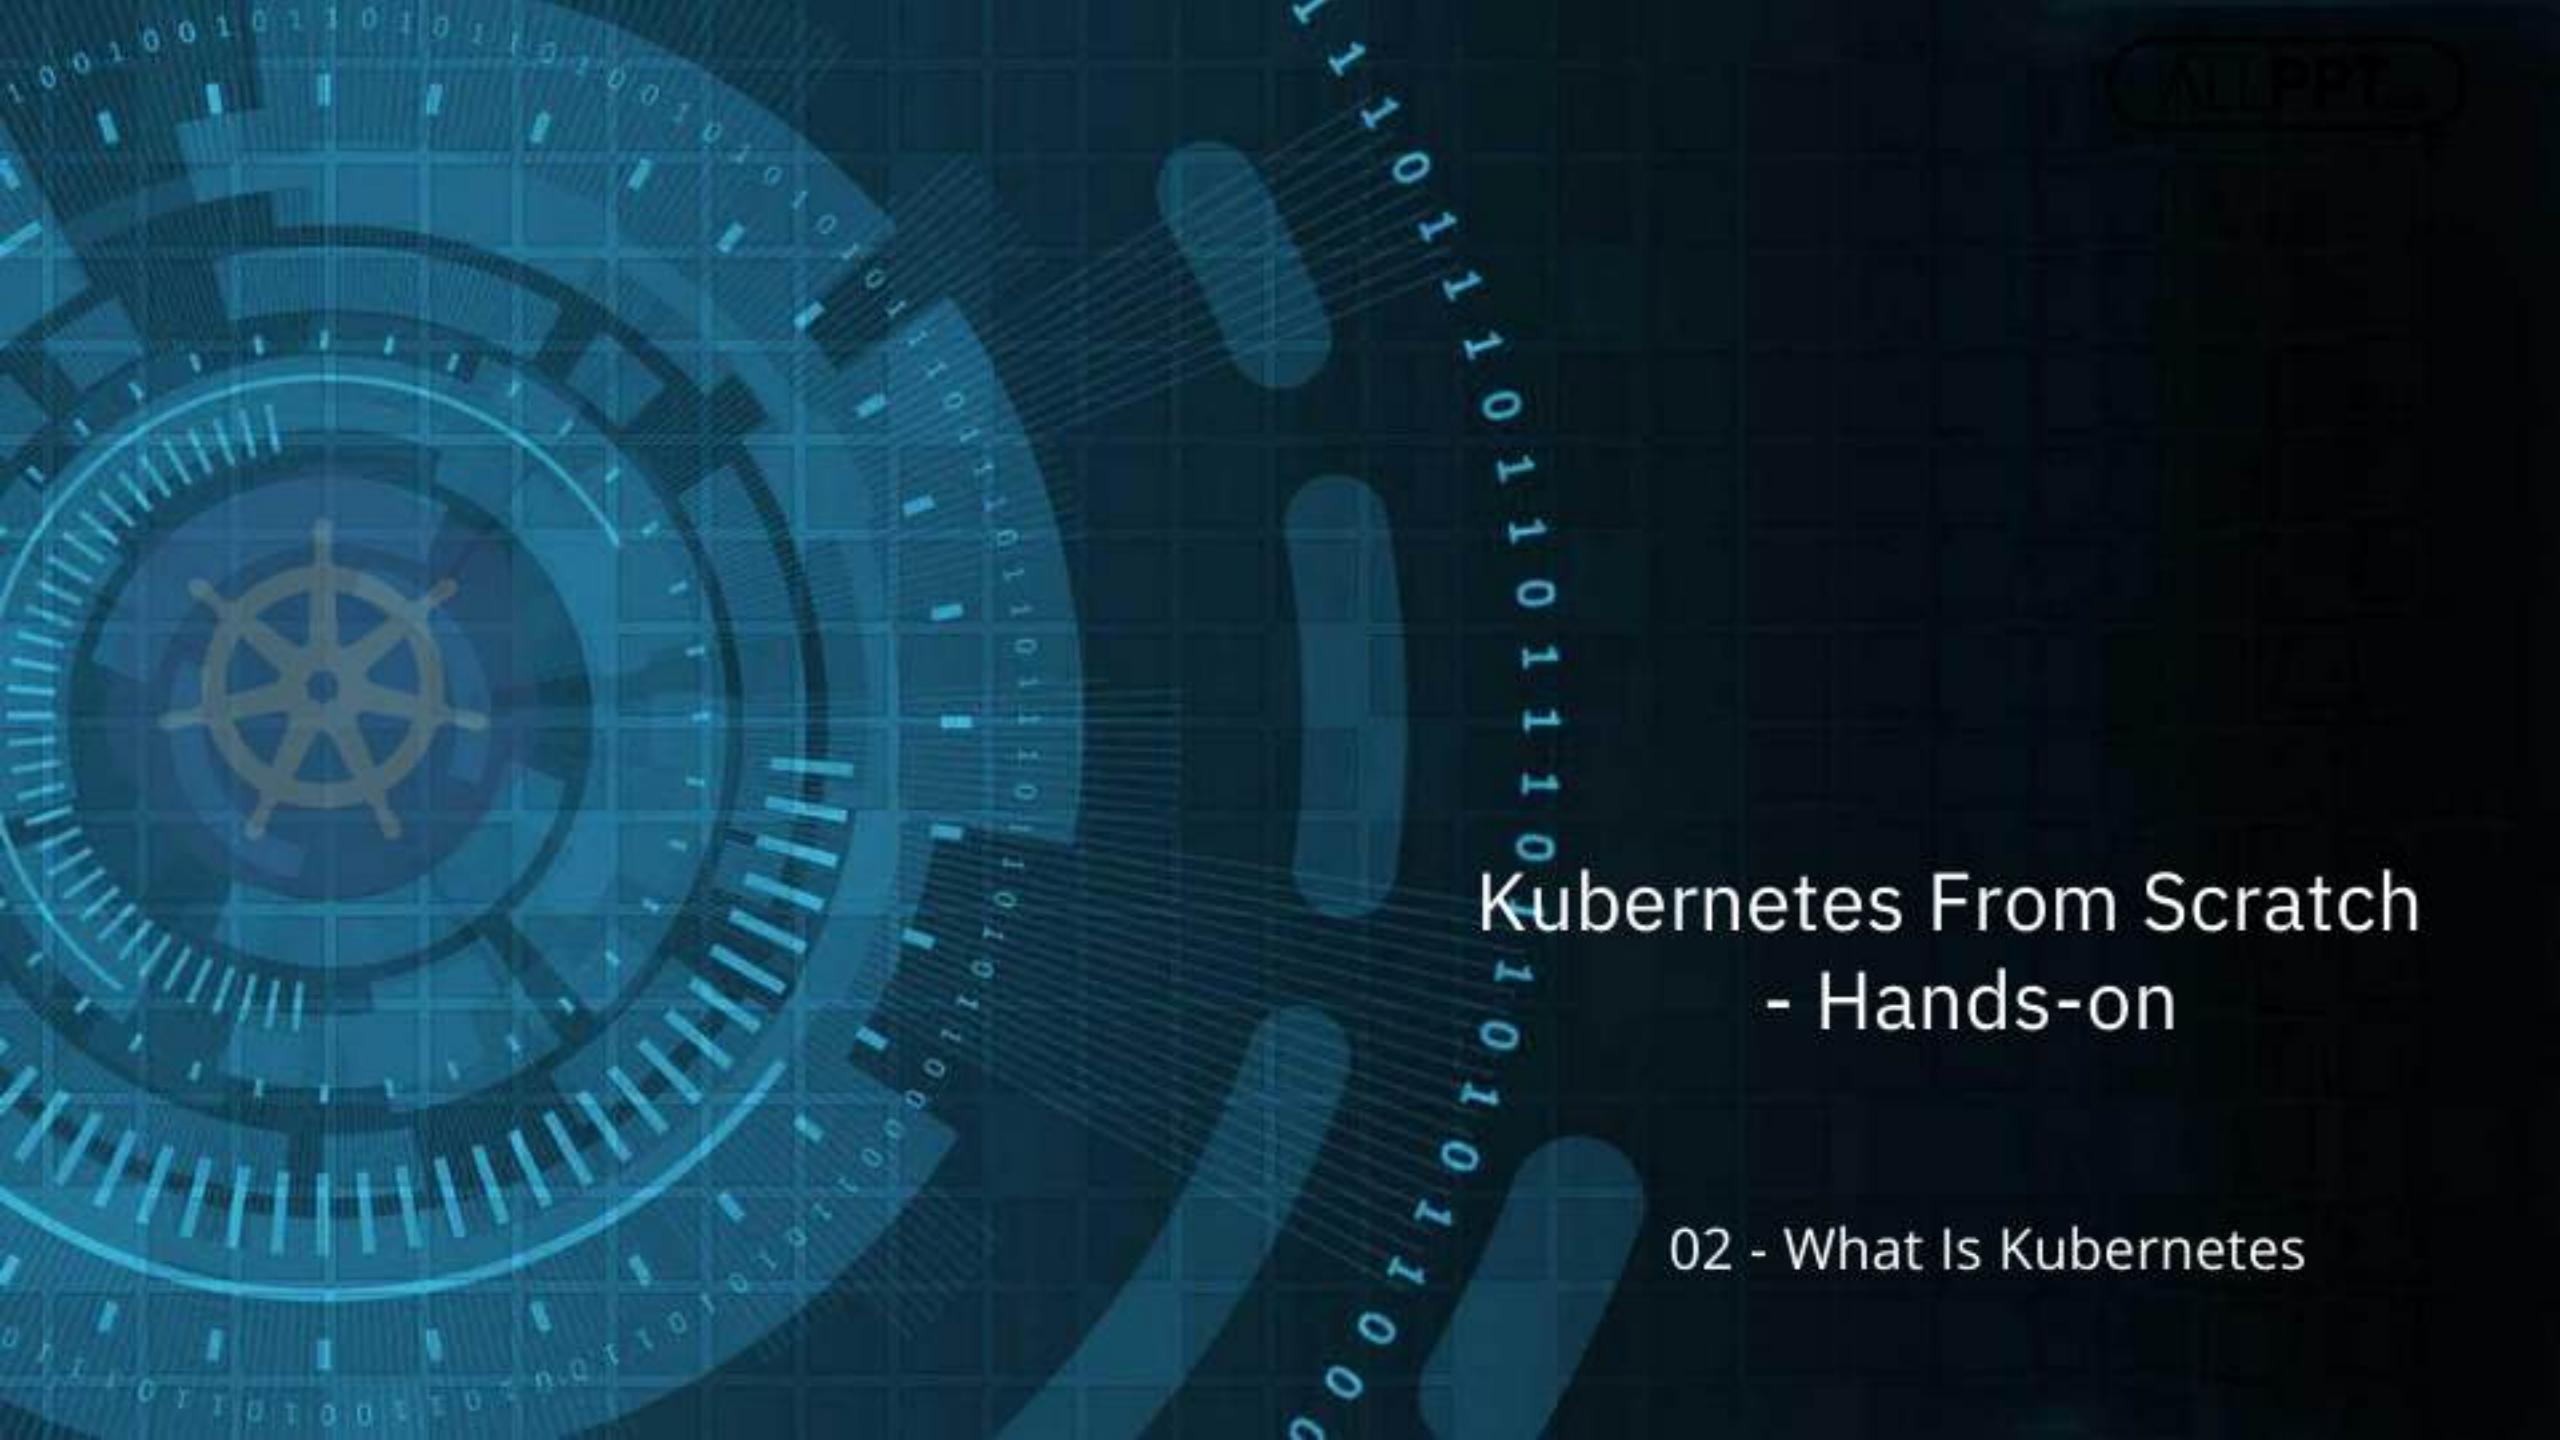
- Service Explained
- Service Types
- Service Demo

6 Ingress

- Ingress Explained
- Deploy Nginx Ingress Controller
- Ingress Demo

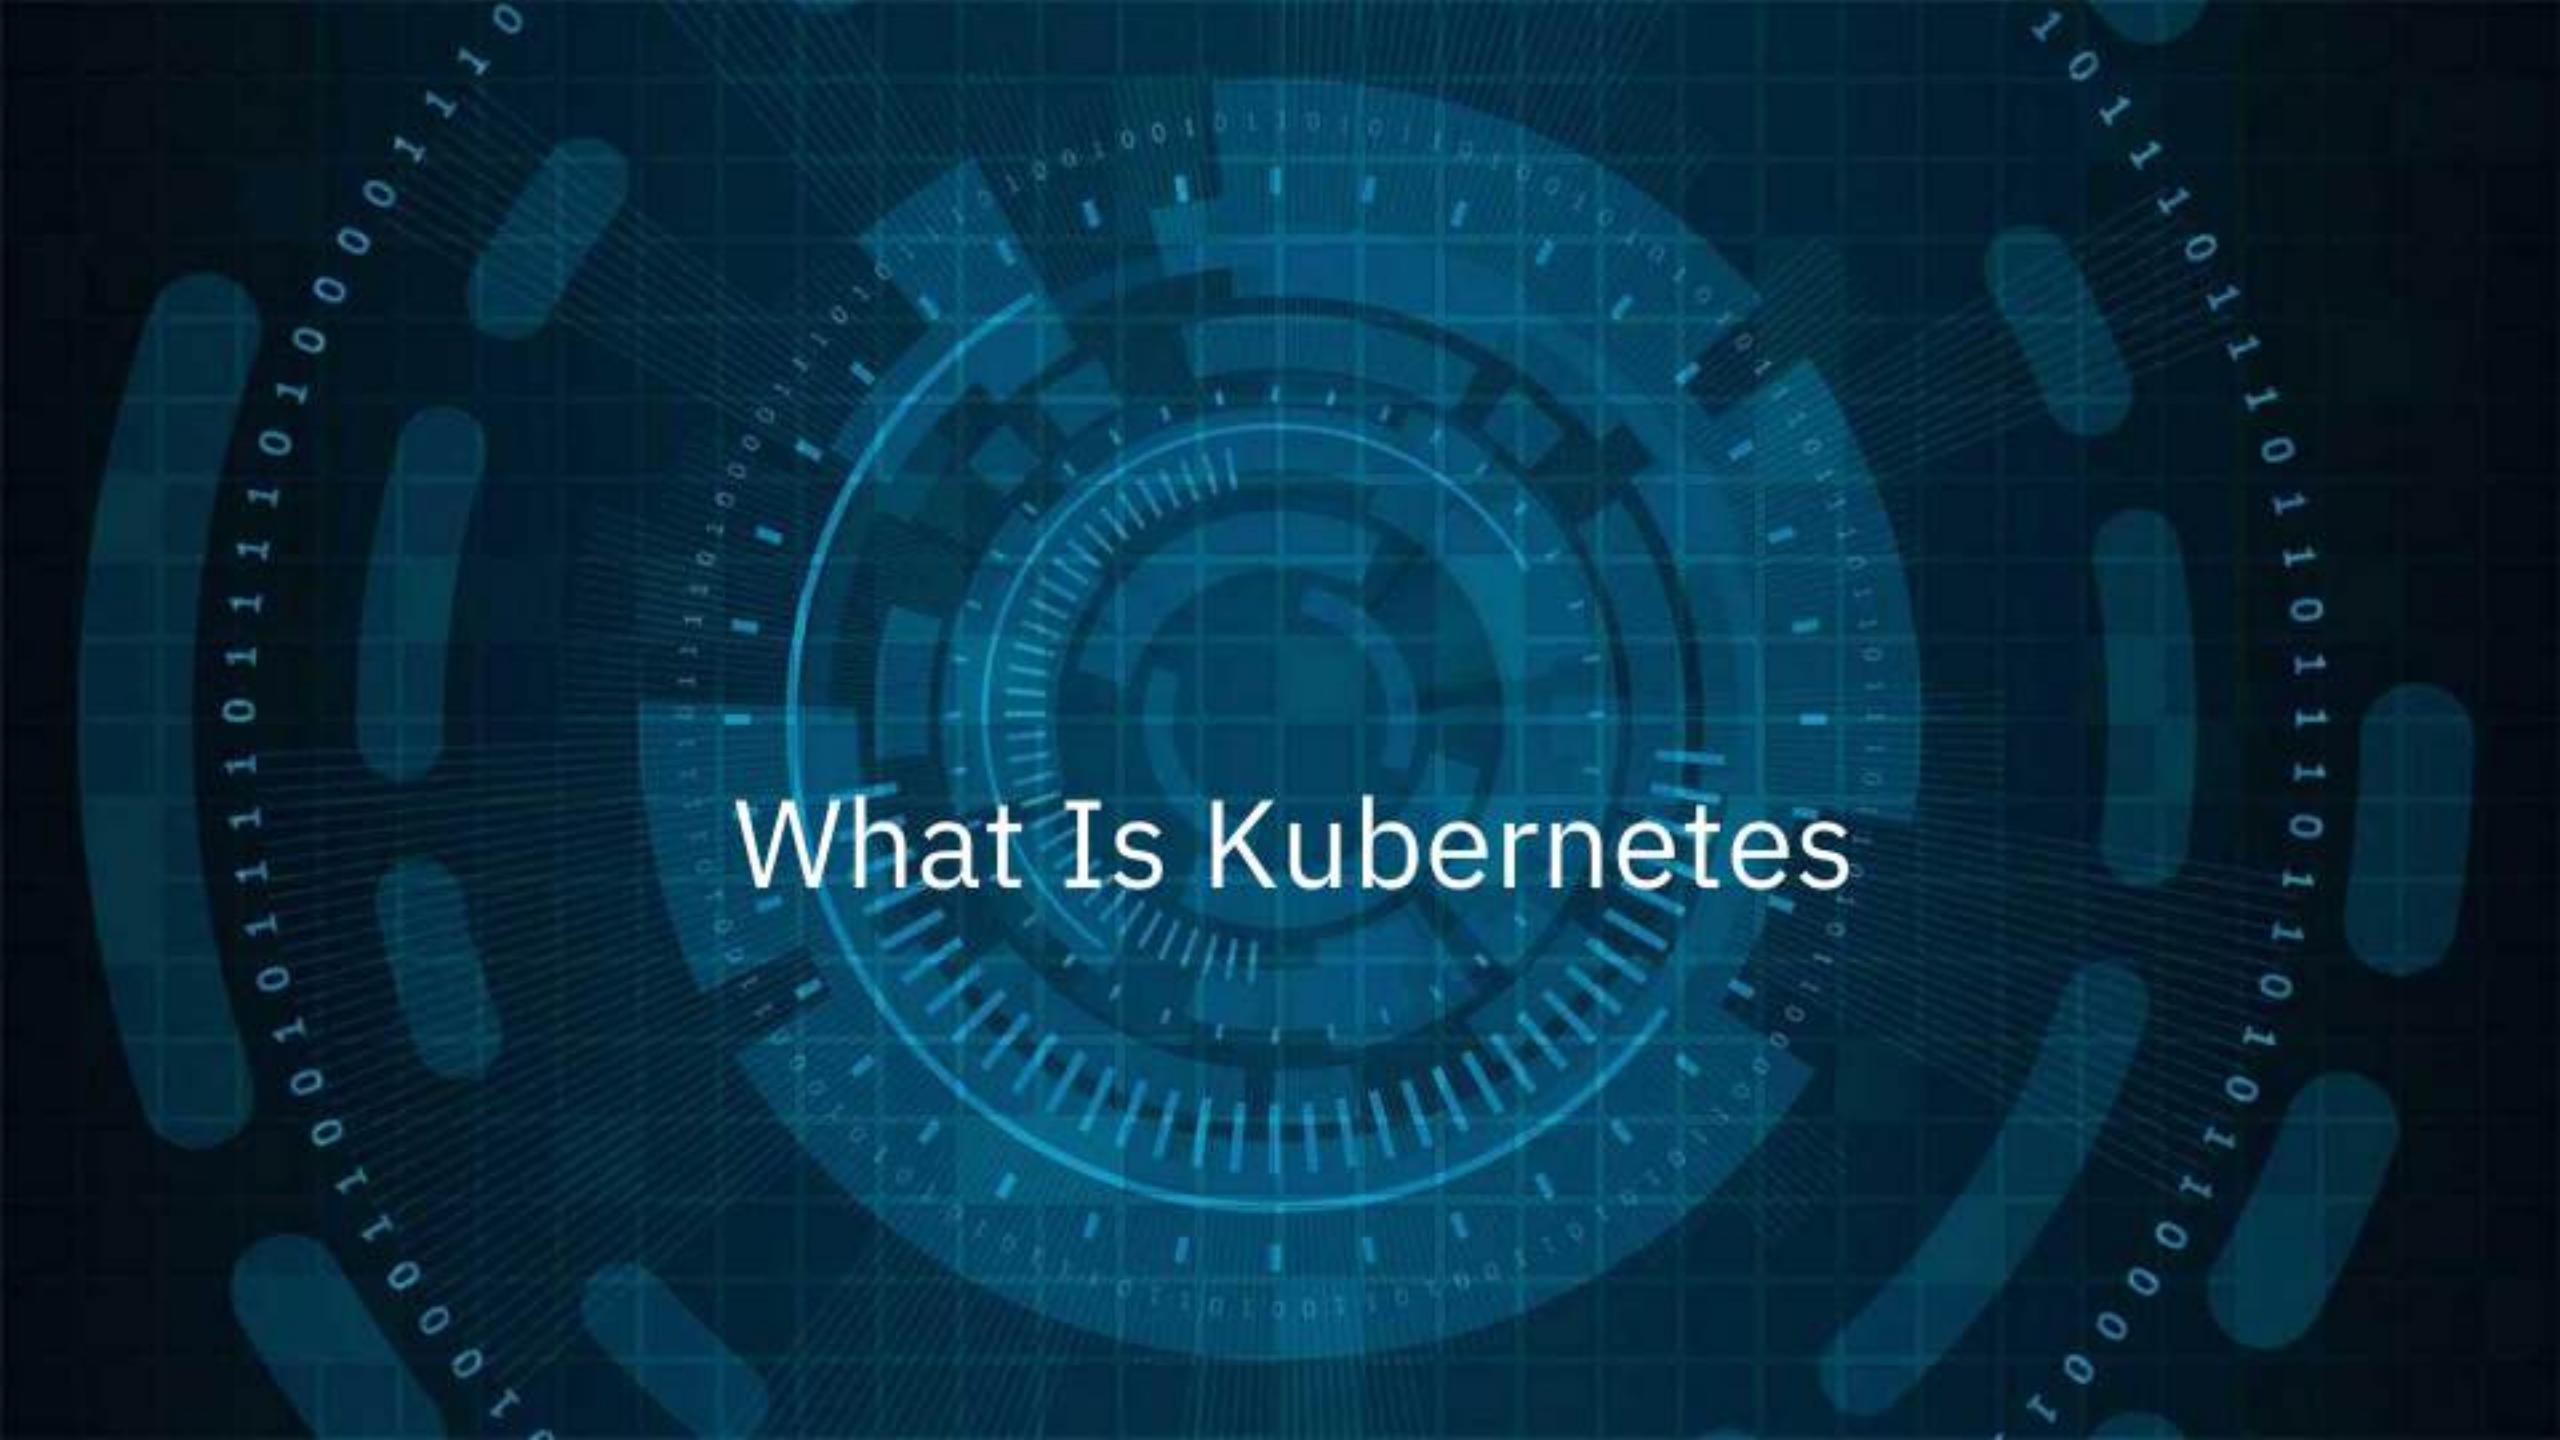
0





Kubernetes From Scratch - Hands-on

02 - What Is Kubernetes

The background of the slide features a dark blue to black gradient with a subtle grid pattern. Overlaid on this are several glowing, semi-transparent blue and white nodes of varying sizes, some with dashed outlines. Binary code (0s and 1s) is displayed in various fonts and orientations around these nodes, particularly along the top and right edges.

What Is Kubernetes

What is Kubernetes

Container Orchestration

It's an open source container orchestration platform, to automate and manage containerized applications.

Developed at Google

It was developed at Google and released as open source in 2014. So it combines the best practices of running containerized workloads at google for 15 years.

Cloud Agnostic

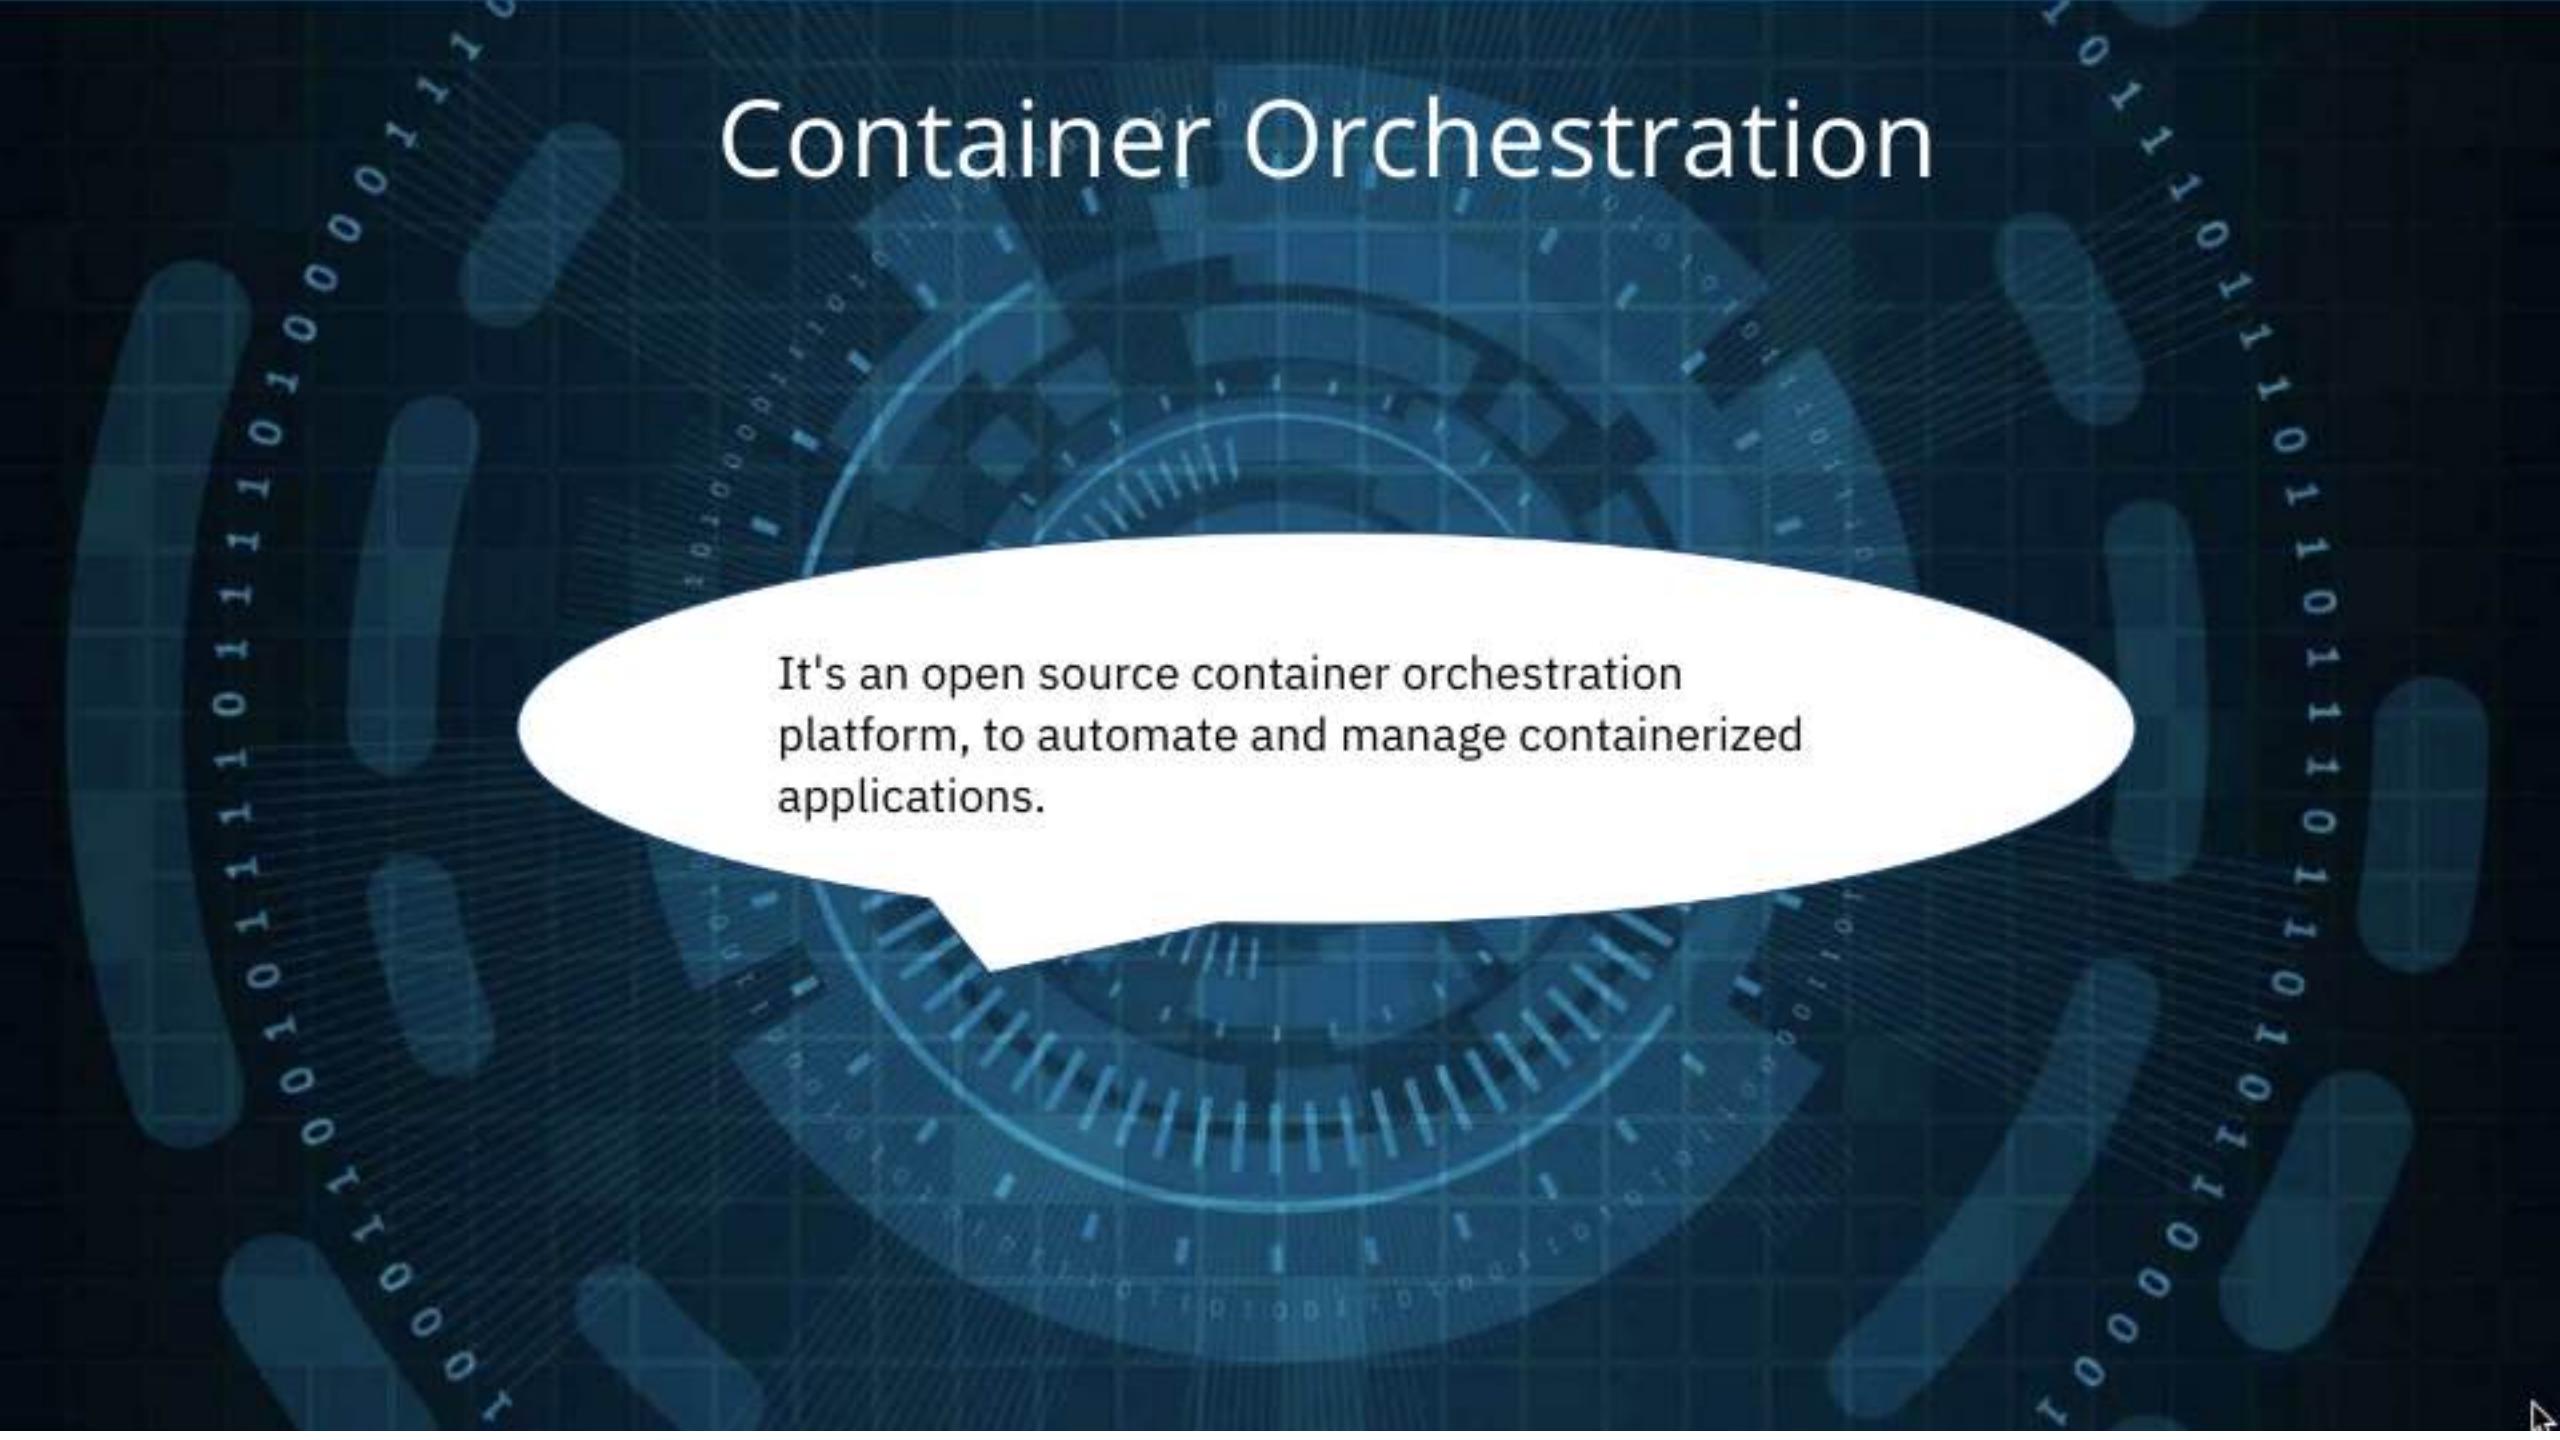
you can run it anywhere on bare metal or in any cloud provider infrastructure.



Zero Downtime Deployment

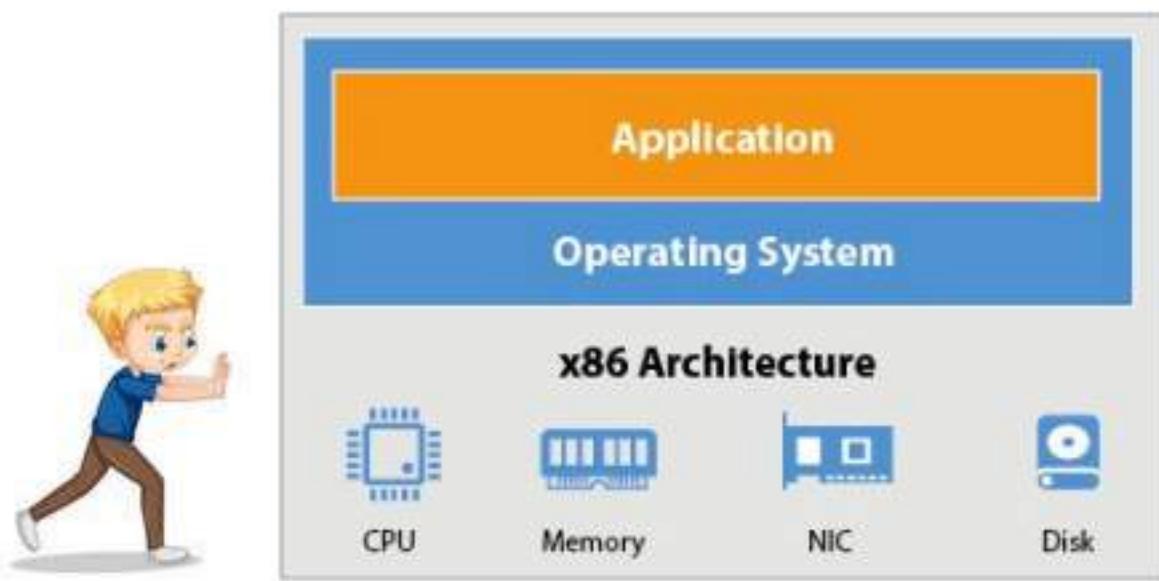
It became the industry standard for deploying containers in production

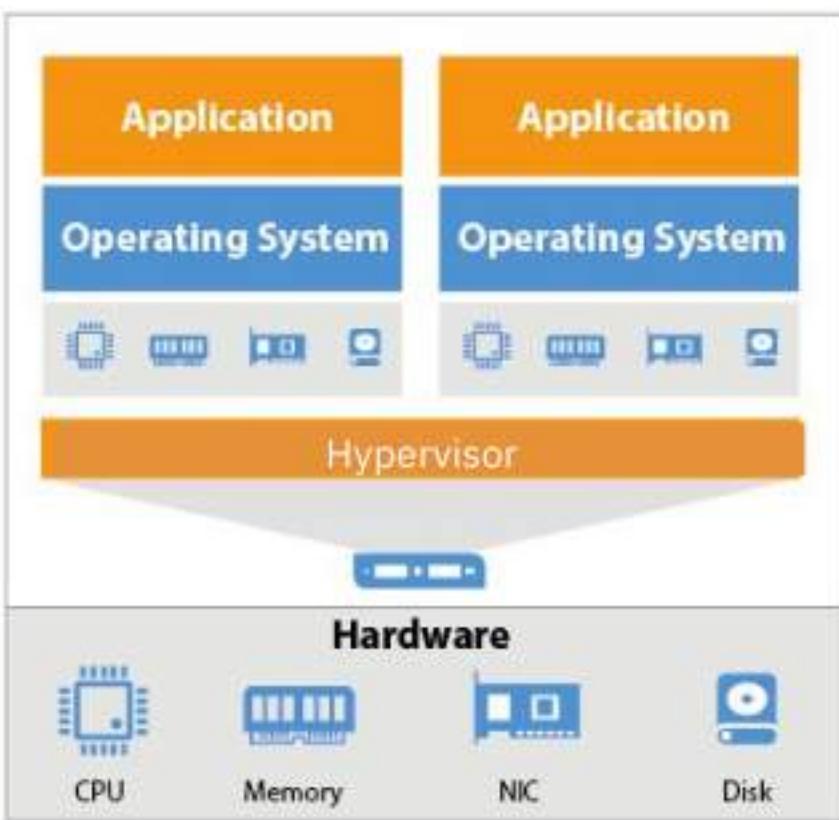
Container Orchestration

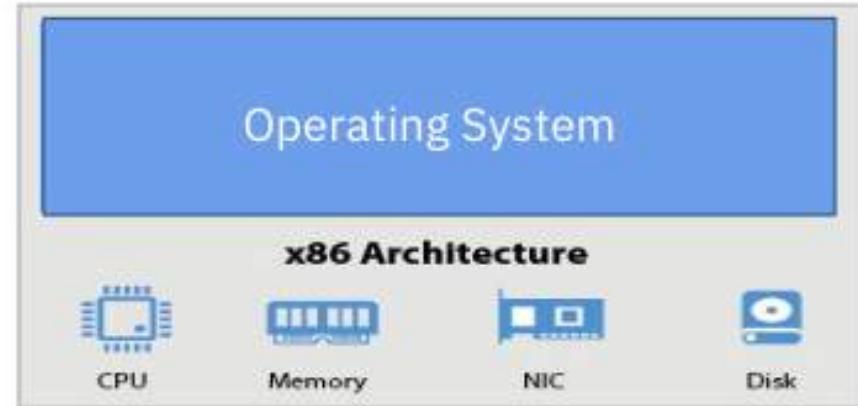
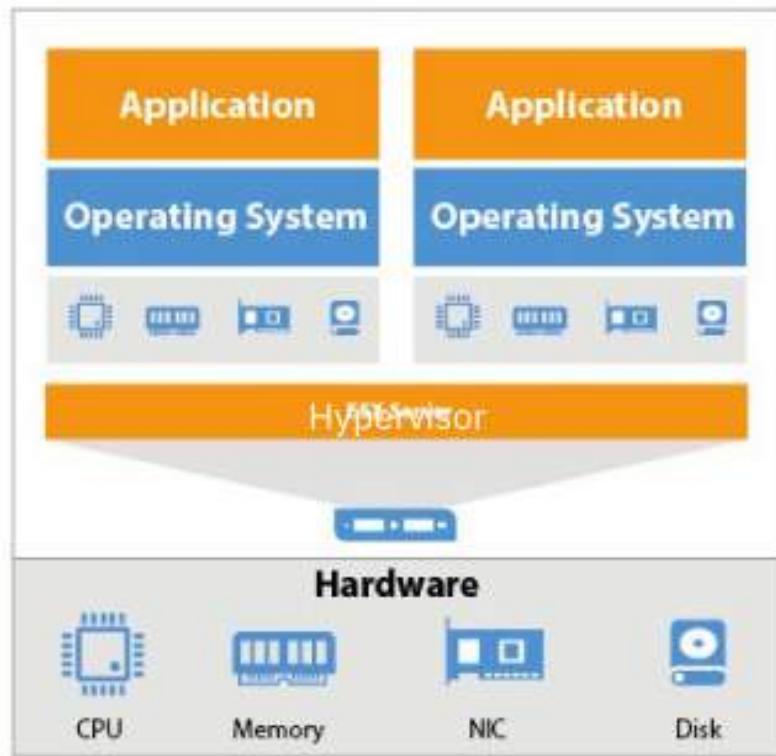


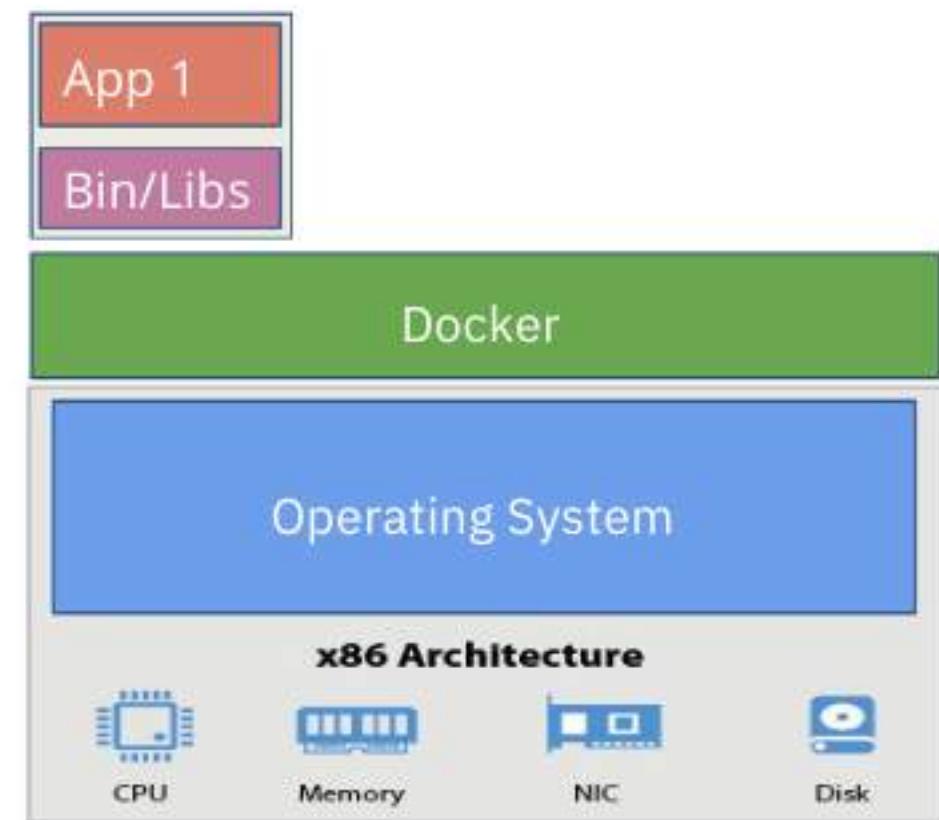
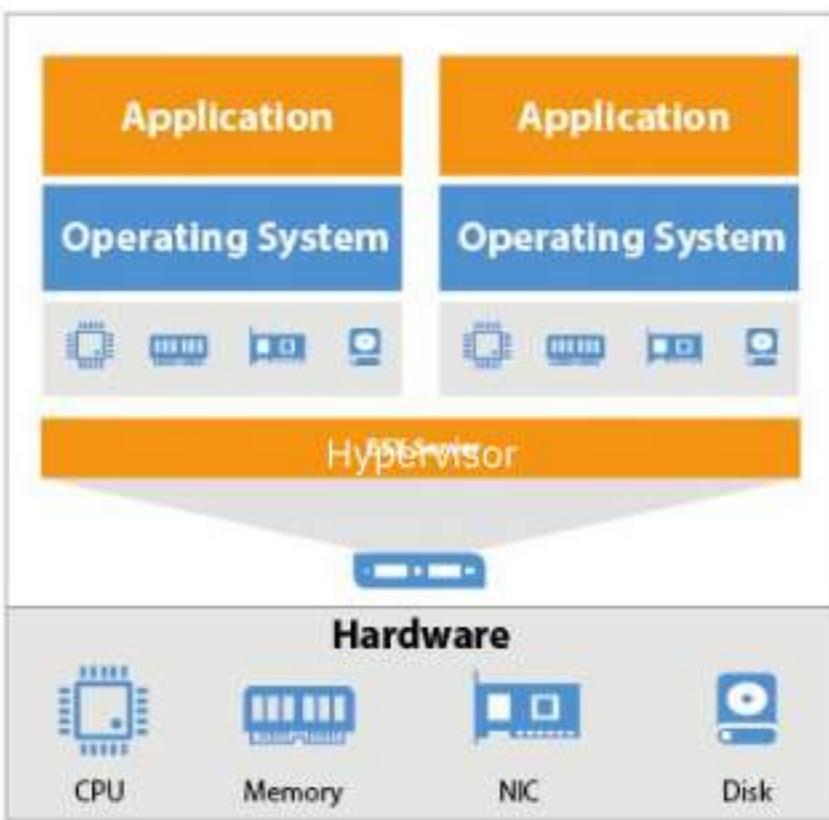
It's an open source container orchestration platform, to automate and manage containerized applications.

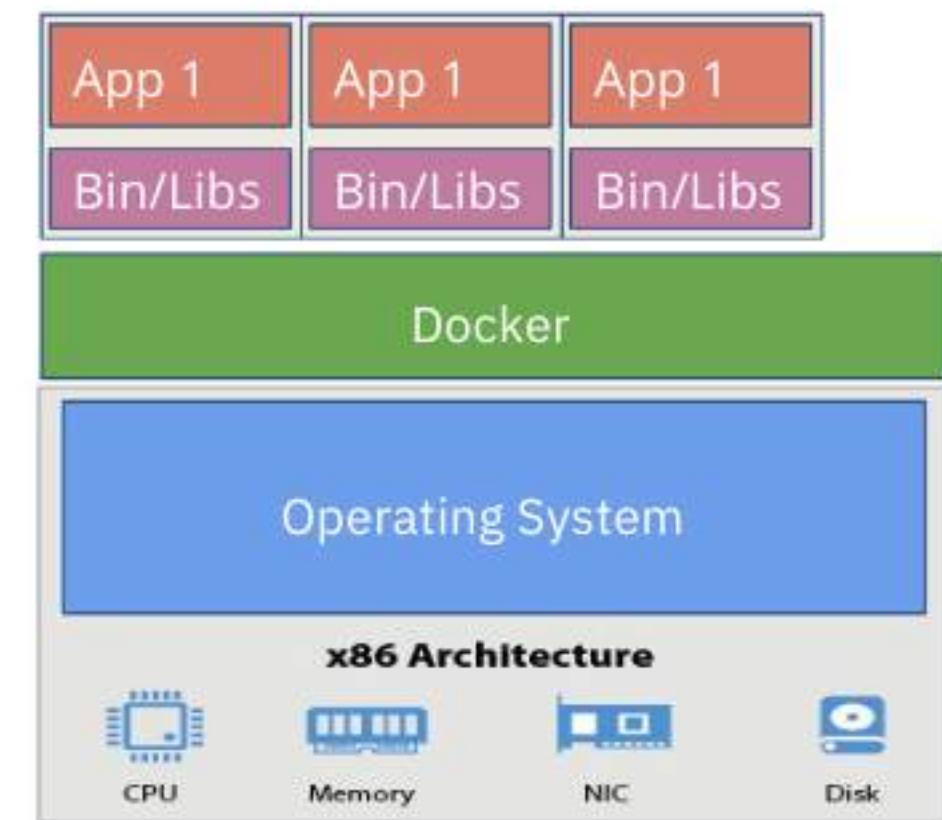
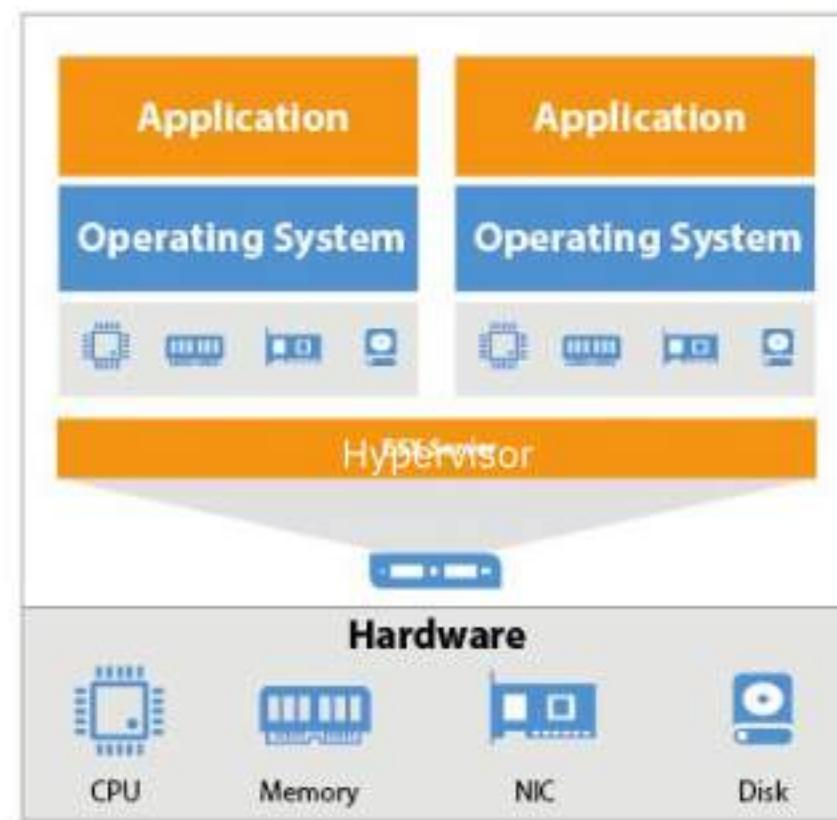




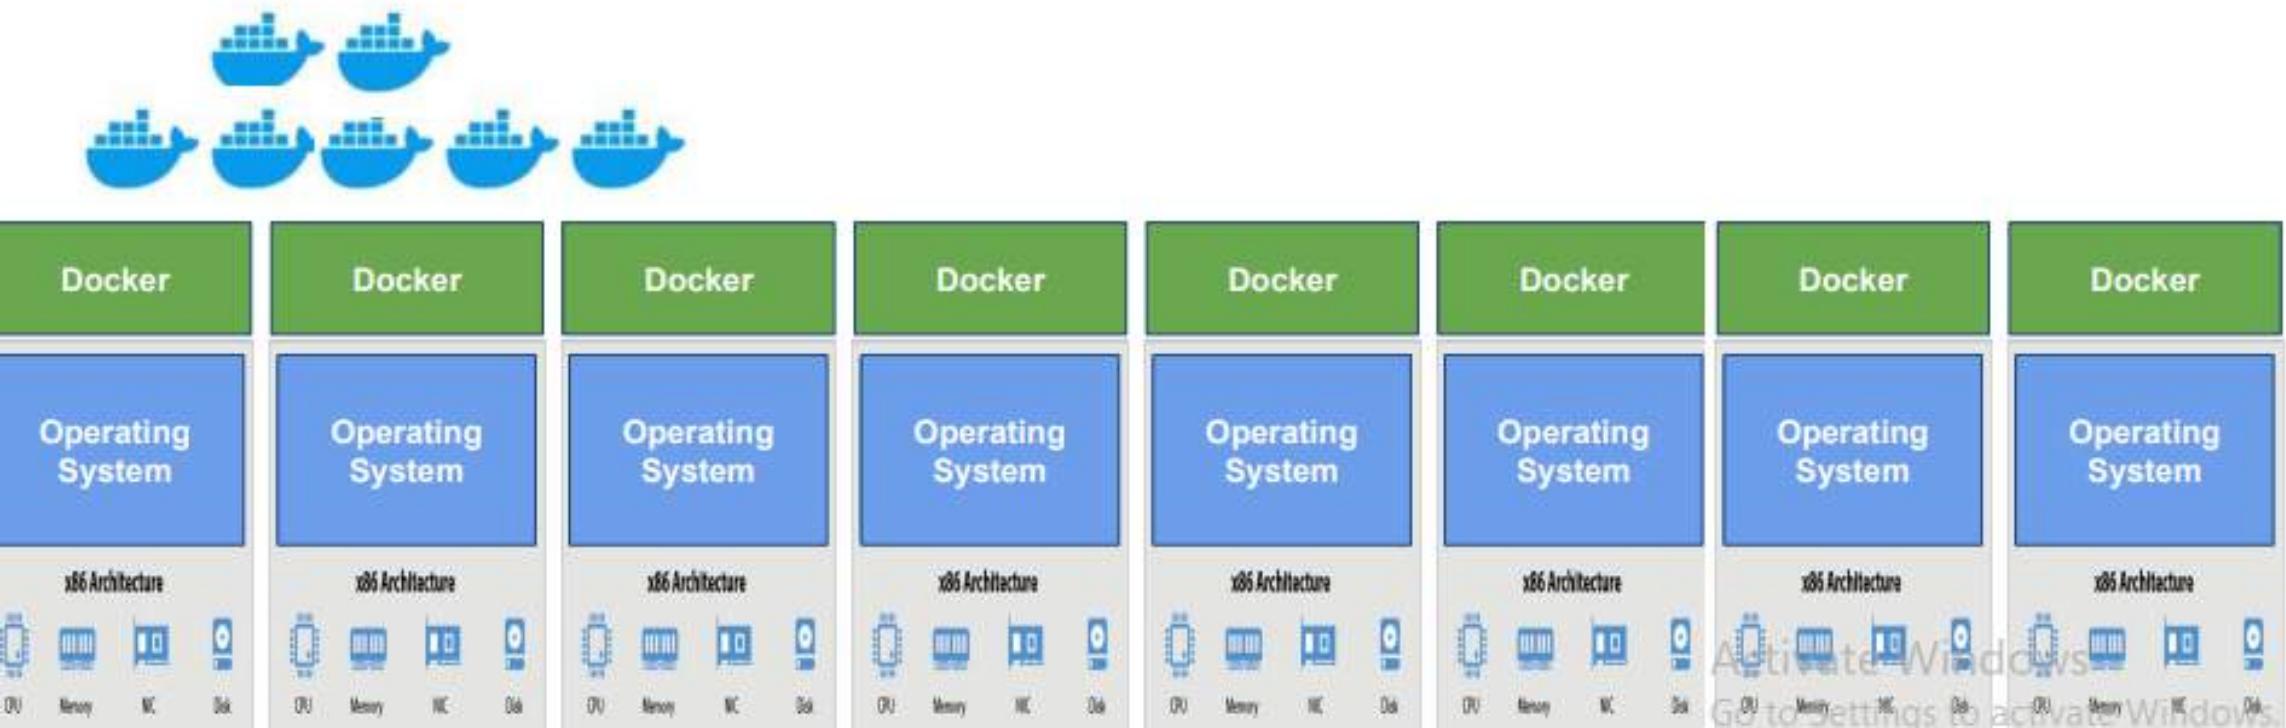




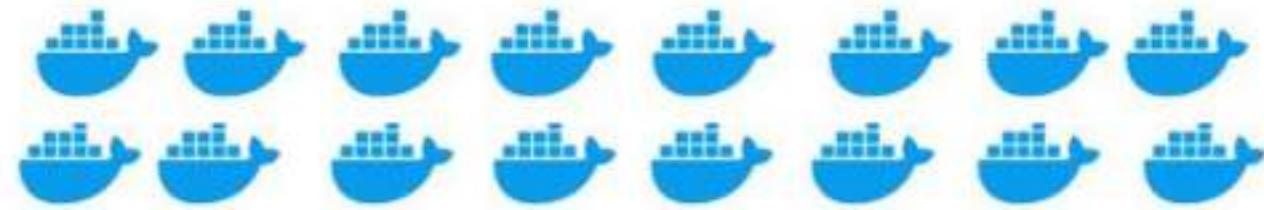
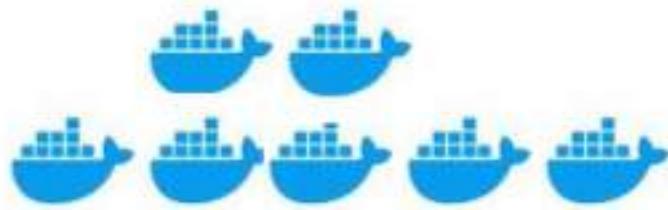












Docker

Operating System

Operating System
Operating System

Operating System
Operating System

Operating System
Operating System

Operating System
Operating System

Operating System
Operating System

Operating System
Operating System

Operating System
Operating System

Operating System

x86 Architecture

0

0

0

0

0

0

0

0

0

0

0

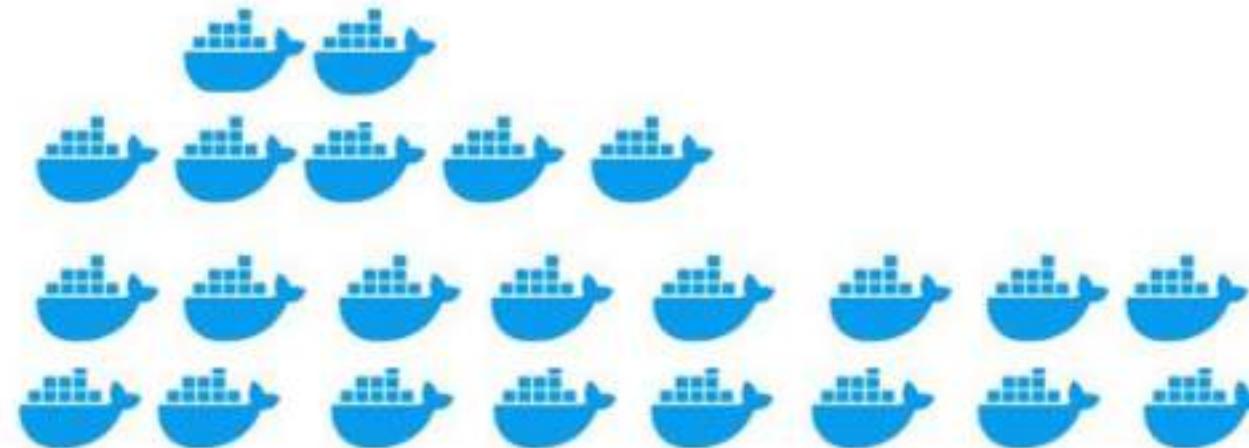
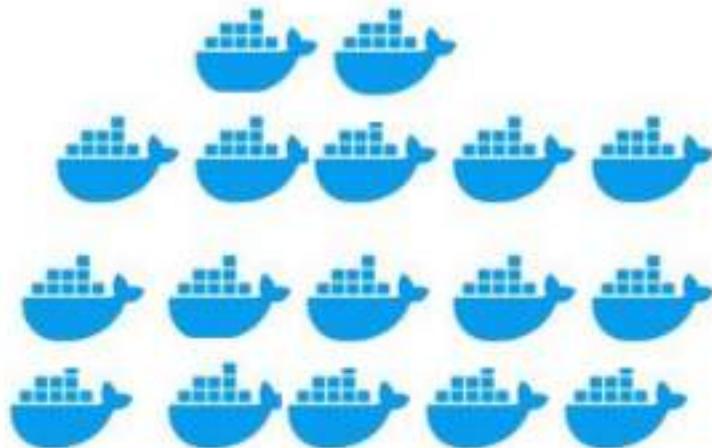
0

0

0

0

0



Docker

Operating System

x86 Architecture

Memory

Memory

Memory

Memory

Memory

Memory

Memory

Memory

Processor

Processor

Processor

Processor

Processor

Processor

Processor

Processor

Storage

Storage

Storage

Storage

Storage

Storage

Storage

Storage

Network

Network

Network

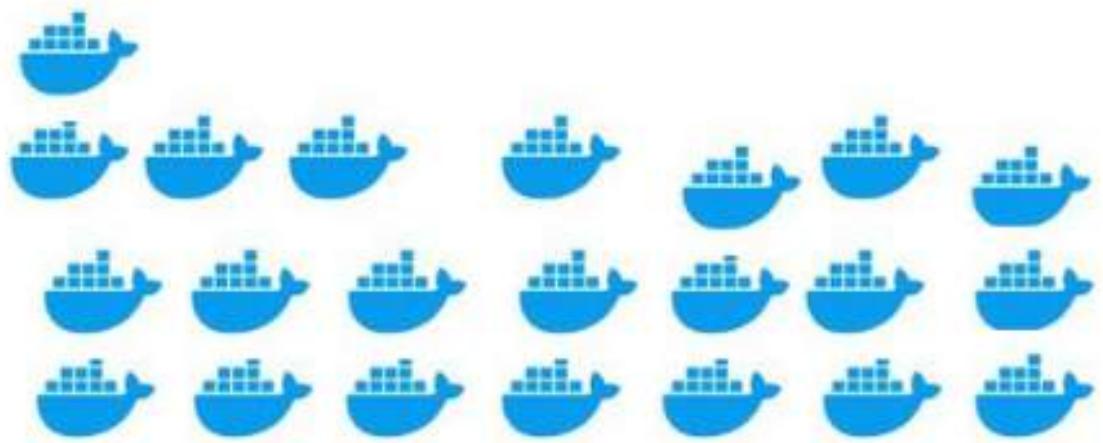
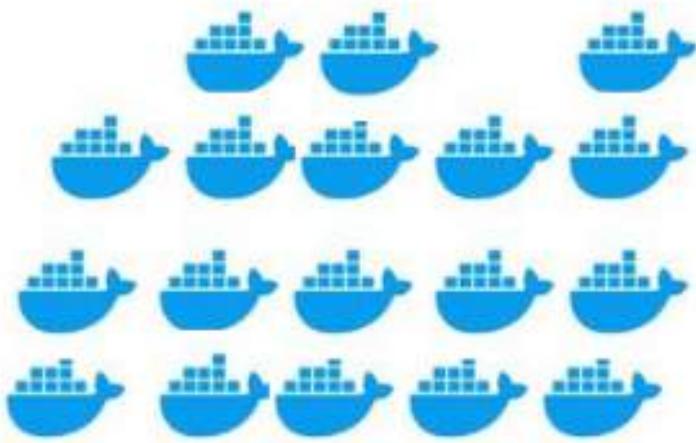
Network

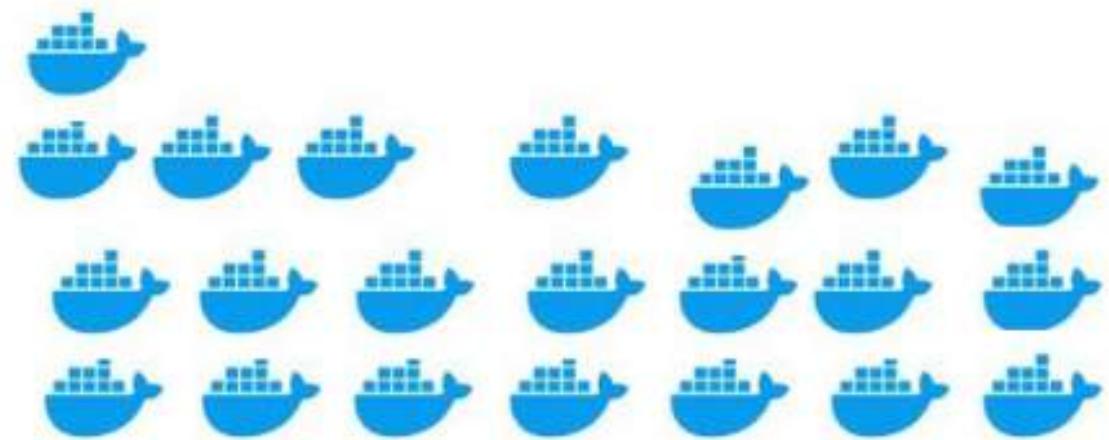
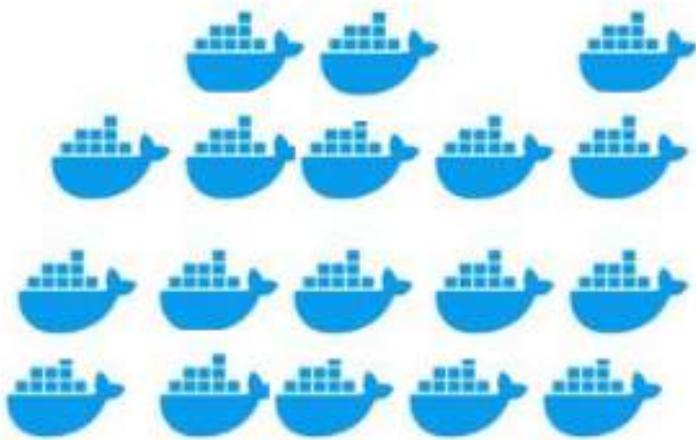
Network

Network

Network

Network





Docker

Operating System	Operating System
x86 Architecture	x86 Architecture
32	64

Operating System	Operating System
x86 Architecture	x86 Architecture
32	64

Operating System	Operating System
x86 Architecture	x86 Architecture
32	64

Operating System	Operating System
x86 Architecture	x86 Architecture
32	64

Operating System	Operating System
x86 Architecture	x86 Architecture
32	64

Operating System	Operating System
x86 Architecture	x86 Architecture
32	64

Operating System	Operating System
x86 Architecture	x86 Architecture
32	64

Operating System	Operating System
x86 Architecture	x86 Architecture
32	64

Final Note

01

What does Kubernetes Provide for the Infrastructure

Automated Infrastructure

Compute Networking Storage
Firewall Monitoring Alerting

02

What does Kubernetes Provide for the Developer

Automated Deployment

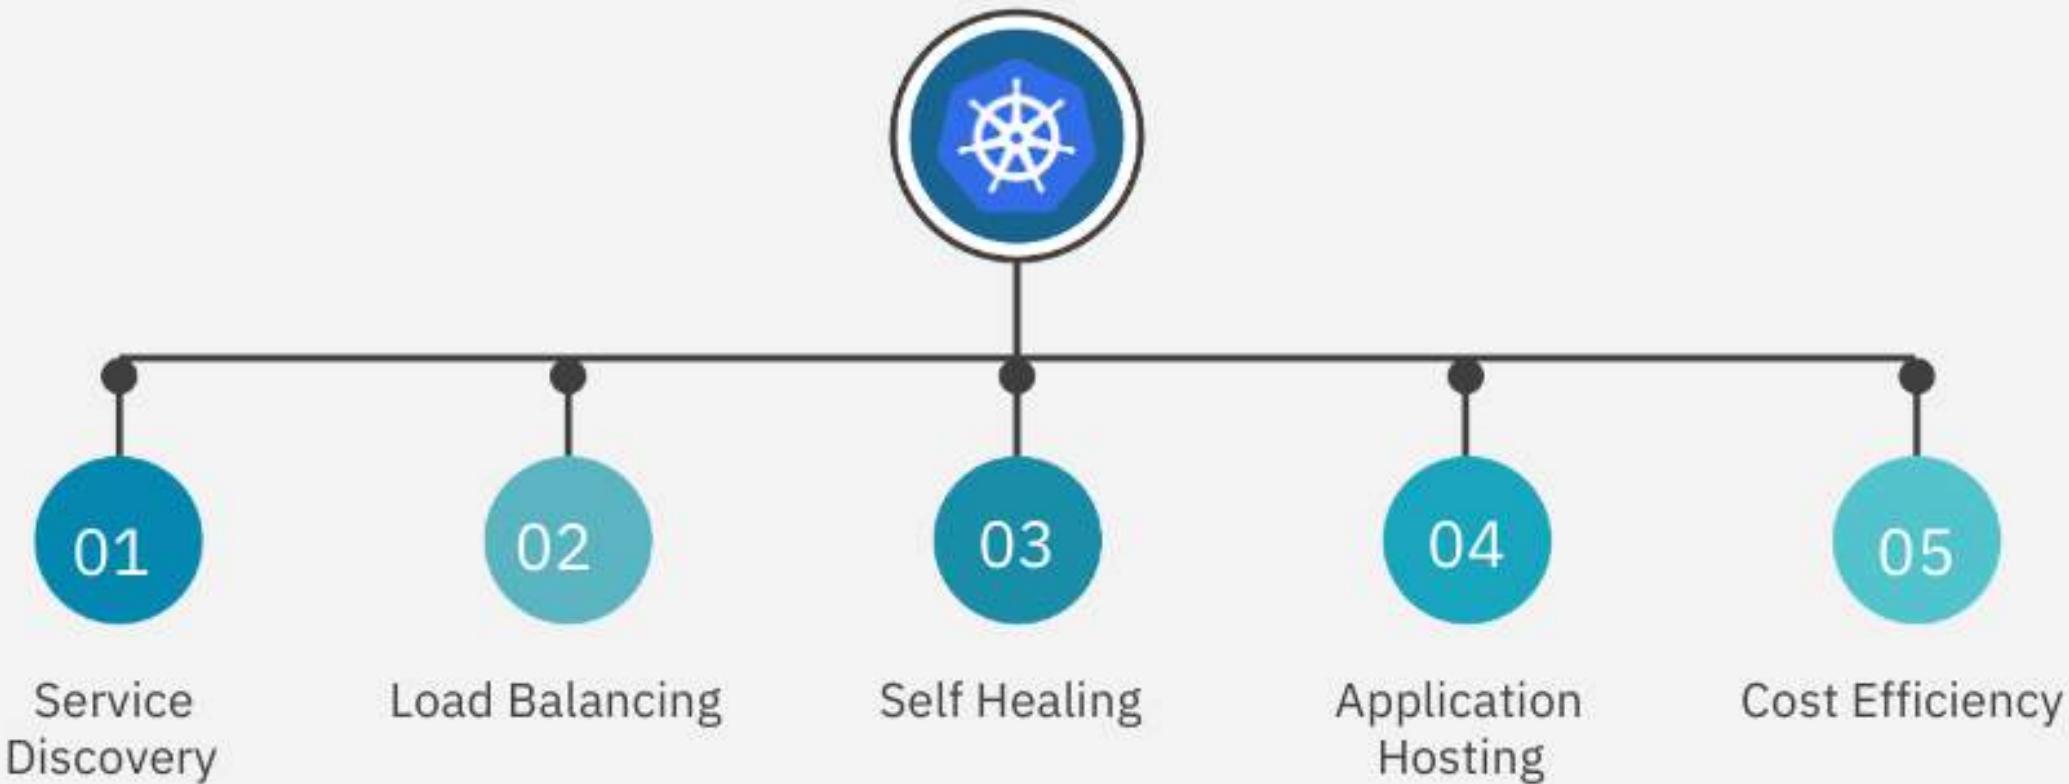
Deployment Self Healing Load
Balancing Service Discovery
Monitoring Logging Alerting

The background of the slide features a dark blue gradient with a faint, abstract technical illustration. It includes several binary code numbers (e.g., 1001001011010101) scattered across the top and right side. In the center-left, there are three interlocking circular gears of varying sizes. A yellow steering wheel icon is positioned inside the largest gear. The overall aesthetic is futuristic and technical.

Kubernetes From Scratch - Hands-on

03 - Kubernetes Features

Kubernetes Features



Kubernetes Features

01

Service Discovery

A way to connect different services or components with a single stable endpoint

02

Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.

03

Self Healing

Kubernetes ensures that the Desired state and Actual state are always in sync by monitoring the container healthcheck and restarting the containers which fail this healthcheck.

04

Application Hosting

Kubernetes hosts containerized applications. So the developers doesn't care whether kubernetes is running on bare metal or on the cloud

Kubernetes Features

05

Faster Deployment

Kubernetes allows you to deliver a self-service Platform-as-a-Service (PaaS) that creates a hardware layer abstraction for development teams.

06

Fine-grained Access Control

Kubernetes RBAC allows the cluster maintainers to configure fine-grained access control to Kubernetes resources

07

Batch Jobs

Kubernetes supports running automated jobs with the CronJob object. They are similar to the Unix/Linux Cron tasks.

08

Autoscaling

One of the cool features of Kubernetes is that it can monitor your workload and scale it up or down based on the CPU utilization or memory consumption

Kubernetes Features

09

Cost Efficiency

Kubernetes optimizes infrastructure cost by efficiently managing the infrastructure resources
Reduction of administration time, etc.

10

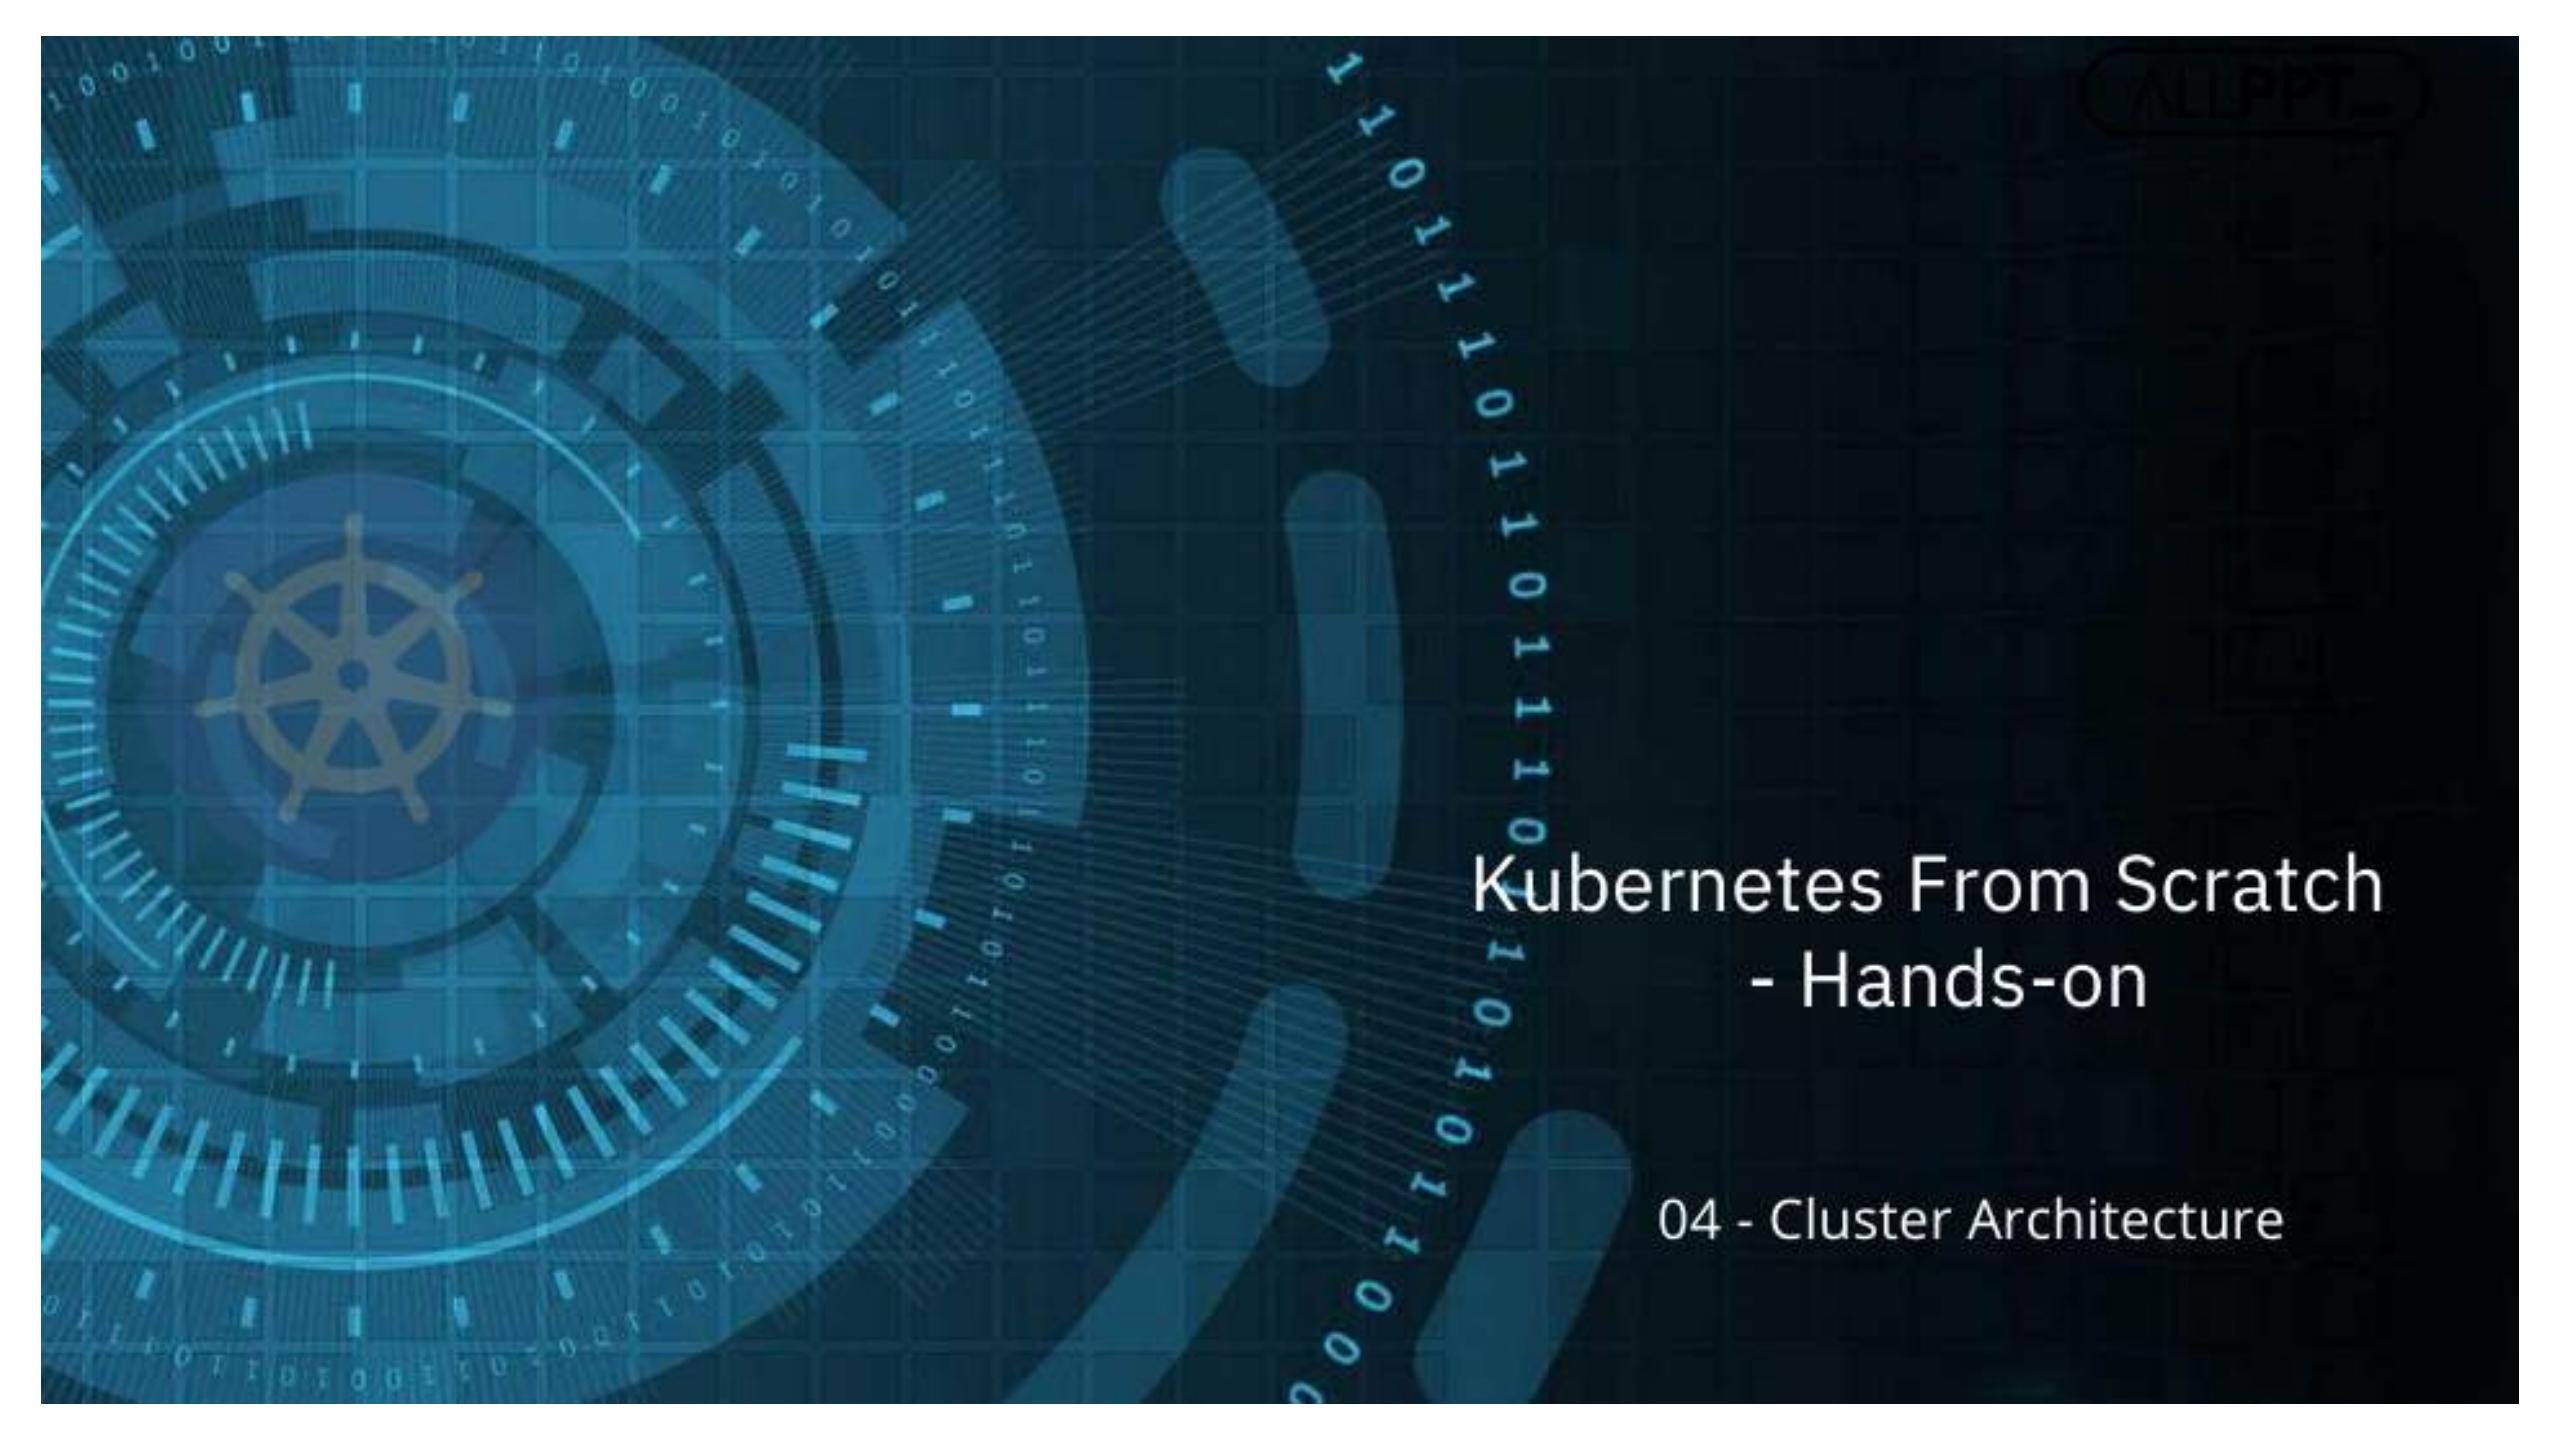
Declarative Configuration

We provide the API server with the manifest files (YAML or JSON) describing how we want the Kubernetes cluster to look.



Module 2

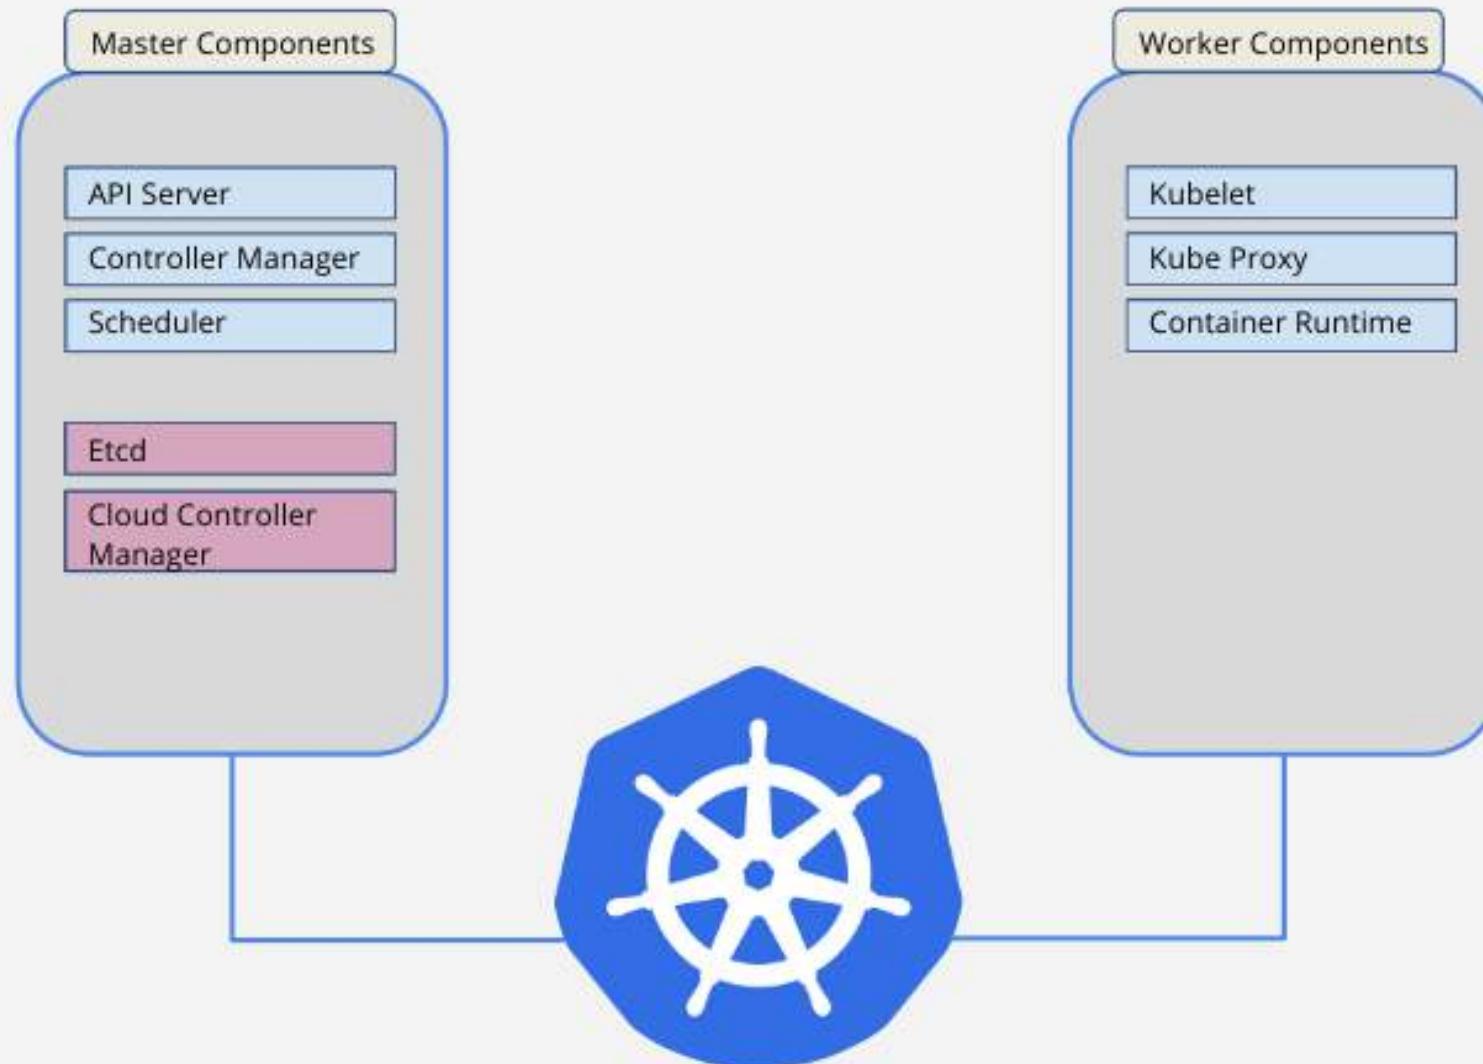
Kubernetes Architecture



Kubernetes From Scratch - Hands-on

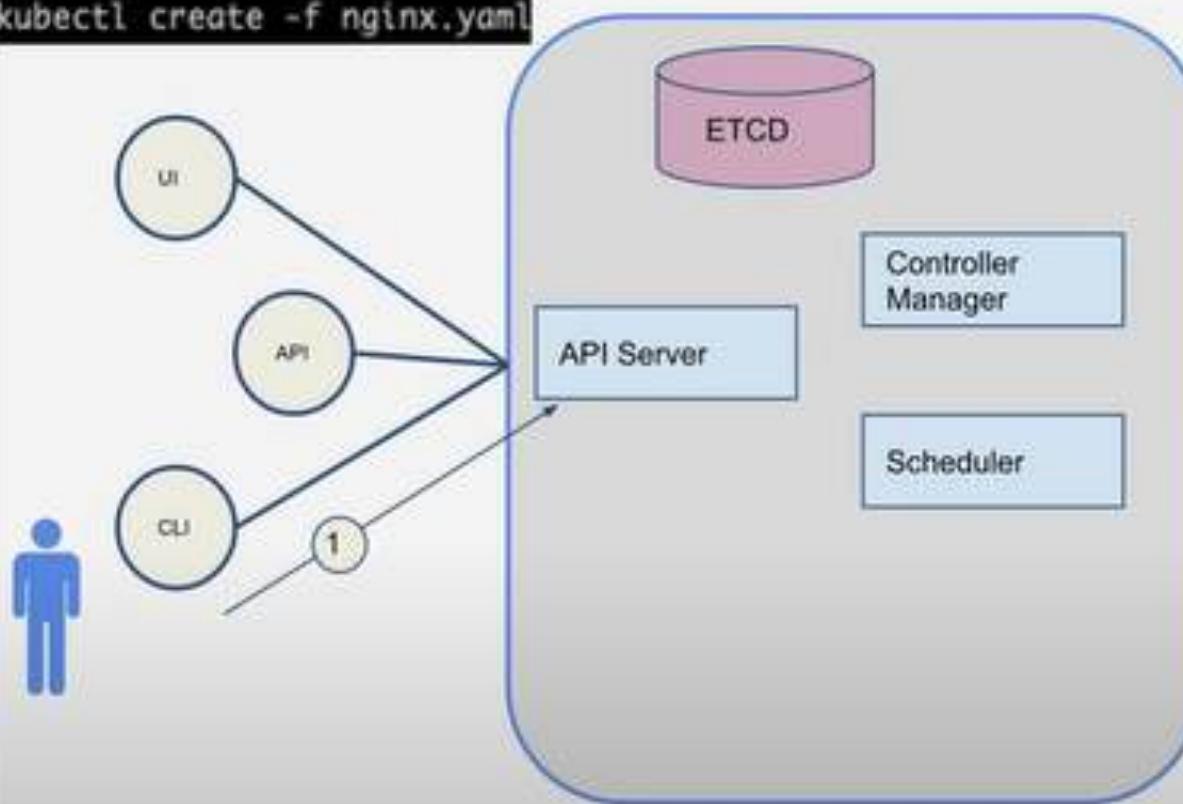
04 - Cluster Architecture

Cluster Architecture



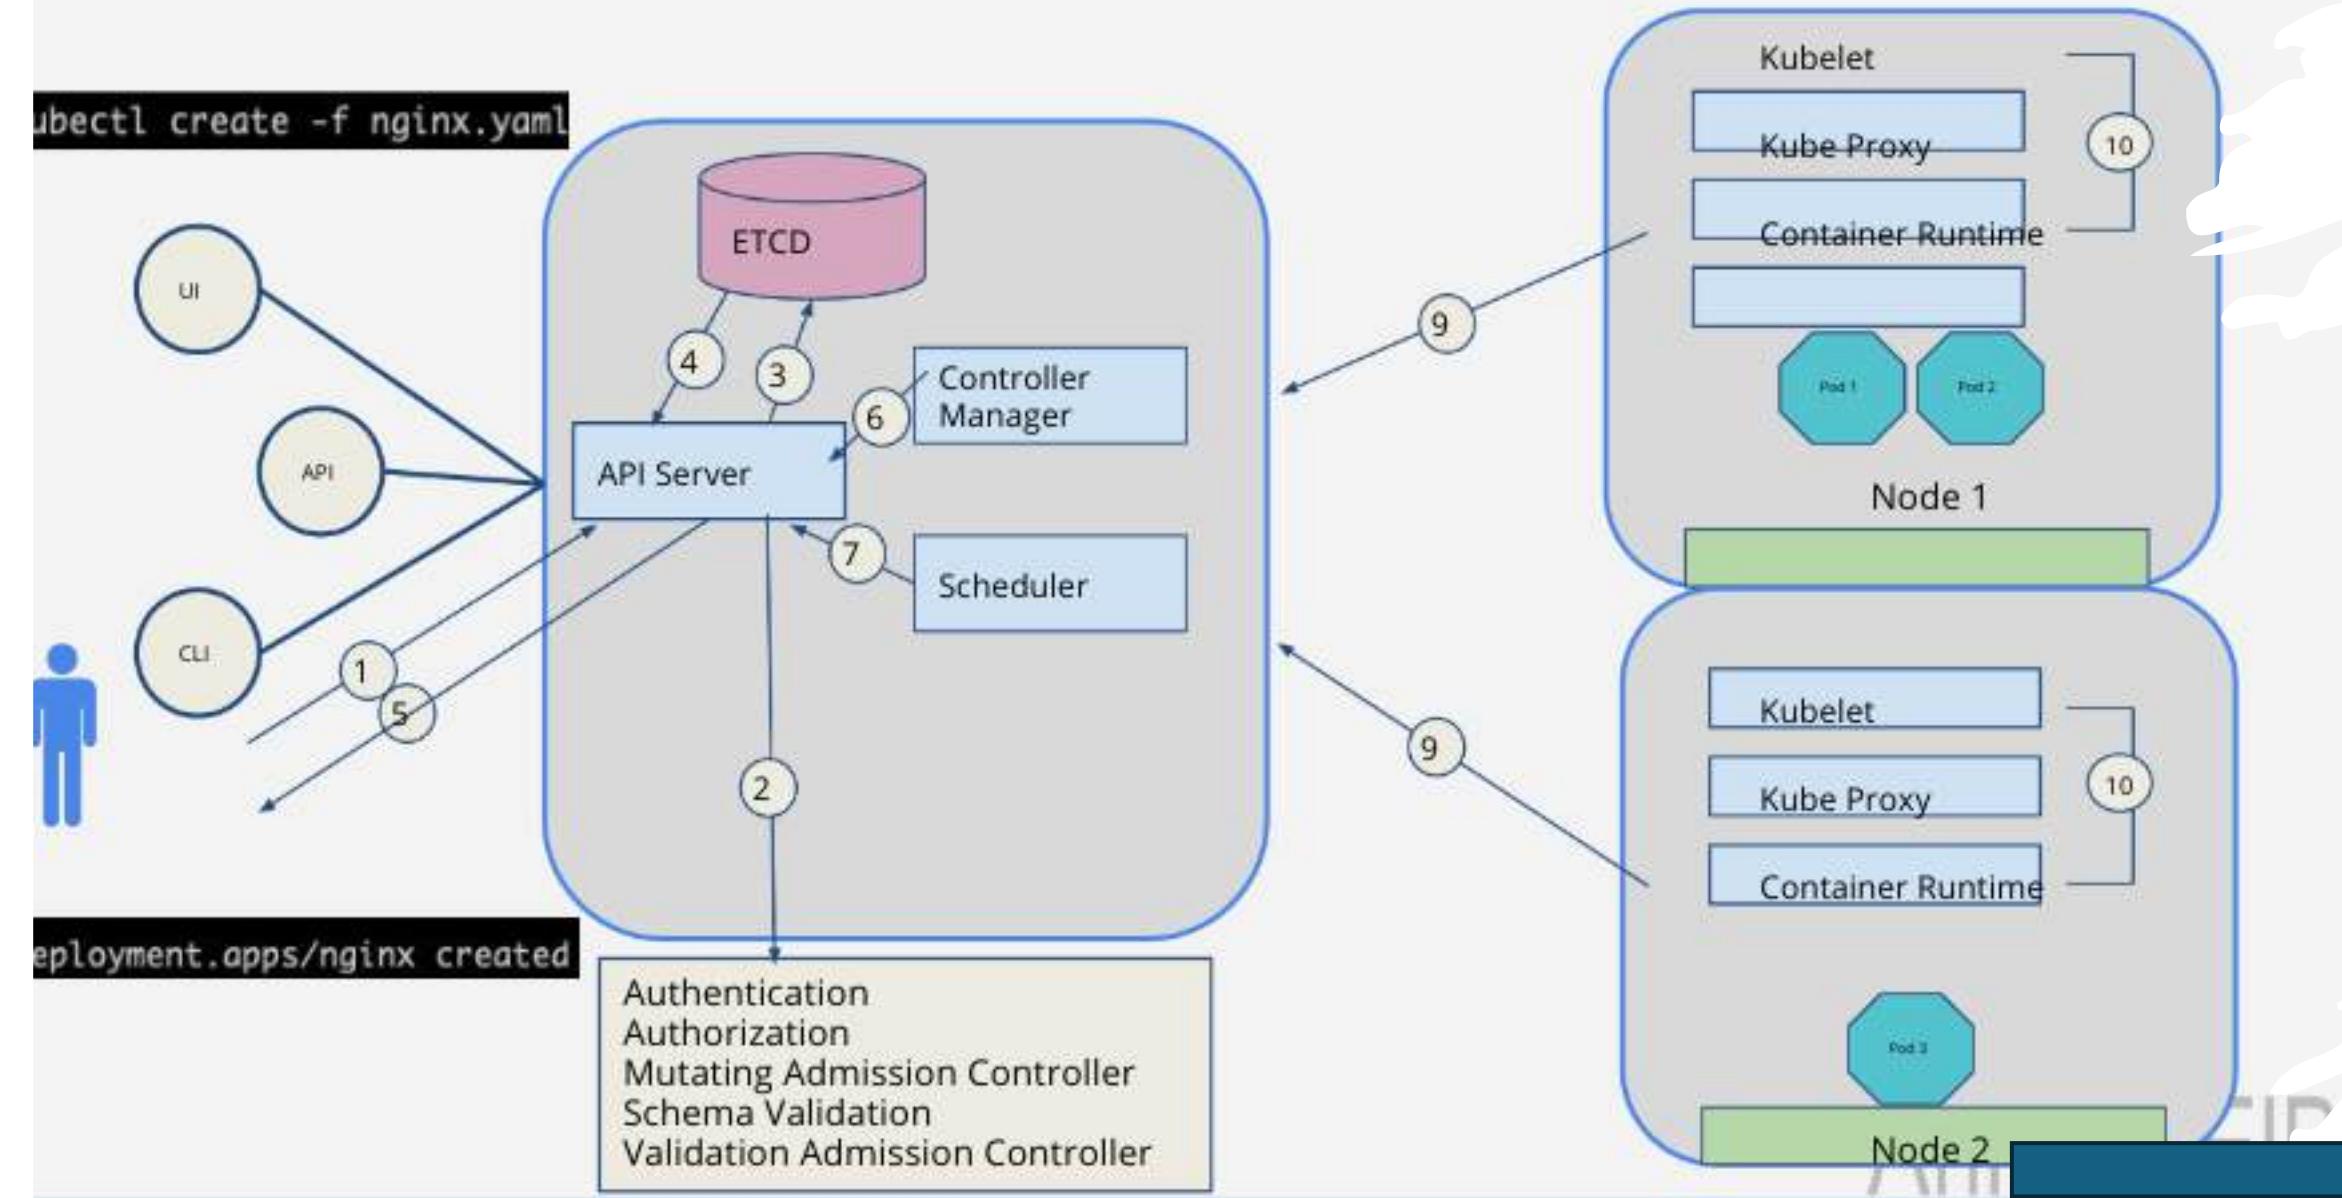
Cluster Architecture

```
kubectl create -f nginx.yaml
```



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.14.2
      ports:
```

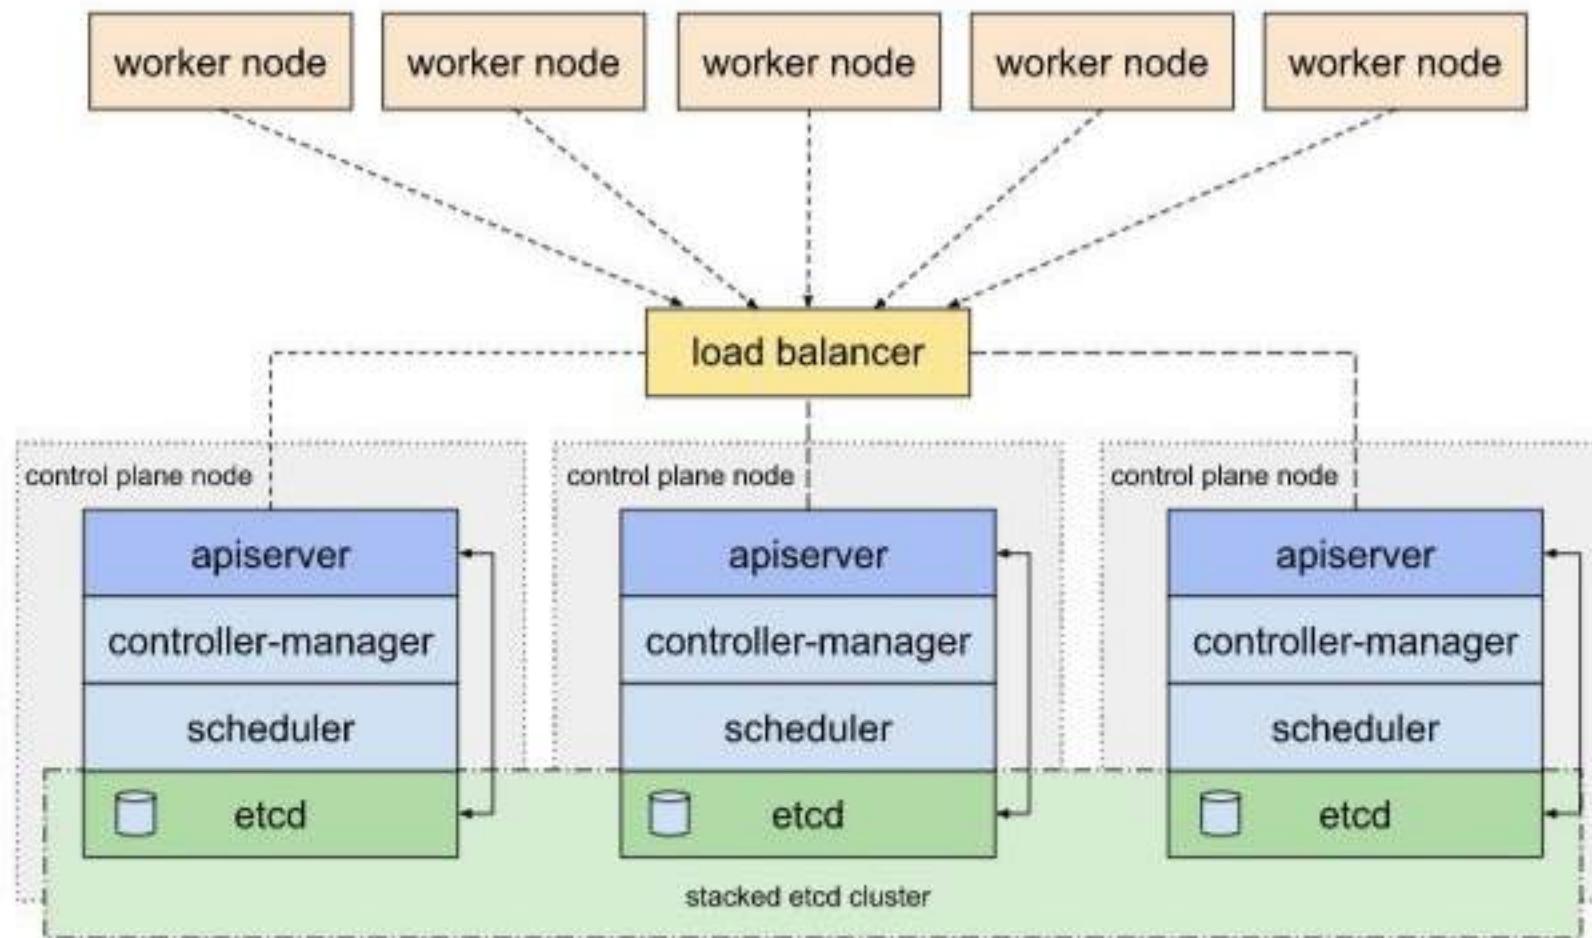
Cluster Architecture



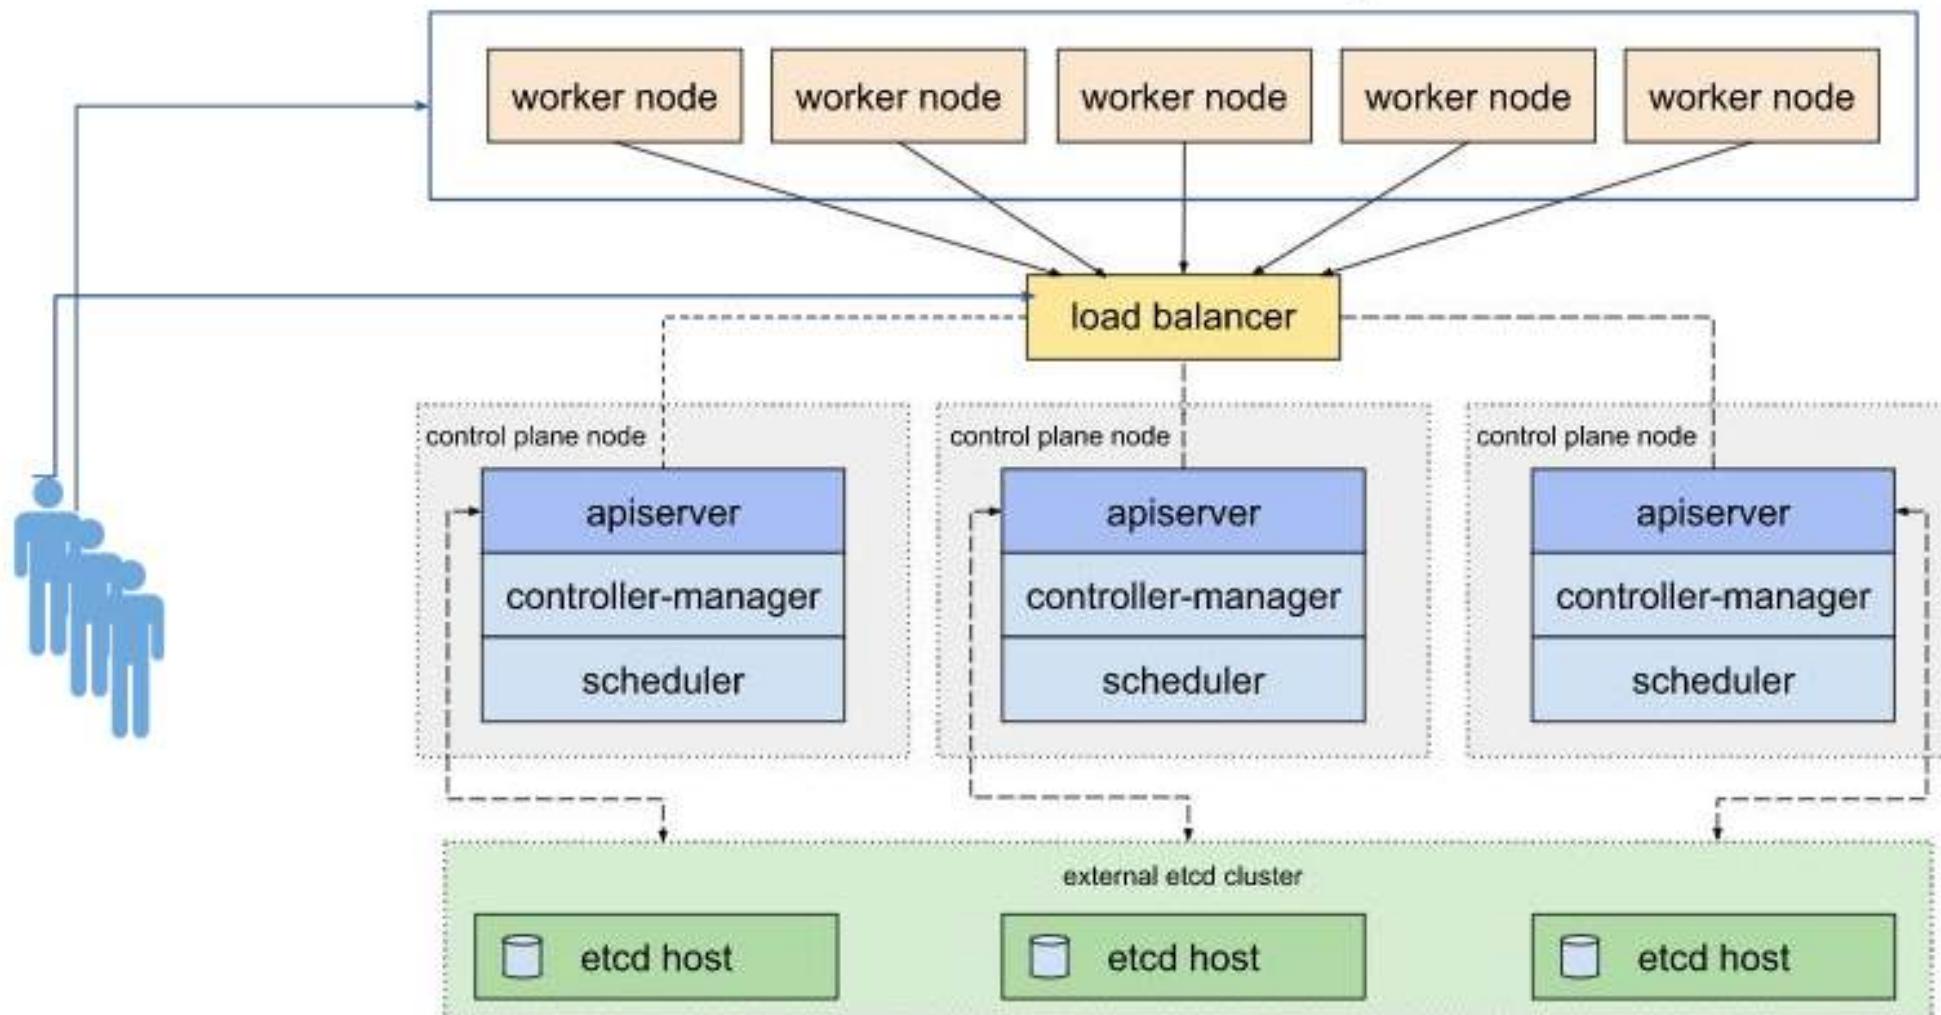


THANK YOU

Master Components

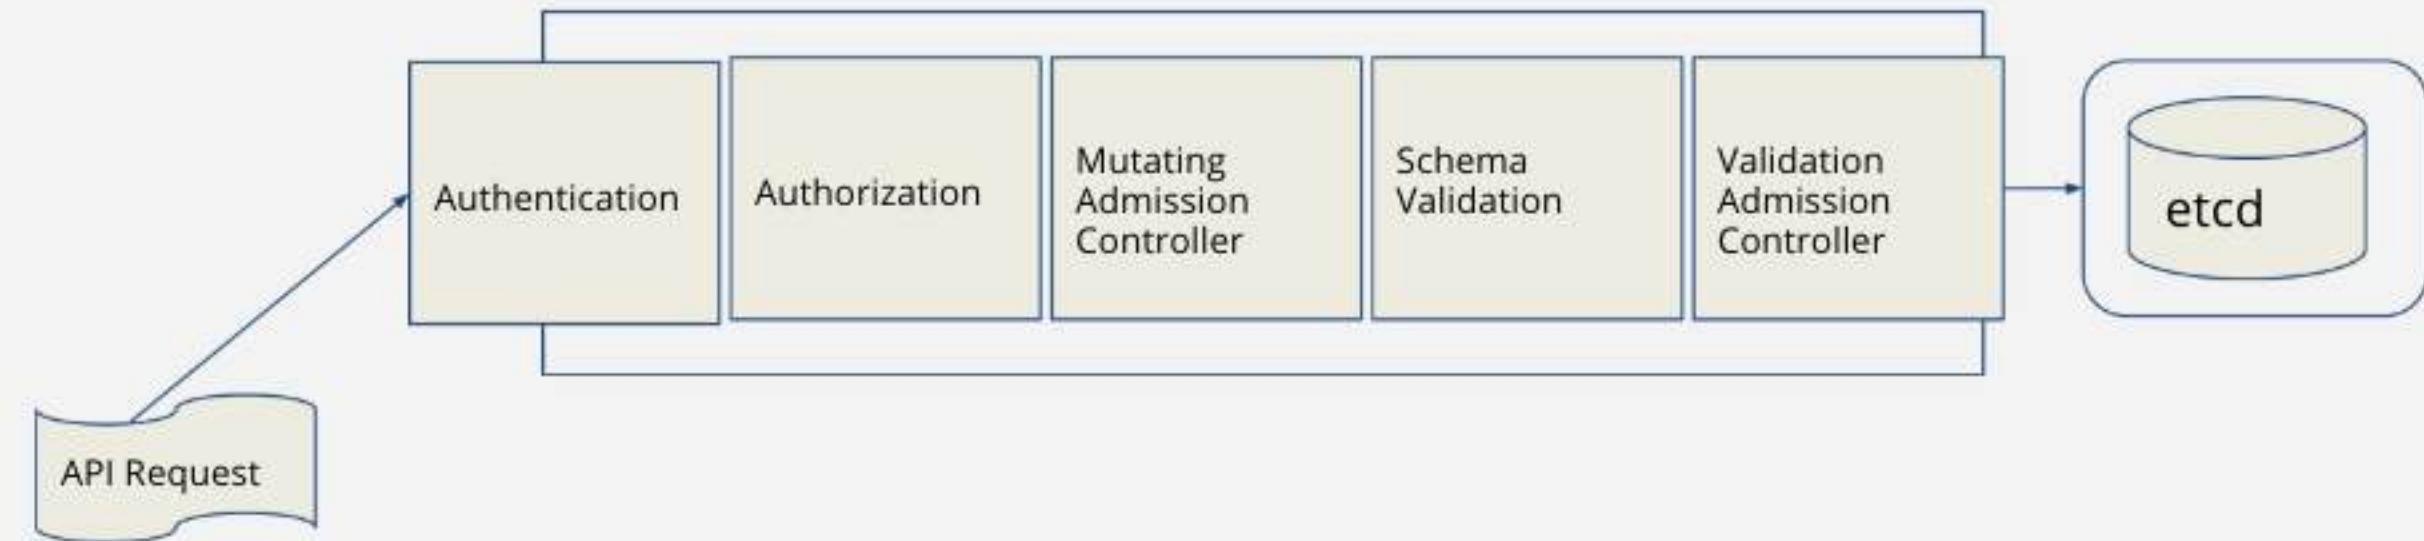


Master Components



Master Components

Authentication, Authorization, and Admission Control

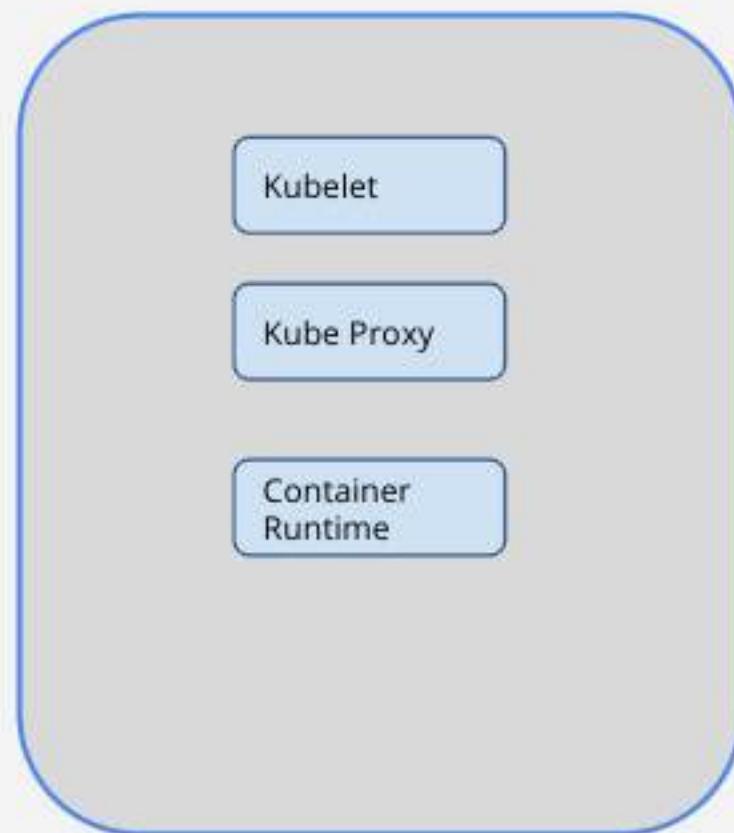


The background of the slide features a dark blue gradient with a faint, abstract watermark. This watermark consists of several large, semi-transparent binary numbers ('0's and '1's) scattered across the frame, along with stylized icons of a steering wheel and interlocking gears, all set against a light blue grid.

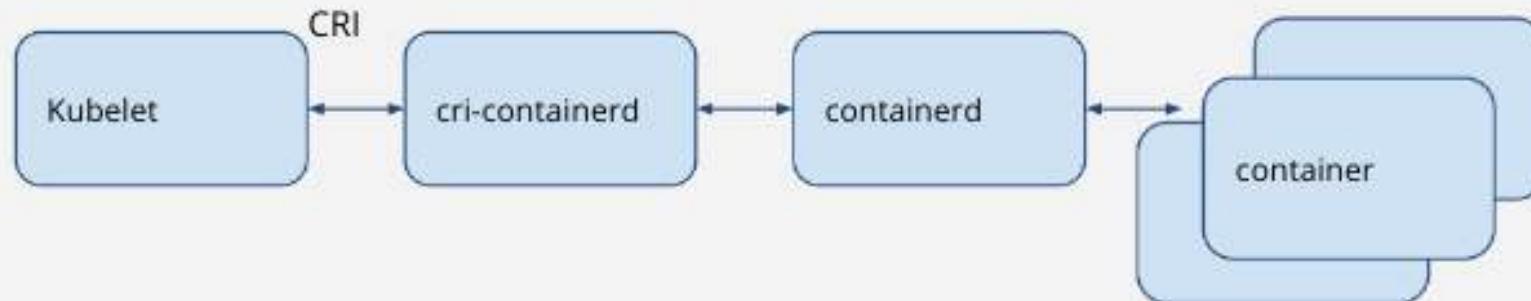
Kubernetes From Scratch - Hands-on

06 - Worker Components

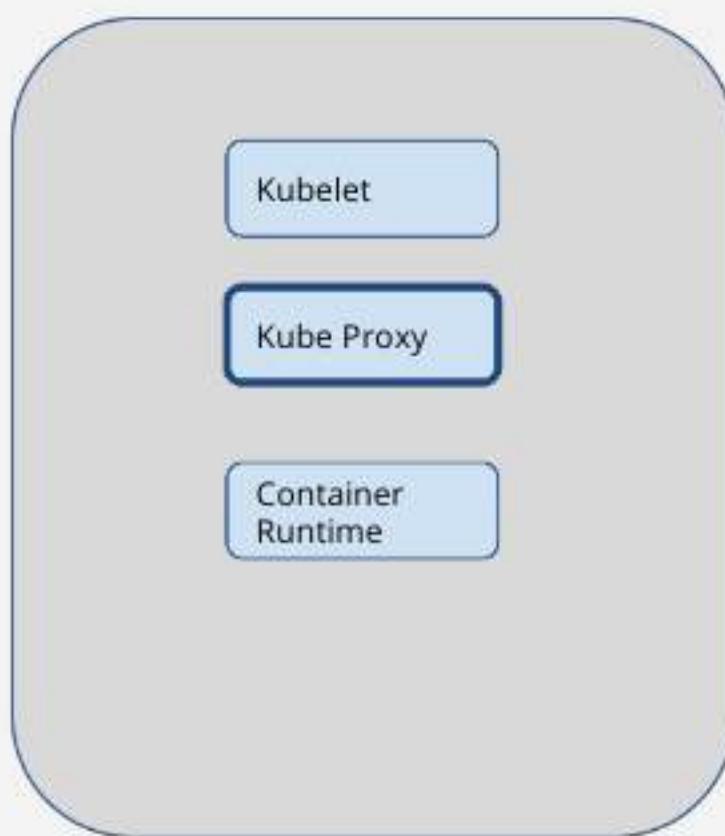
Worker Components



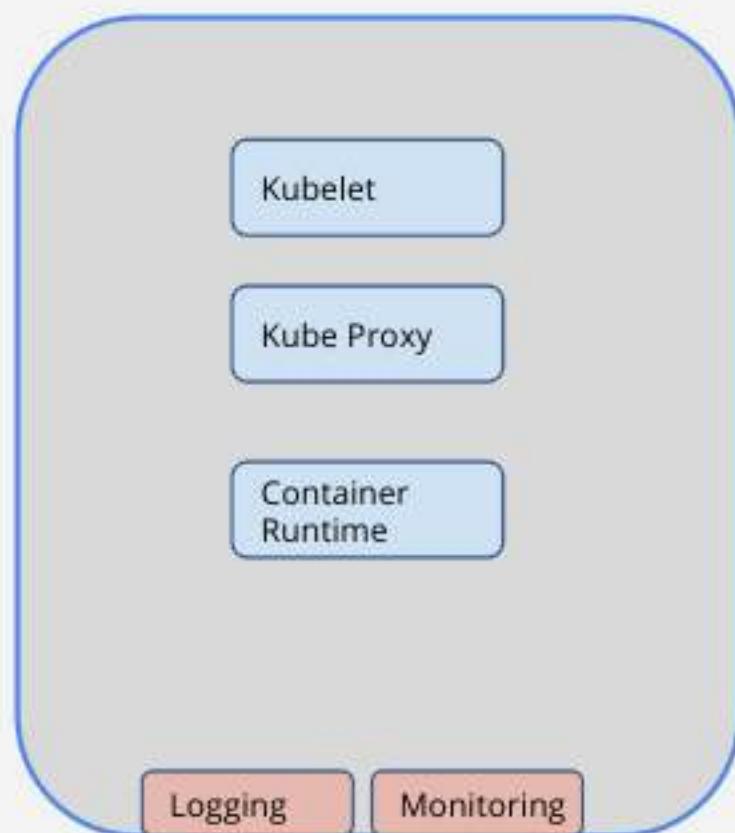
Worker Components



Worker Components



Worker Components





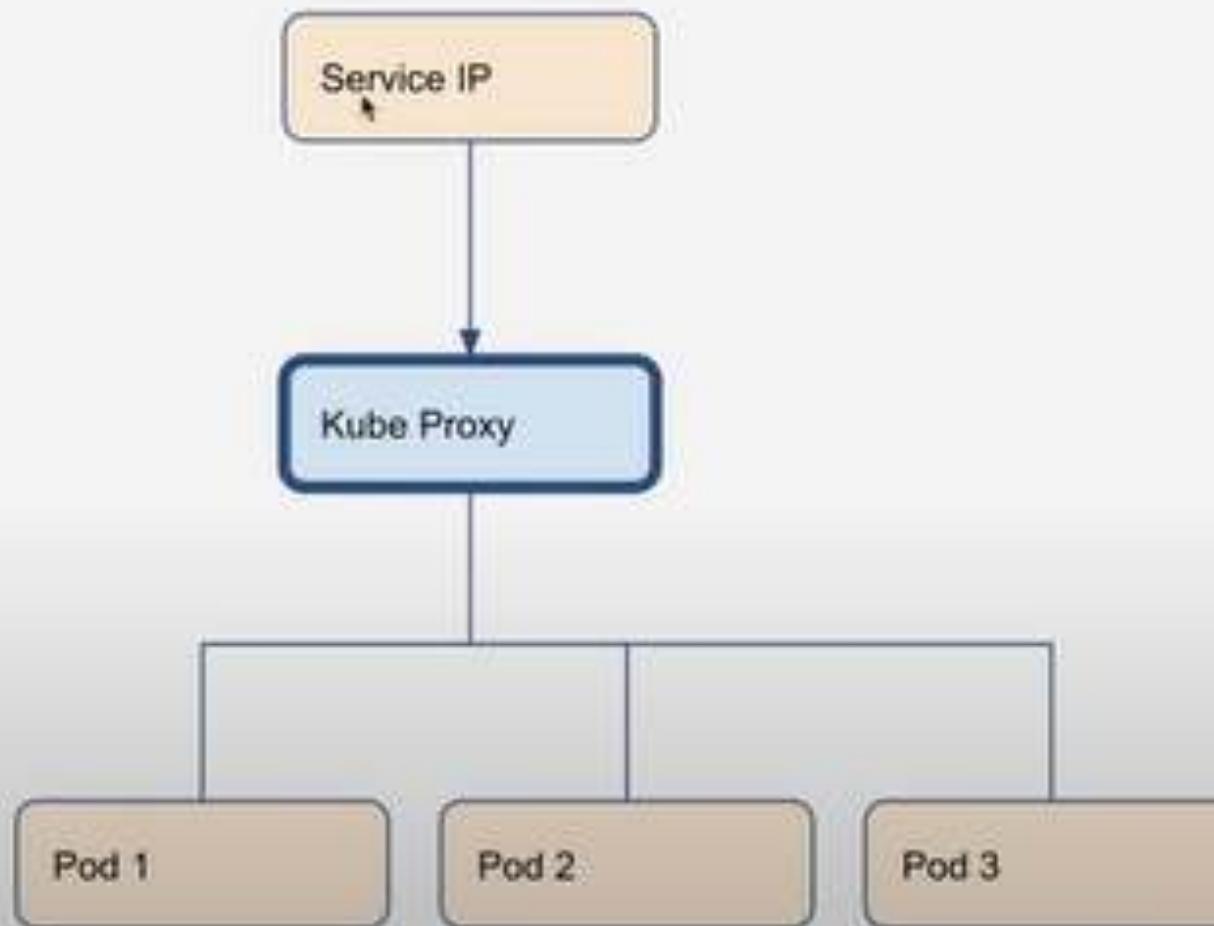
THANK YOU



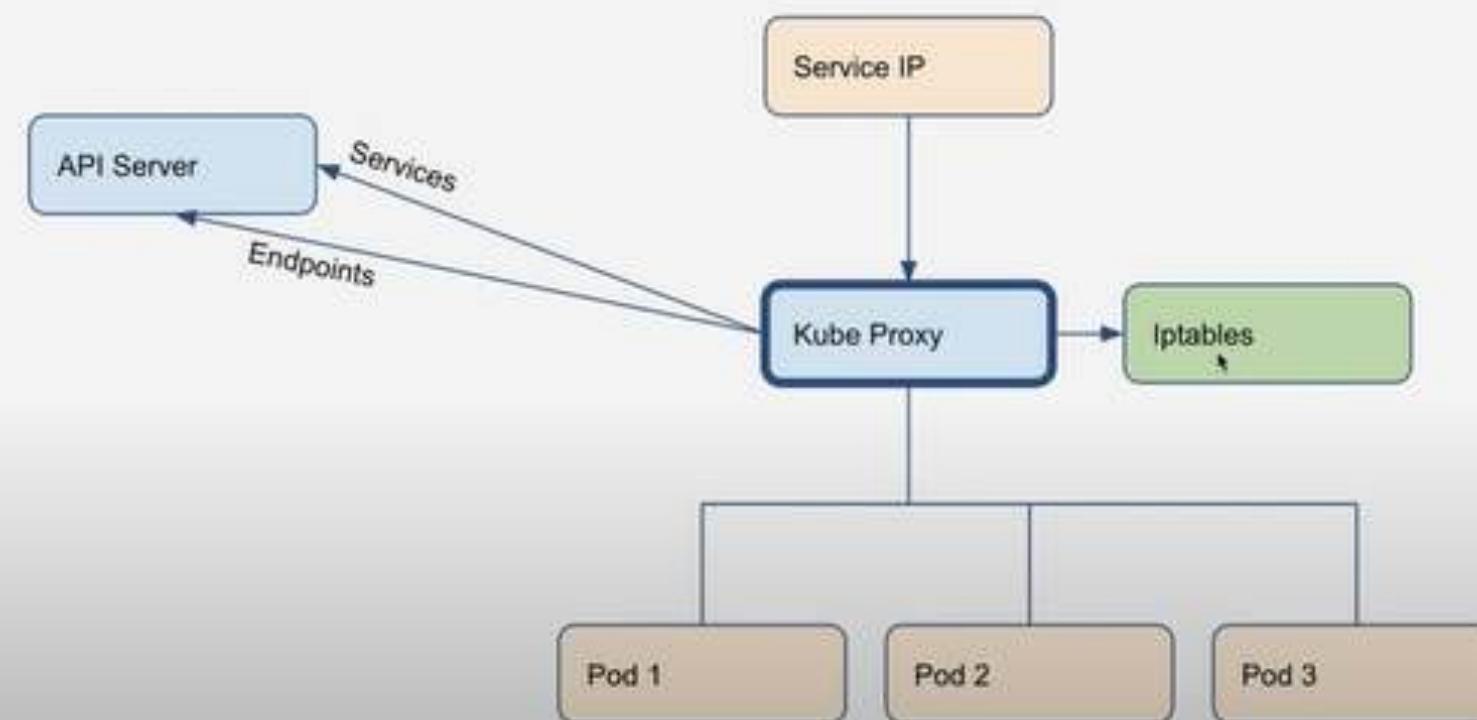
Kubernetes From Scratch - Hands-on

07 - Cluster Architecture Demo

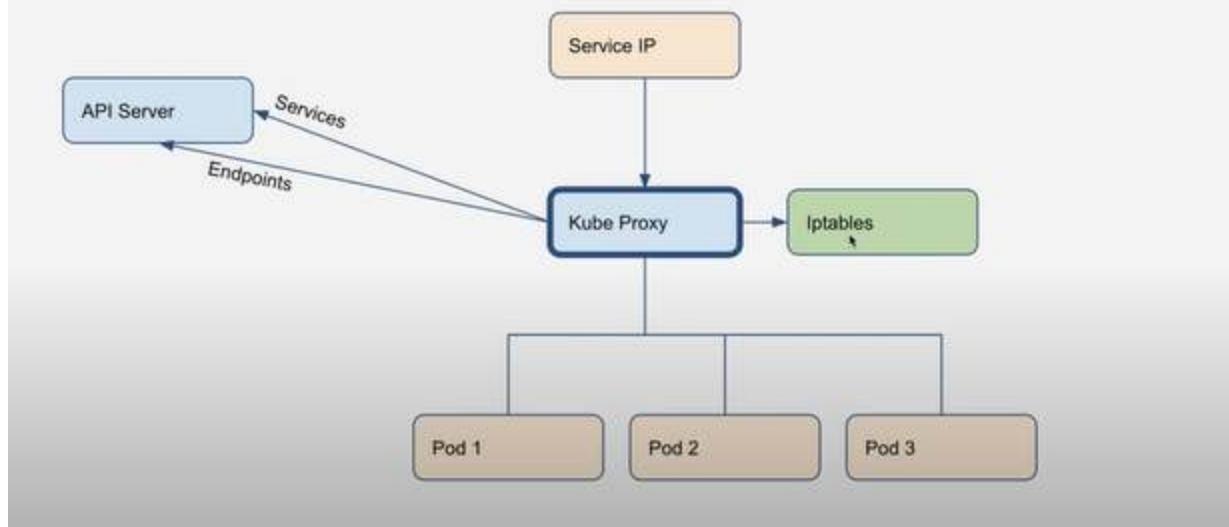
Worker Components



Worker Components



Worker Components



What is Minikube

- Minikube is a tool that makes it easy to run Kubernetes locally
- Minikube runs a single-node Kubernetes cluster inside a Virtual Machine (VM)
- Supports the following features
 - DNS
 - NodePorts
 - ConfigMaps and Secrets
 - Dashboards
 - Container Runtime: Docker, CRI-O, and containerd
 - Enabling CNI (Container Network Interface)
 - Ingress
- Requirements:
 - 2 CPU or more
 - 2GB of free memory
 - 20GB of free disk space
 - Internet connection
 - Container or virtual machine manager, such as: Docker, Hyperkit, Hyper-V, KVM, Parallels, Podman, VirtualBox, or VMWare



Minikube

Master processes

- ❖ api-server
- ❖ controller-manager
- ❖ scheduler
- ❖ etcd

Worker processes

- ❖ kube-proxy
- ❖ kubelet
- ❖ container runtime

What is Kind

- kind is a tool for running local Kubernetes clusters using Docker container
- “nodes”. kind supports Linux, macOS and Windows Run multiple clusters on different versions. Can be used for
- - Testing Kubernetes upgrades
 - Testing components with Kubernetes new versions
 - Testing with CI/CD

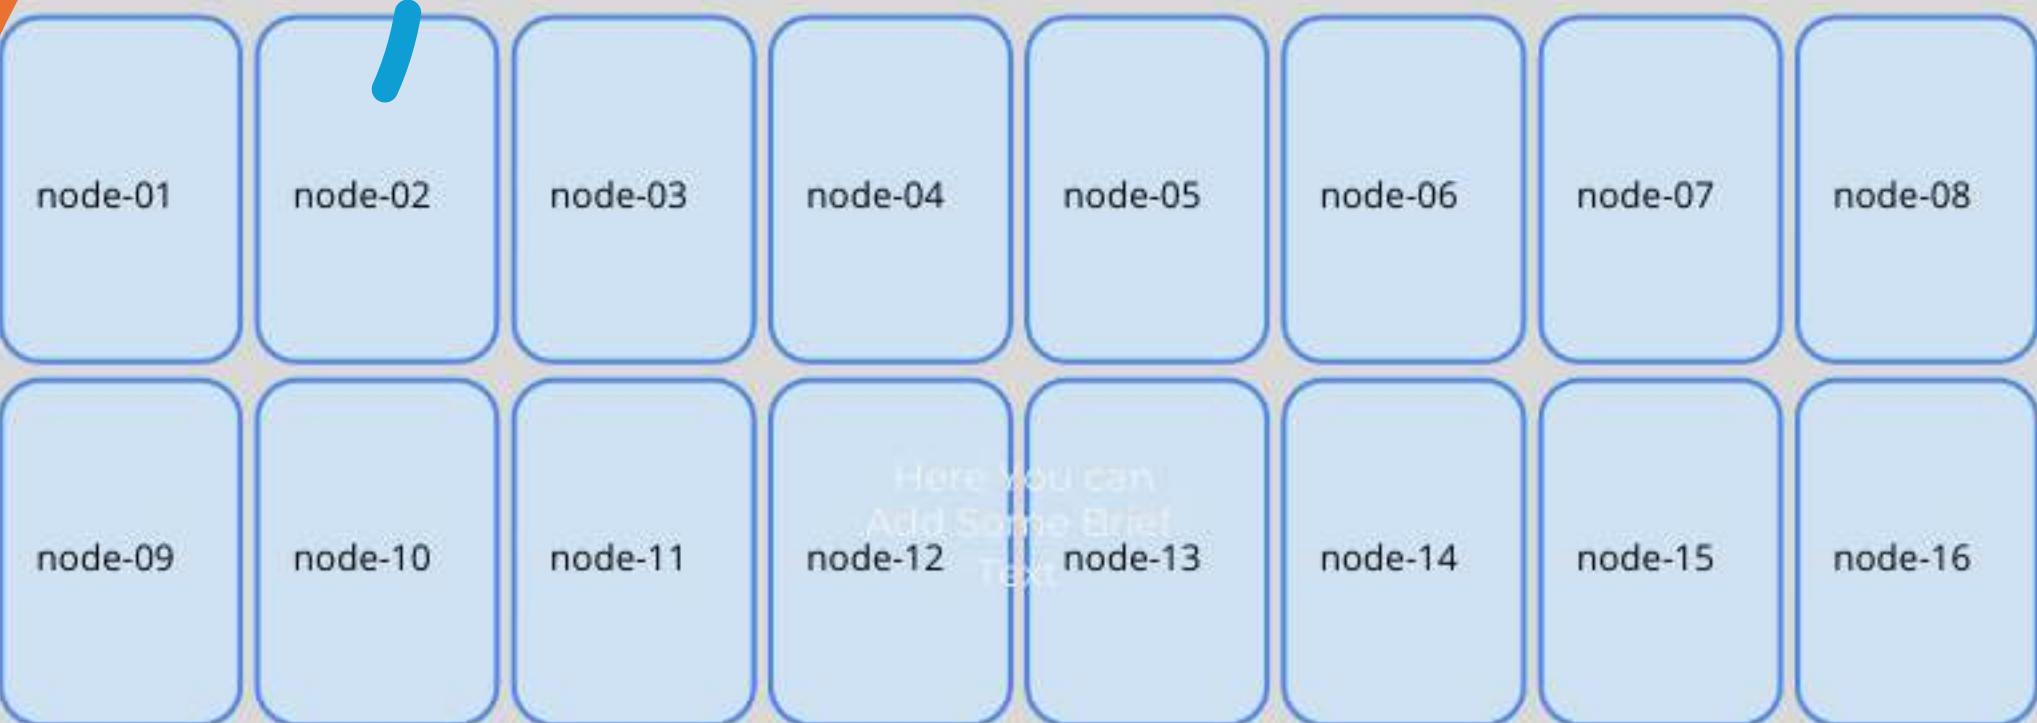


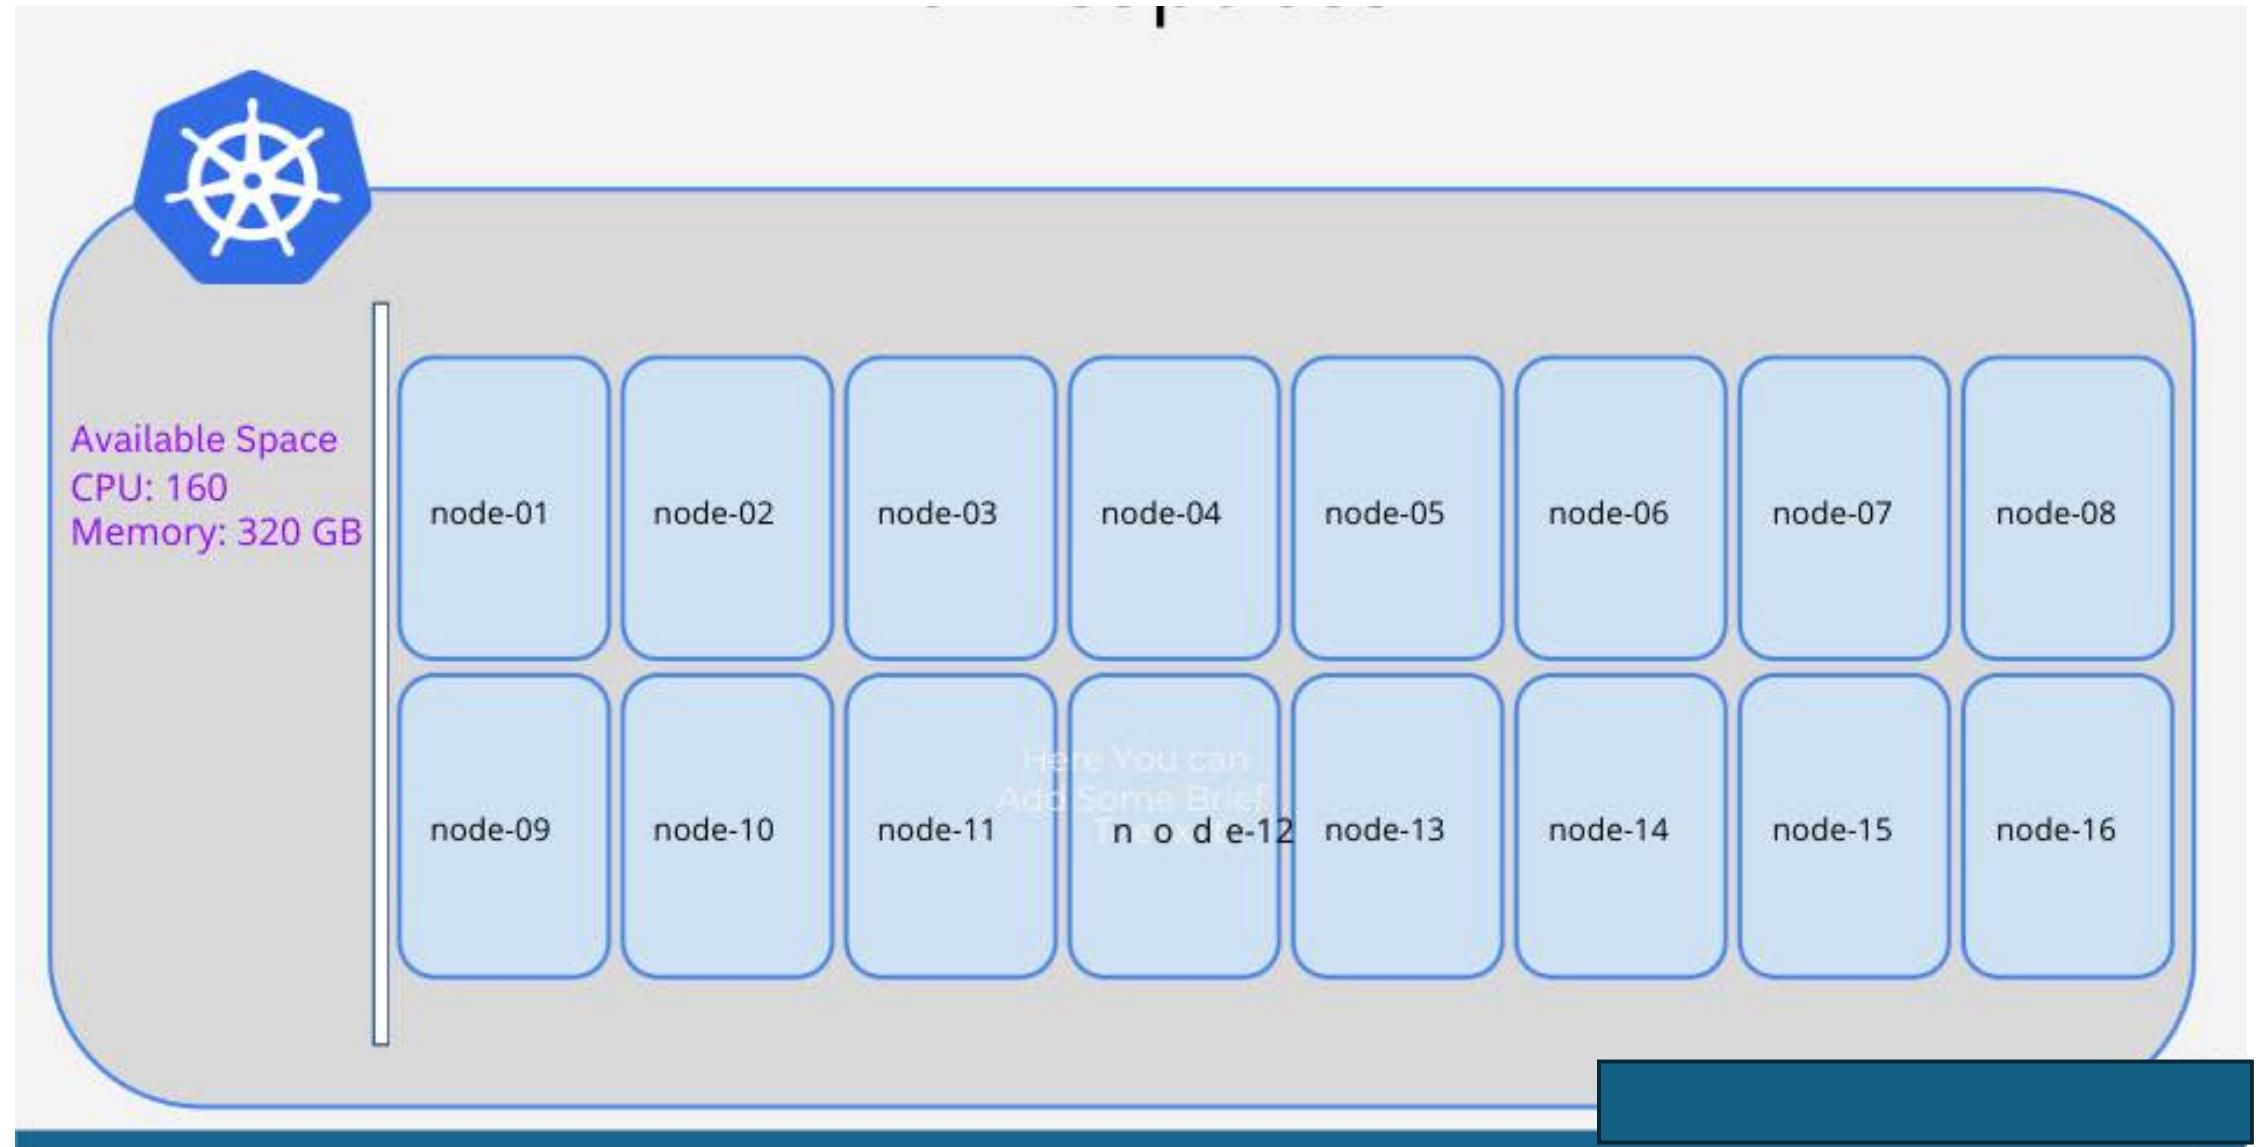
Namespaces



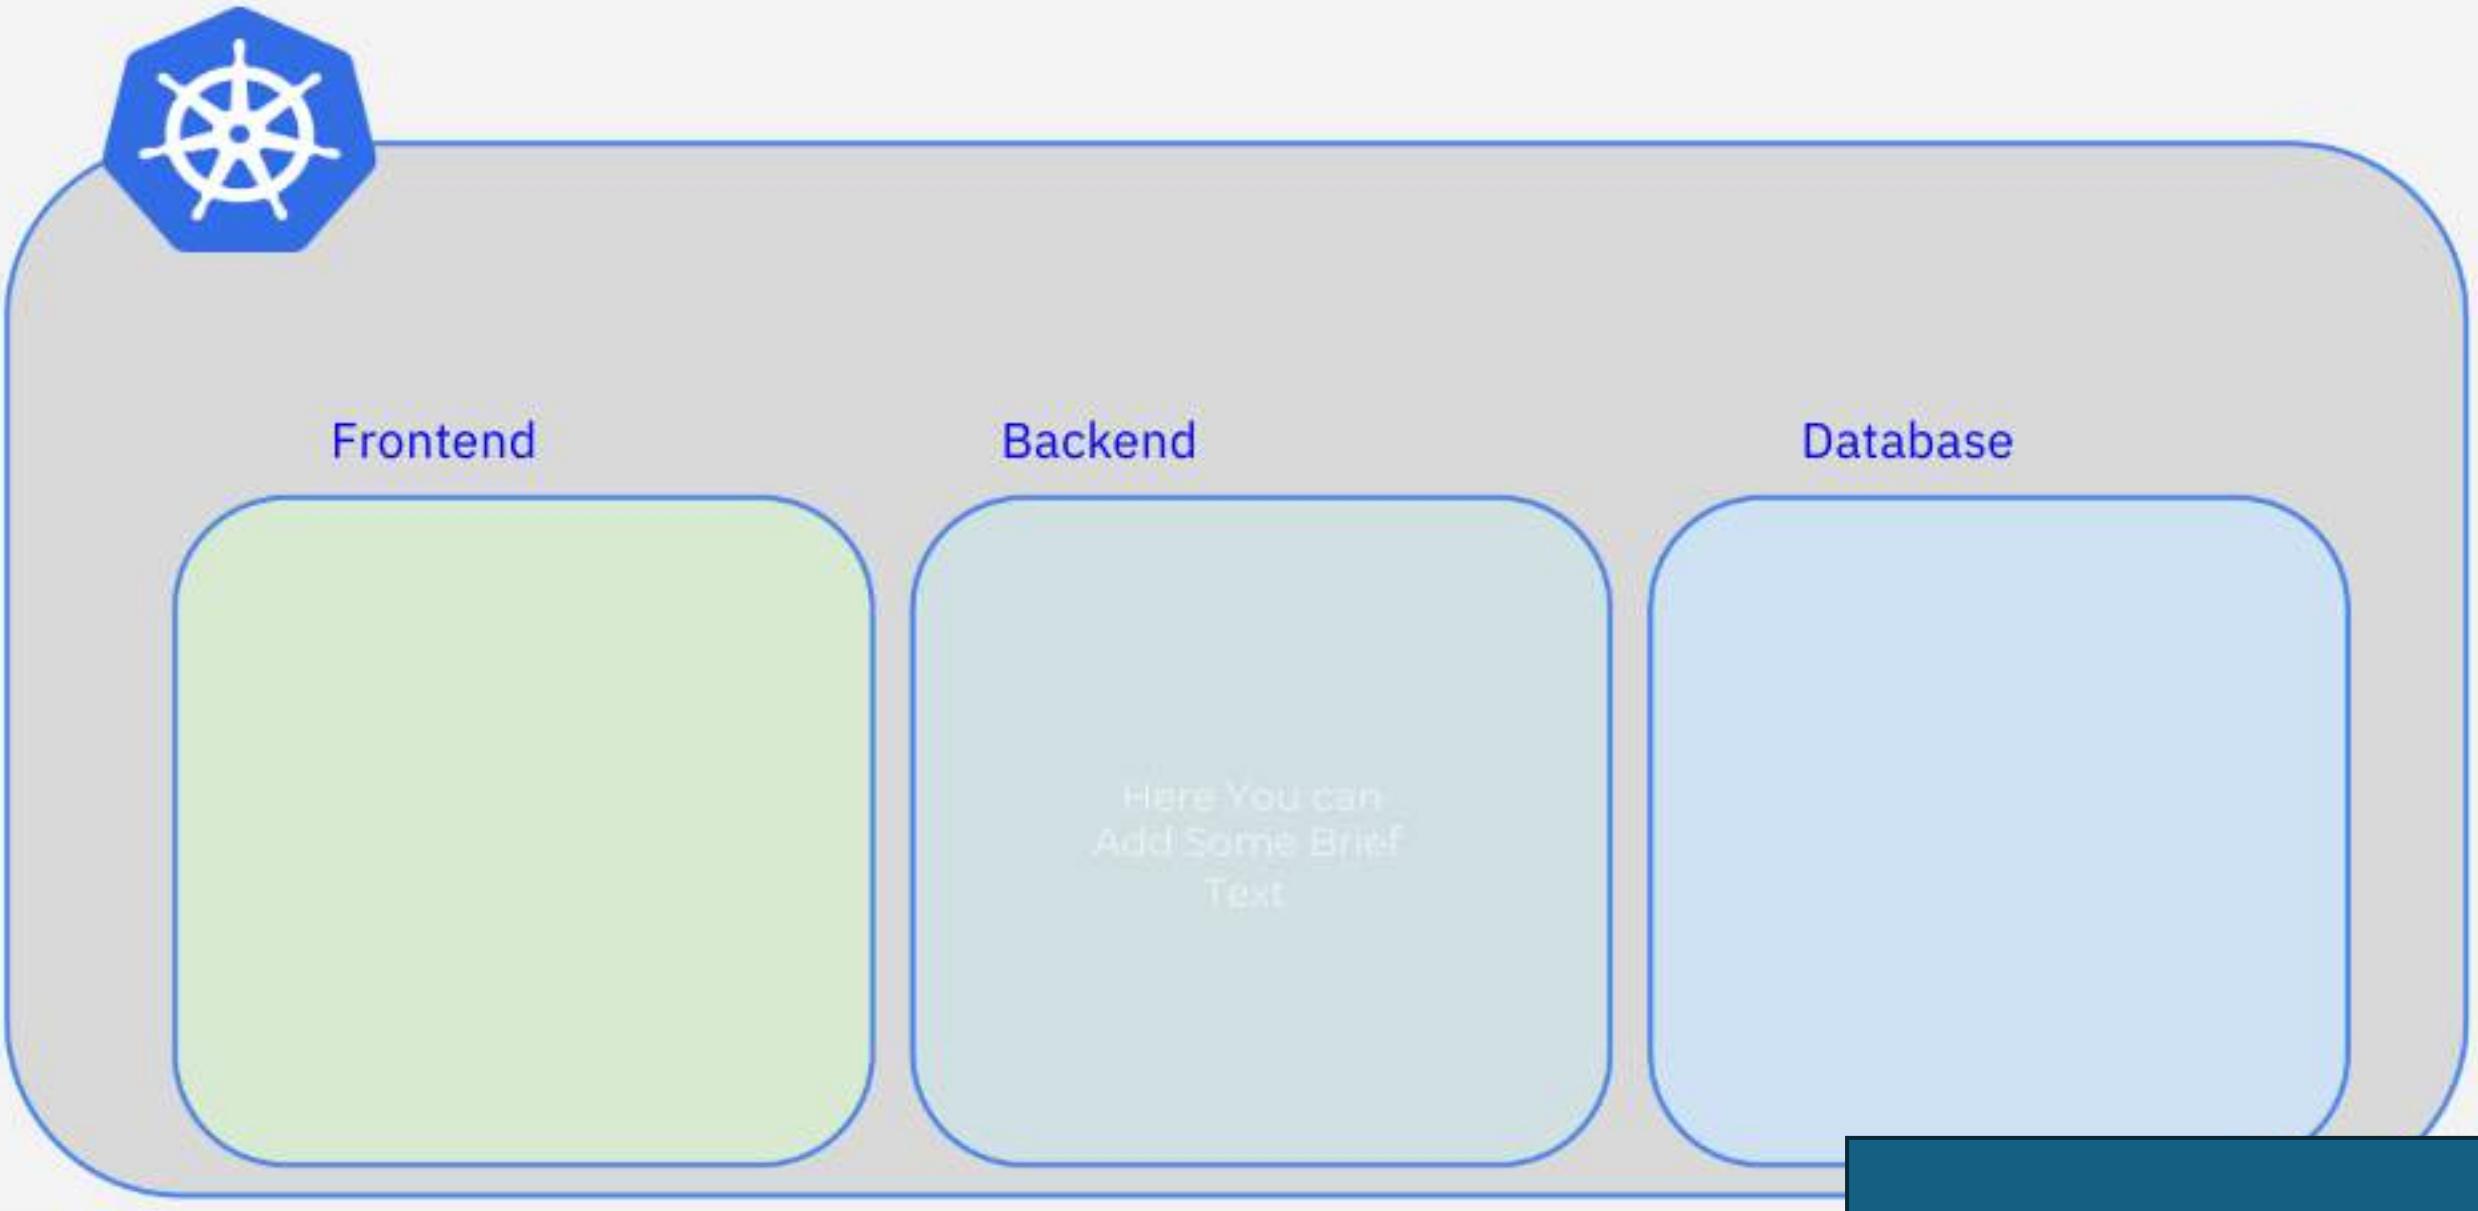
Here You can
Add Some Brief
Text

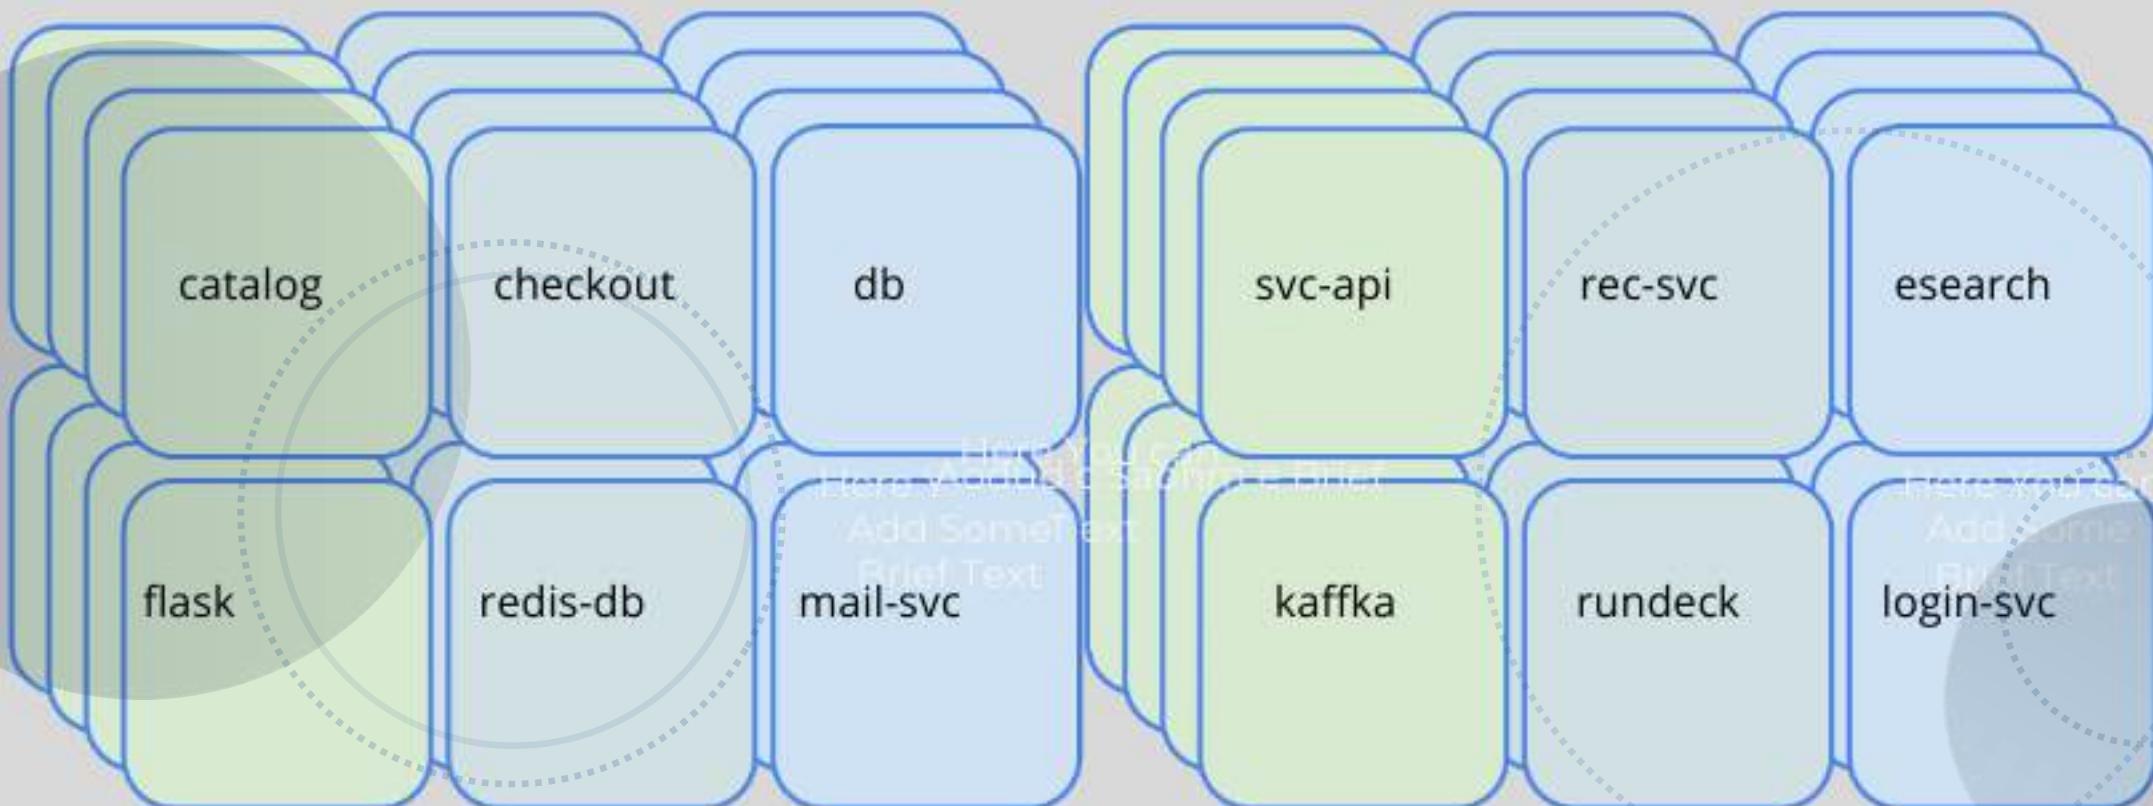
Namespaces





Namespaces



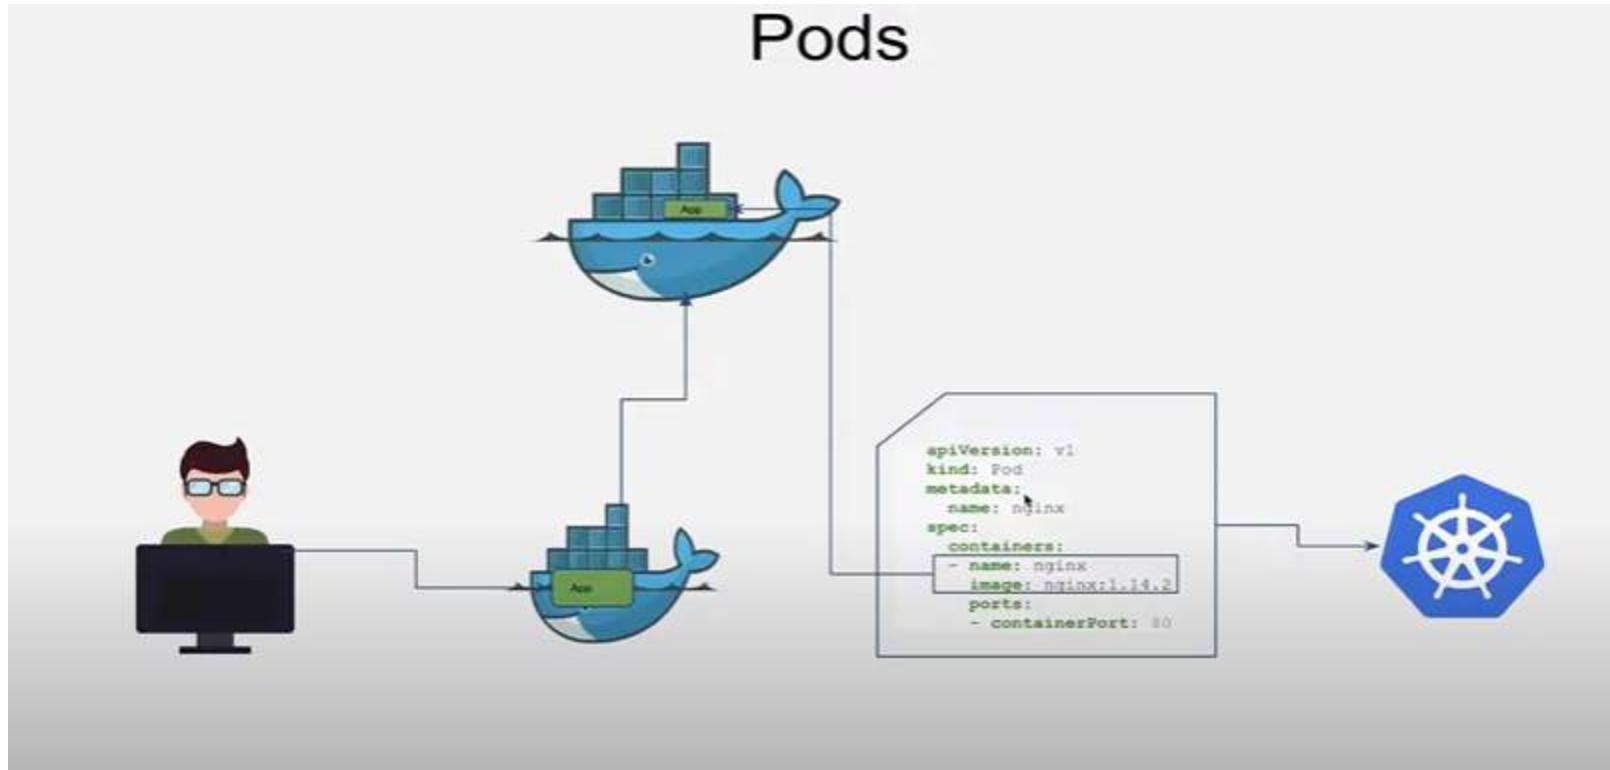


Pods

- Pods are the smallest deployable units in Kubernetes.
- A Pod can be one or more containers.



Pods



Foos

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
```

Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

(⎈ |kind-c11:default)→ kubectl create -f pod.yaml

(⎈ |kind-c11:default)→ kubectl get pod



Pods

```
# This is the pod
# Definition yaml file
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

Comment

String

Dictionary

List

Pods

```
example-pod.yaml
# This is the pod
# Definition yaml file
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

The diagram illustrates the structure of a YAML pod definition. It shows the file name and its contents, with annotations pointing to specific elements:

- File name:** example-pod.yaml
- Comment:** # This is the pod
Definition yaml file
- String:** apiVersion: v1
kind: Pod
- Dictionary:** metadata:
name: nginx
- List:** spec:
containers:
- name: nginx
image: nginx:1.14.2
ports:
- containerPort: 80

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
  ports:
    - containerPort: 80
```

Pods

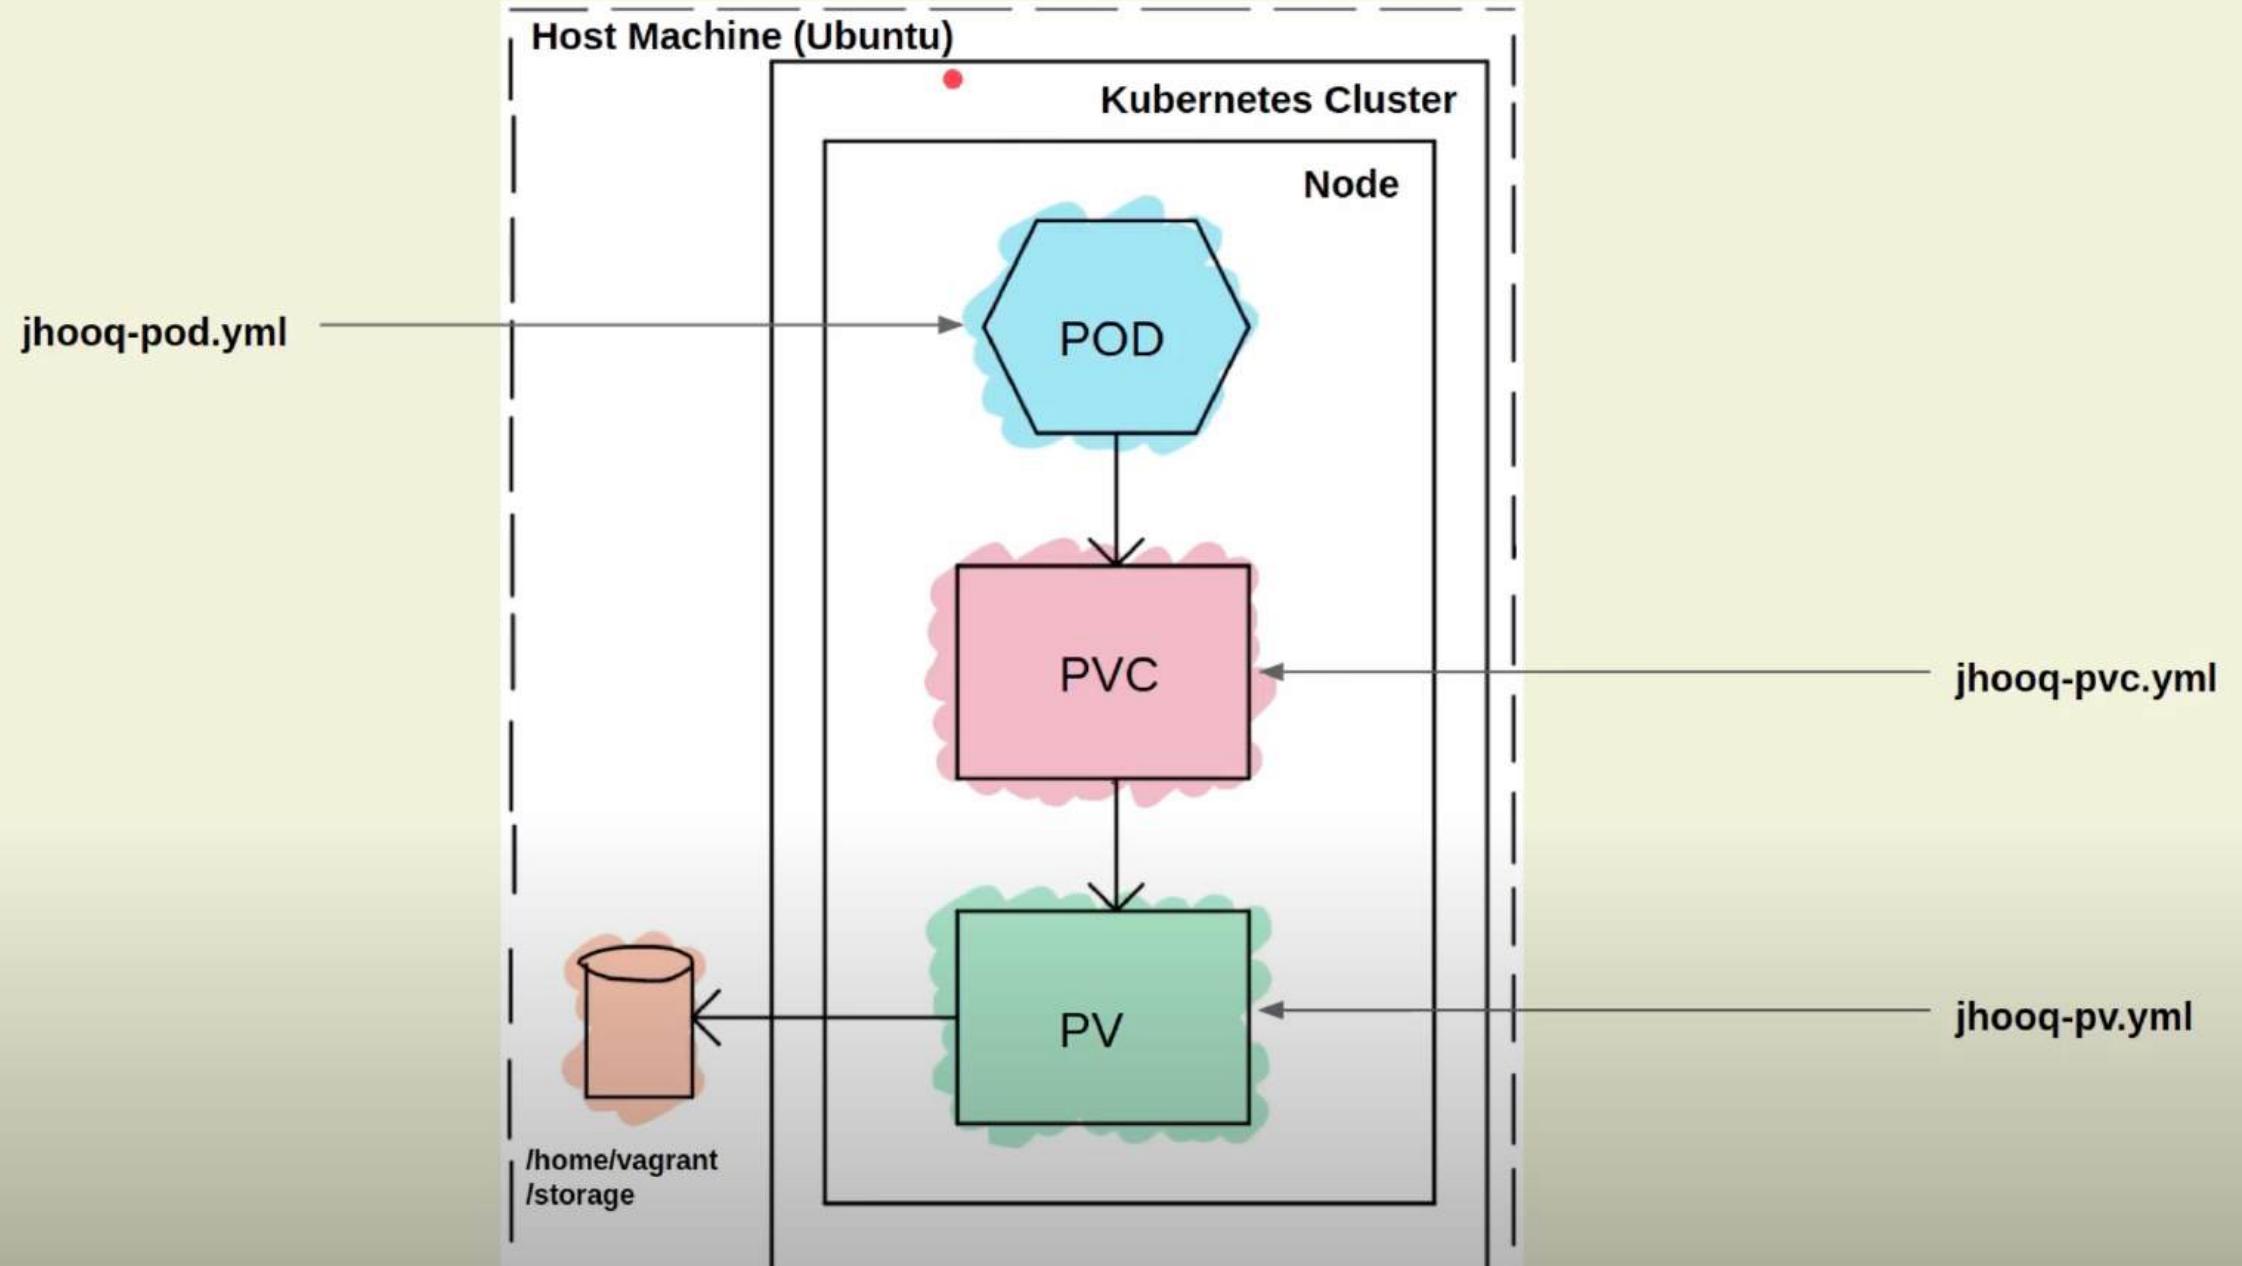
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
    ports:
      - containerPort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - image: nginx:1.14.2
      name: nginx
    ports:
      - containerPort: 80
```

Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
    - name: php-fpm
      image: php-fpm:1.12
  resources:
    limits:
      memory: "200Mi"
    requests:
      memory: "100Mi"
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: php-fpm
      image: php-fpm:1.12
      resources:
        limits:
          memory: "200Mi"
        requests:
          memory: "100Mi"
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```



Persistent Volume (PV)

Persistent Volume (PV)

A **Persistent Volume (PV)** is a piece of storage in a cluster that an administrator has provisioned. It is a resource in the cluster, just as a node is a cluster resource. A persistent volume is a volume plug-in that has a lifecycle independent of any individual pod that uses the persistent volume.

Access Modes

There are currently **four access modes** for PVs in Kubernetes:

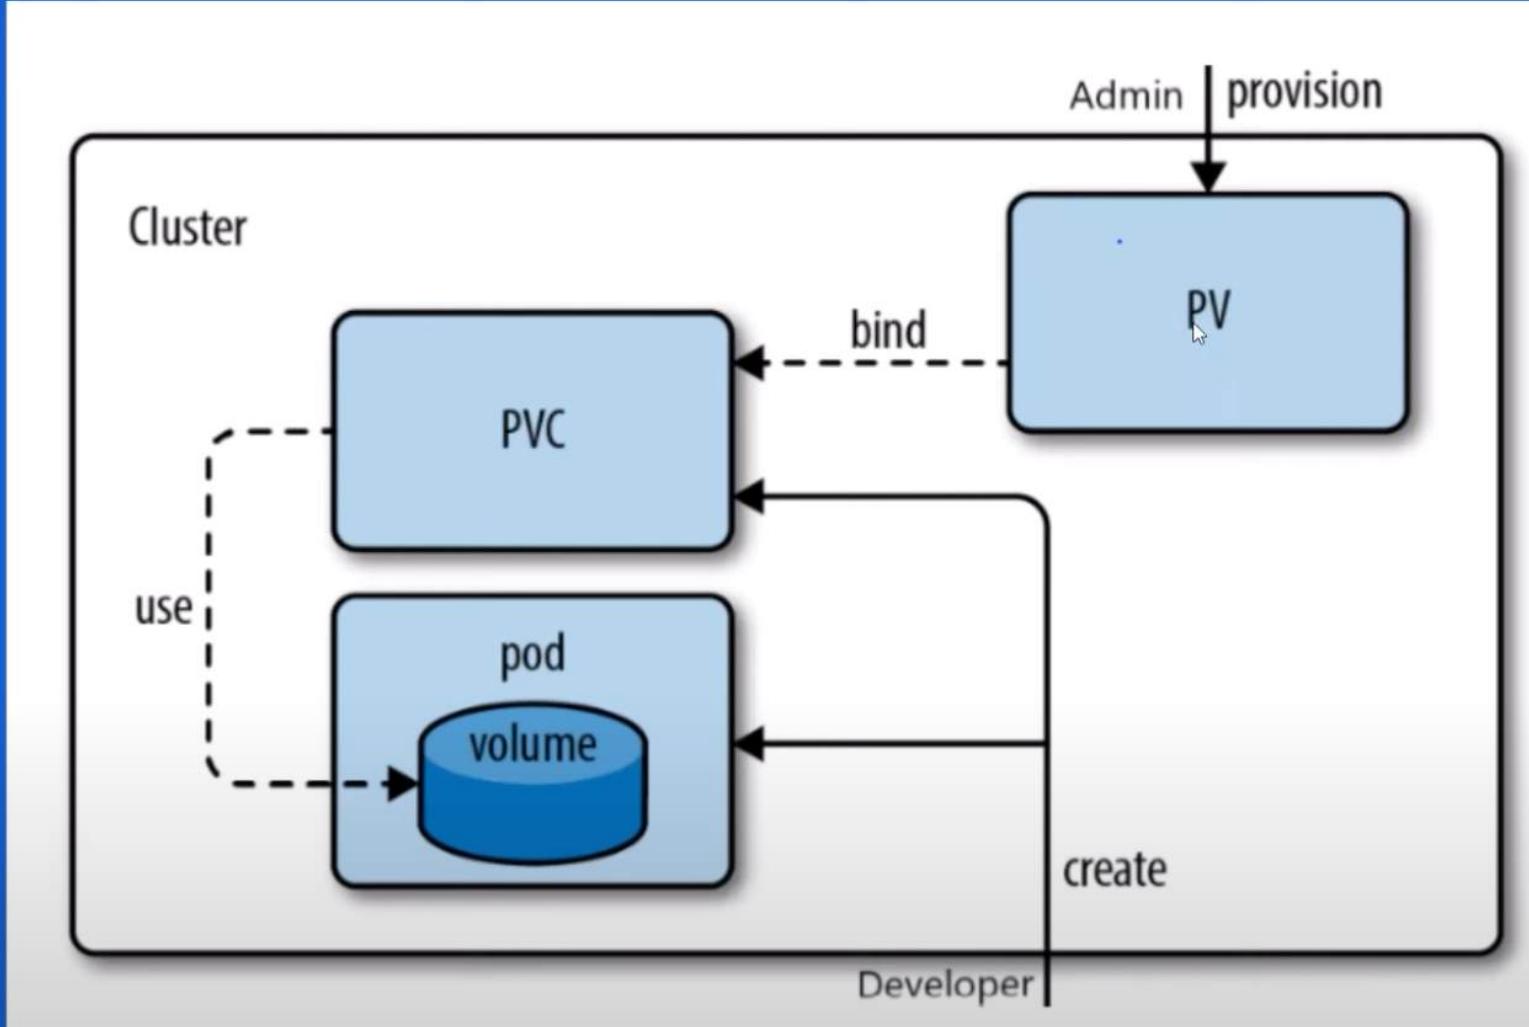
- **ReadWriteOnce**: This allows only a single node to access the volume in read-write mode. Furthermore, all pods in that single node can read and write to such volumes.
- **ReadWriteMany**: Multiple nodes can read and write to the volume.
- **ReadOnlyMany**: This means that the volume will be in a read-only mode and accessible by multiple nodes.
- **ReadWriteOncePod**: Only a single pod can gain access to the volume.

Persistent Volume Claim (PVC)

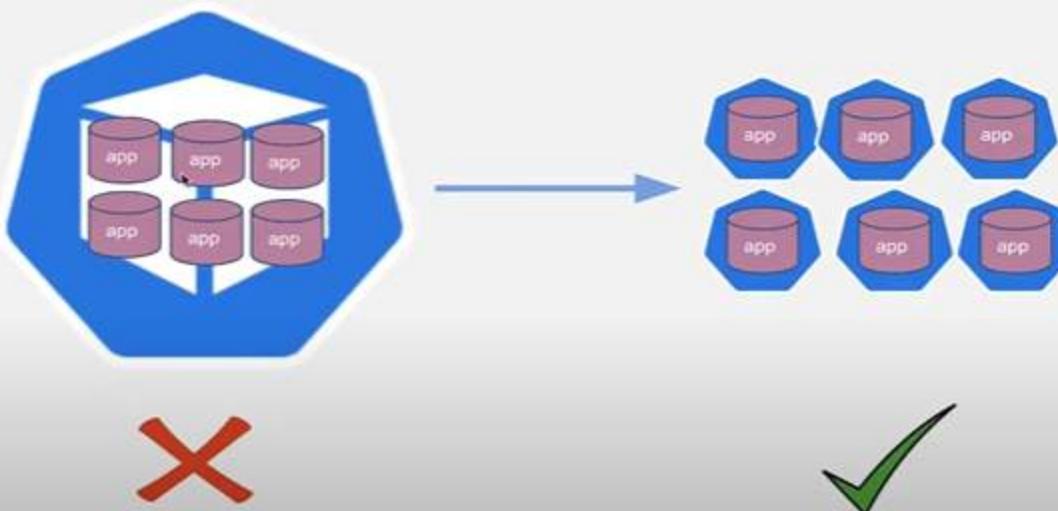
Persistent Volume Claim

A **Persistent Volume Claim (PVC)** is a **request for storage by a user**. It is similar to a Pod. Pods consume node resources and PVCs consume PV resources.

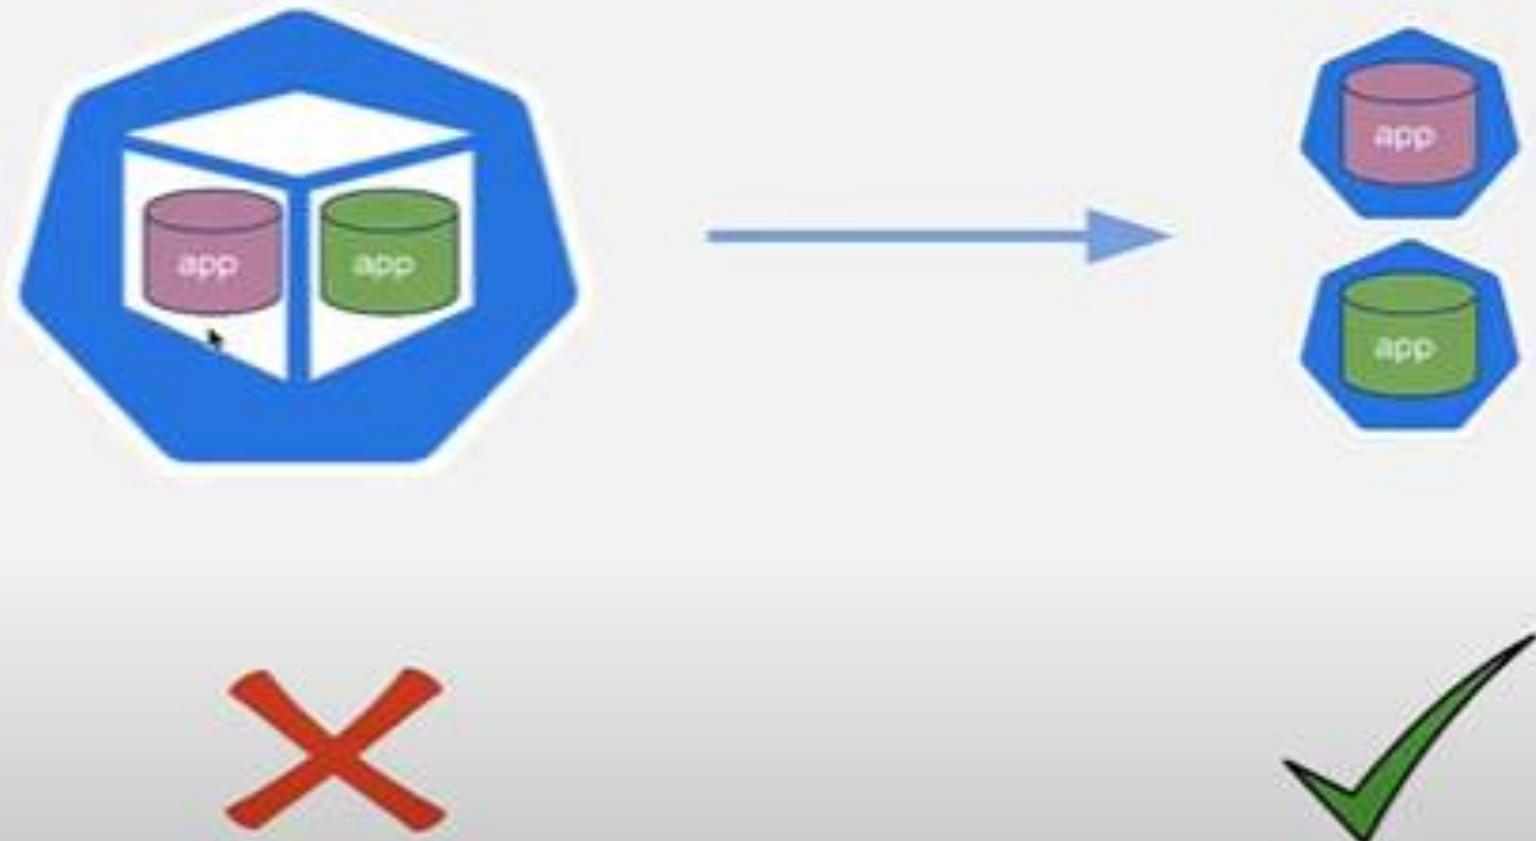
Persistent Volume Claim (PVC)



Multi Container Pod



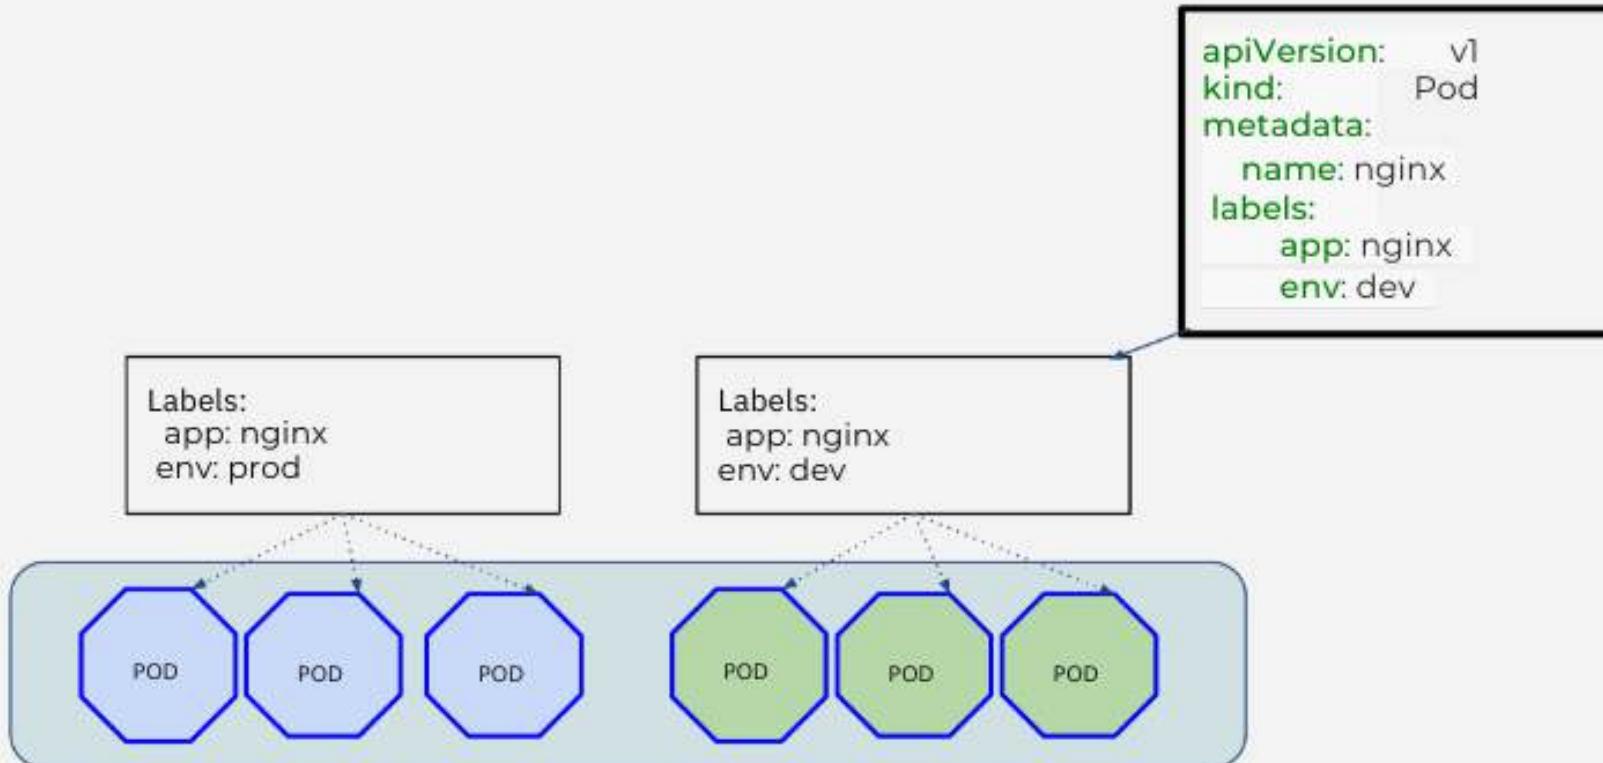
Multi Container Pod



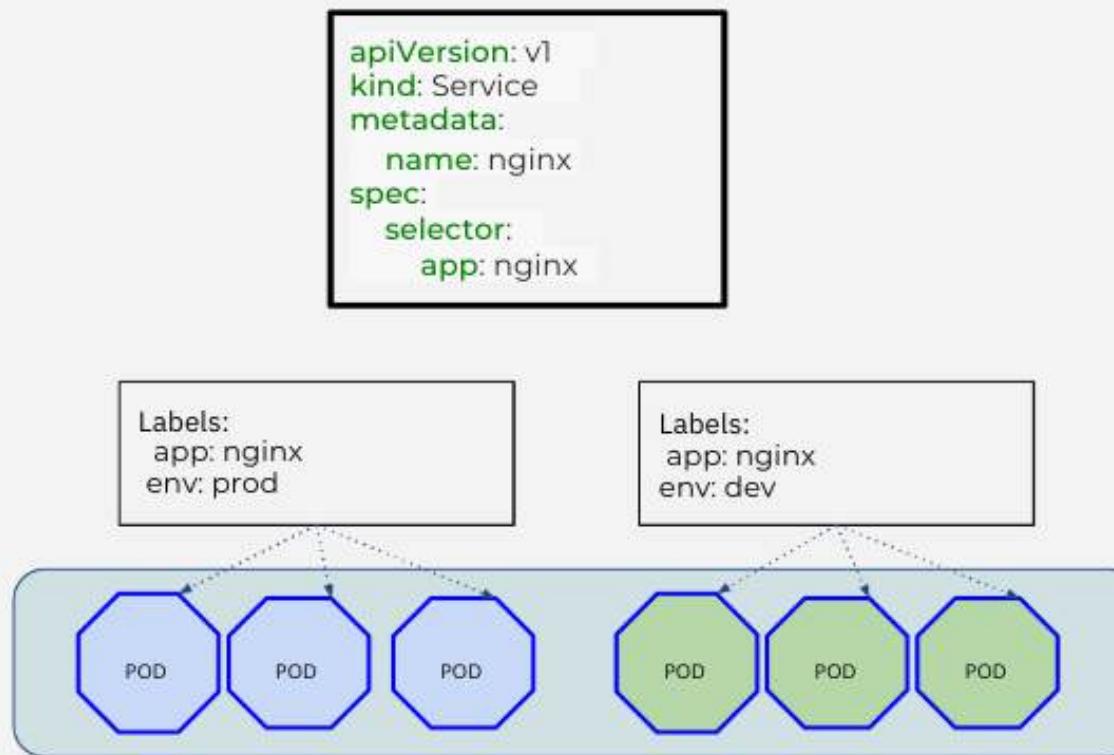
Labels and Selectors

- ❑ Labels are key/value pairs used to identify objects
- ❑ Labels can be attached to almost any Kubernetes Object
- ❑ Labels are mutable, they can be modified at anytime.
- ❑ Selectors are a way to select objects based on their labels
- ❑ Types of Selectors
 - ❑ Equality-based selectors
 - ❑ environment = production
 - ❑ tier != frontend
 - ❑ Set-based selectors
 - ❑ environment in (production, qa)
 - ❑ tier notin (frontend, backend)

Labels and Selectors



Labels and Selectors



Deployments

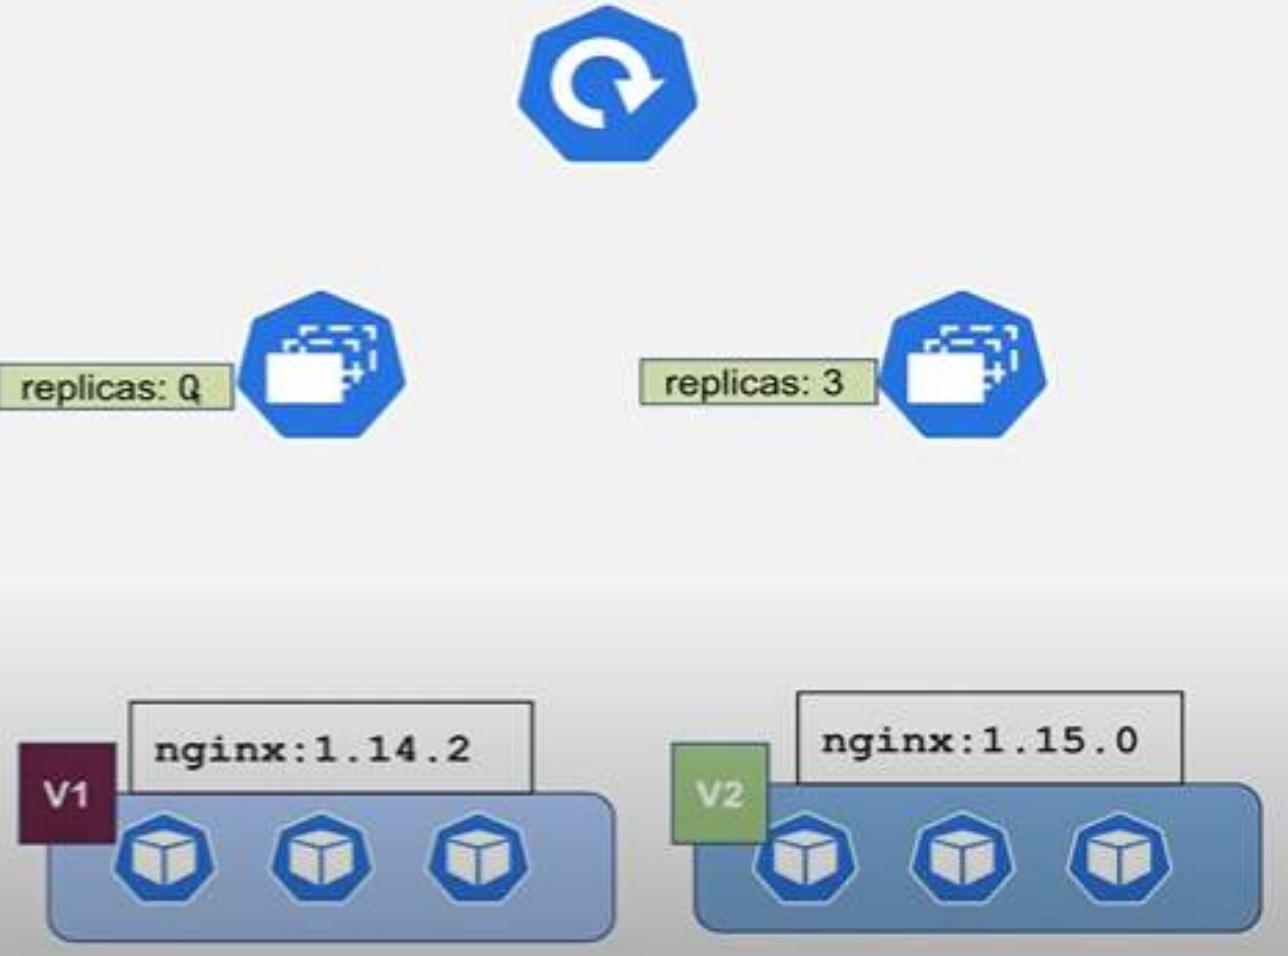
What is Kubernetes Deployment

Managing The Lifecycle of The Application:

- Create the Pods
- Rollout new version of the application
- Rollback to previous version
- Scale the application
- Self healing

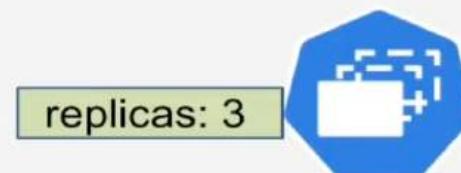
Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
```



Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
      ports:
        - containerPort: 80
```



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.15.0
      ports:
        - containerPort: 80
```



Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

Create Deployment

Declarative

Imperative

(✿ |kind-c11:default)-> kubectl
create -f deploy.yaml



(✿ |kind-c11:default)-> kubectl create
deployment nginx-deployment
--image=nginx:1.14.2
--replicas=3
--port=80

Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
```

Create Deployment

Declarative

Imperative



(✿ |kind-c11:default)-> kubectl
create -f deploy.yaml

(✿ |kind-c11:default)-> kubectl create
deployment nginx-deployment
--image=nginx:1.14.2
--replicas=3
--port=80

Deployments

Demo



- Deployment YAML File
- Create Deployment
 - Declarative
 - Imperative
- Update Deployment
- Rollback Deployment
- Scale Deployment
- Revision History Limit

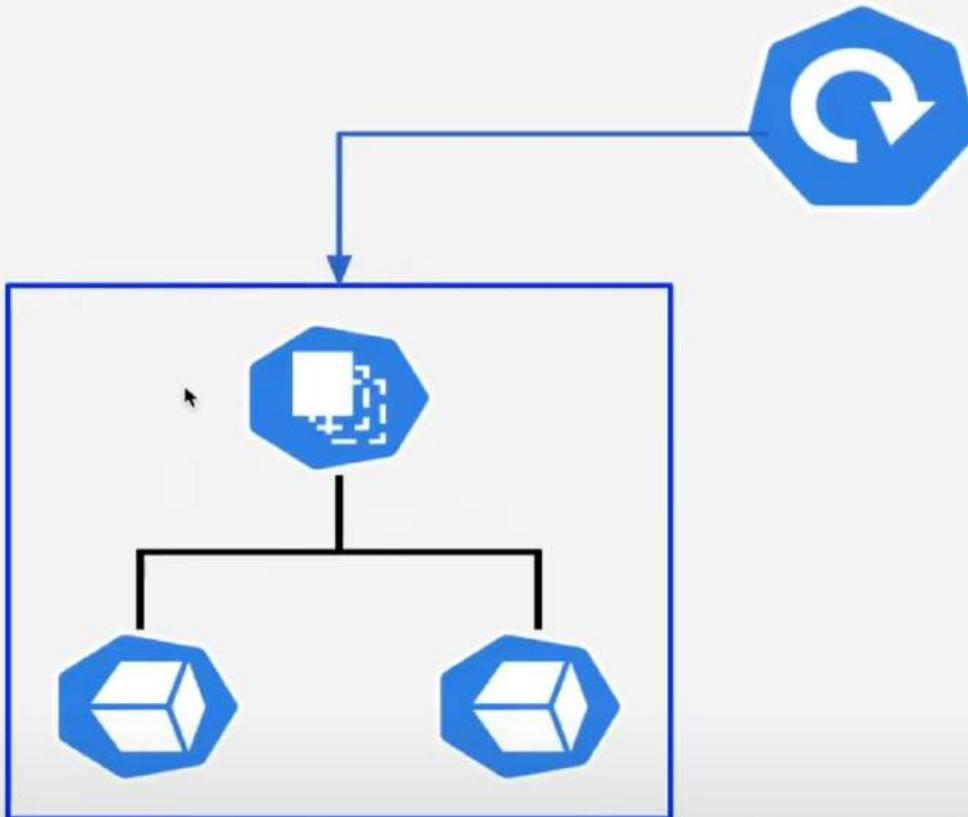
Reference:

<https://kubernetes.io/docs/concepts/workloads/controllers/deployment/>

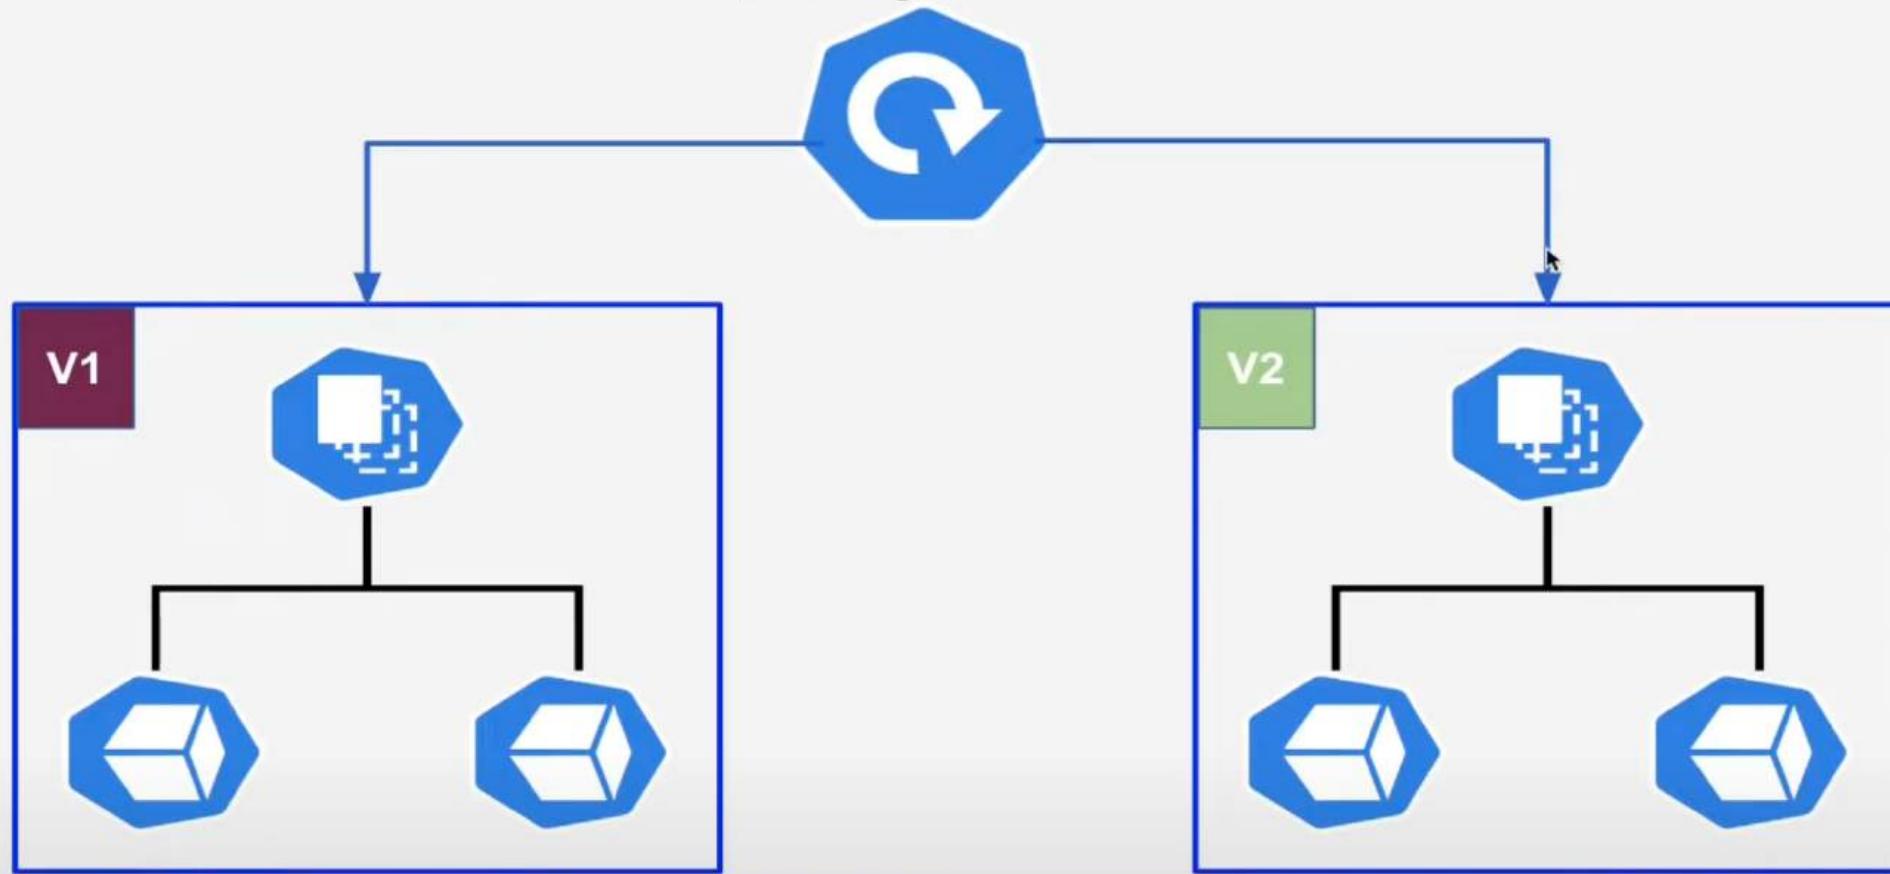
Deployments



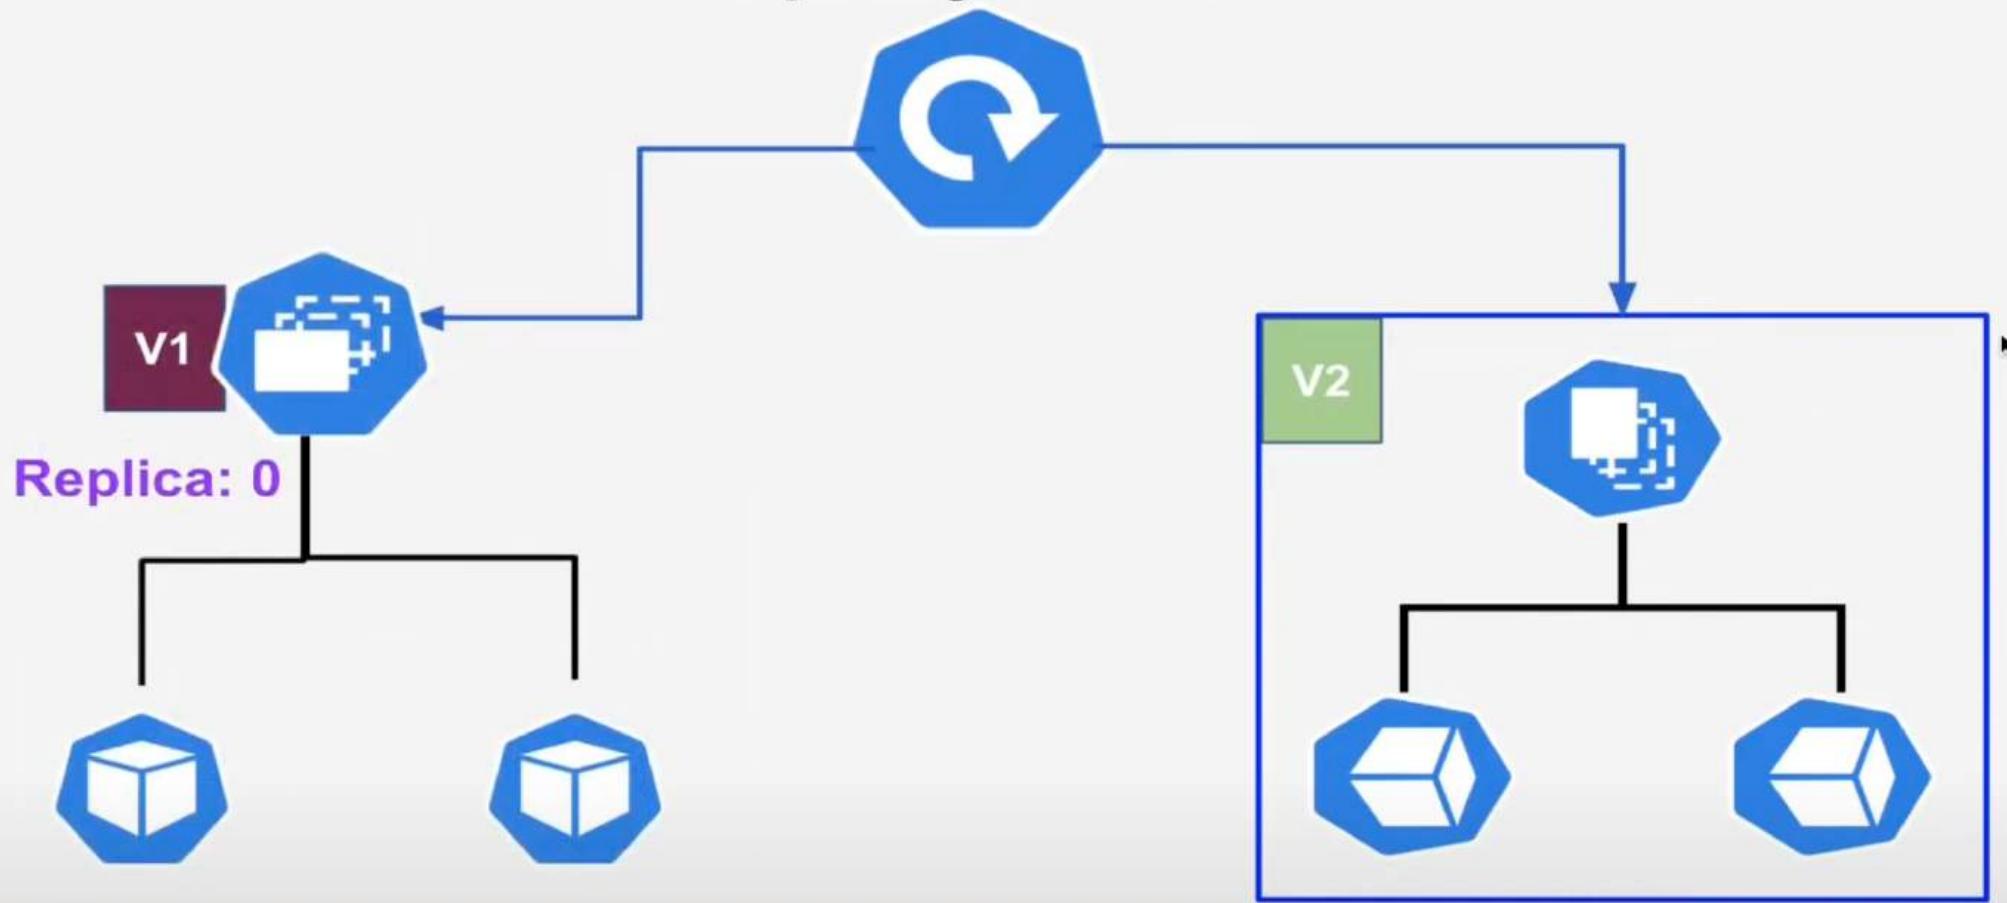
Deployments



Deployments



Deployments



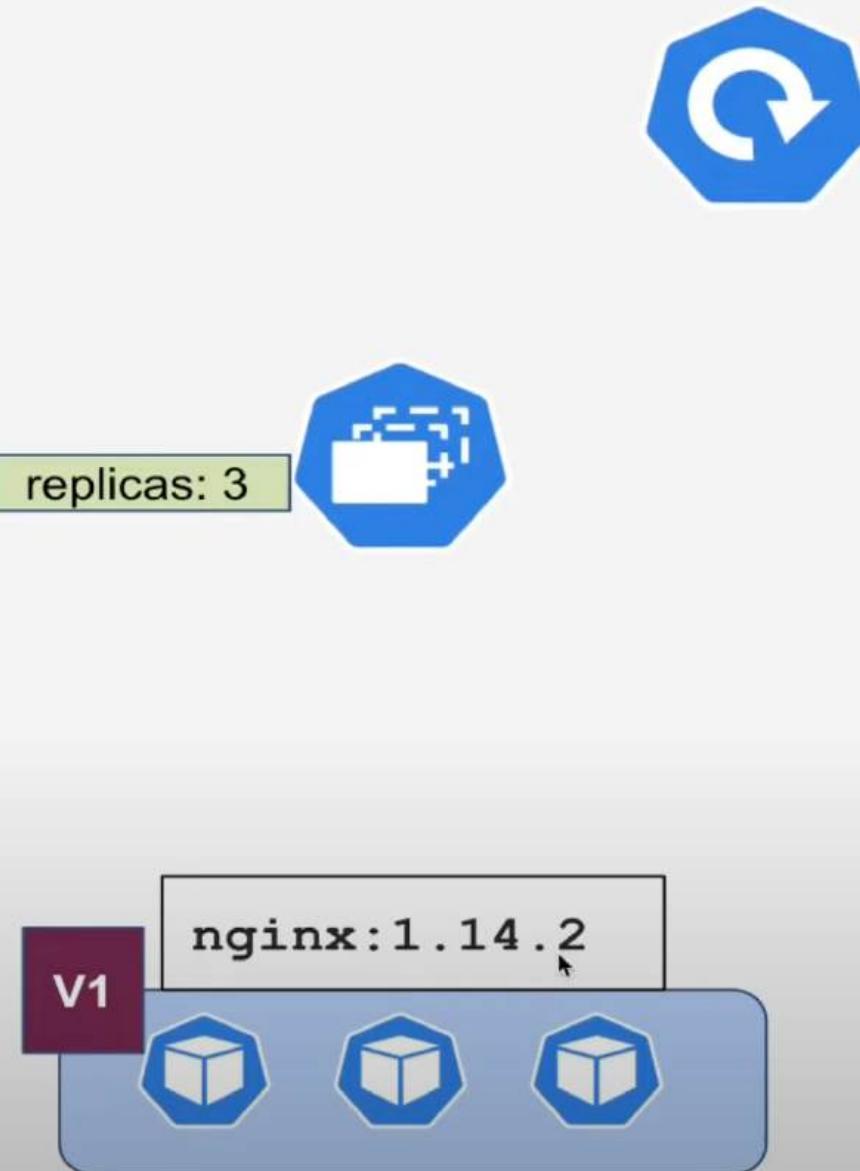
Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
  spec:
    containers:
    - name: nginx
      image: nginx:1.14.2
    ports:
    - containerPort: 80
```



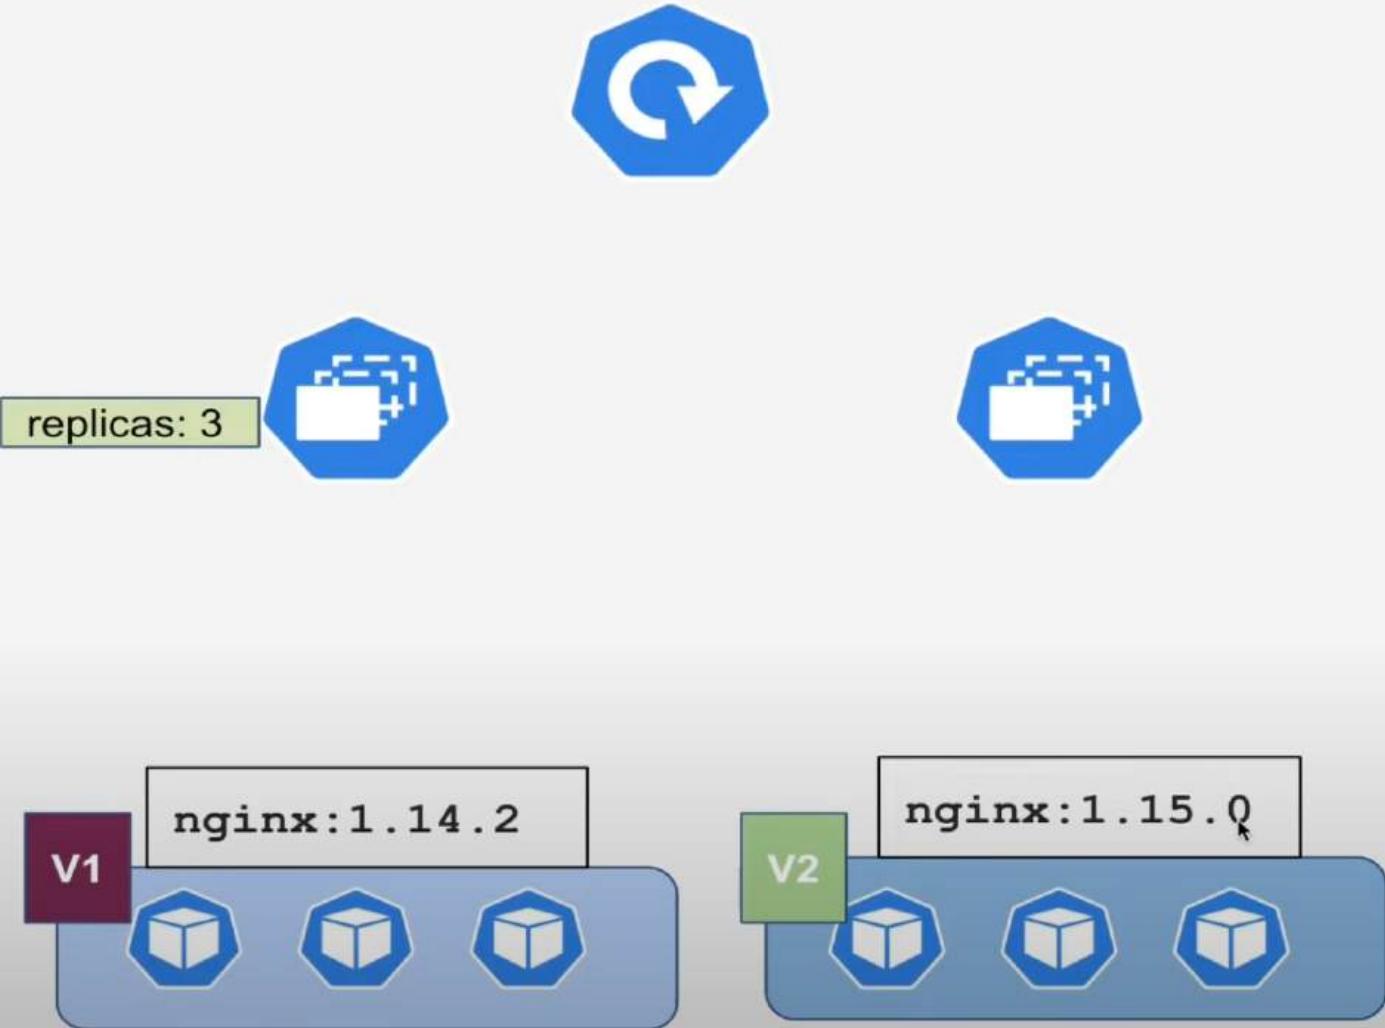
Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
      ports:
        - containerPort: 80
```



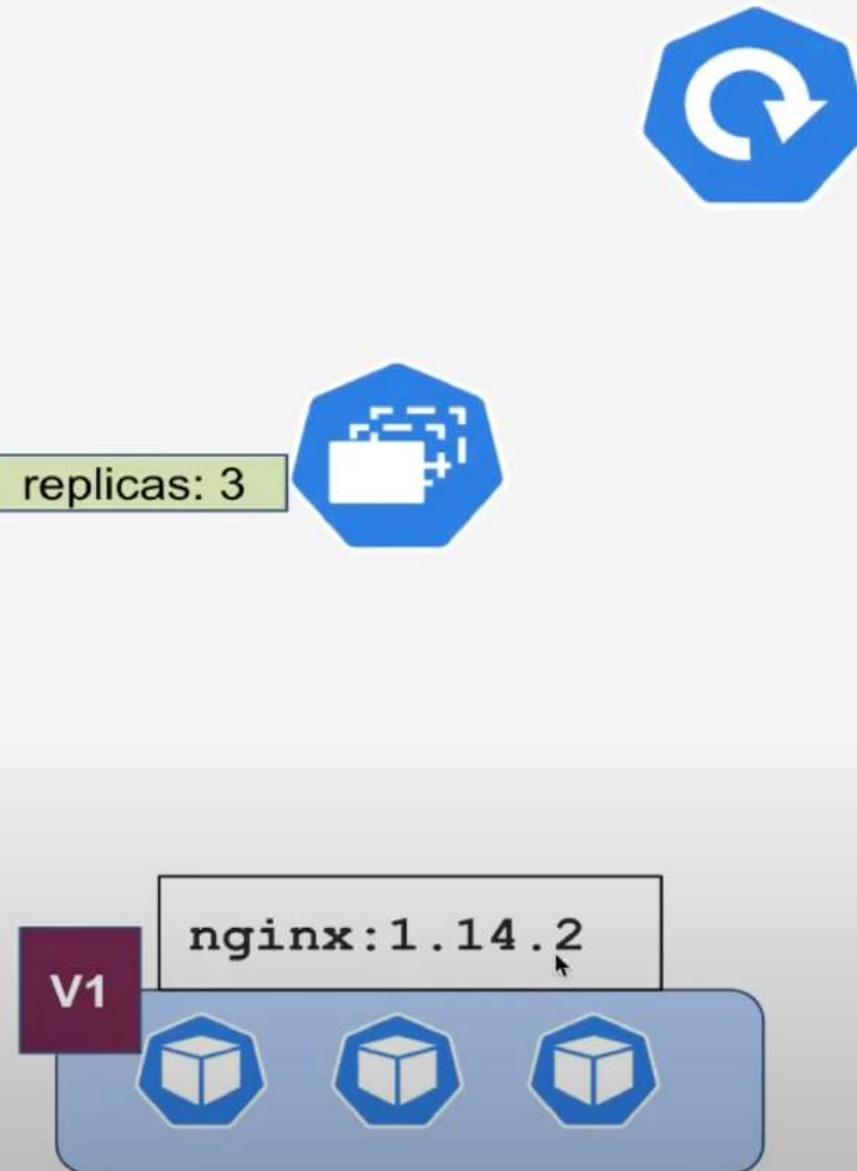
Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```



Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
      ports:
        - containerPort: 80
```



Deployment Strategies

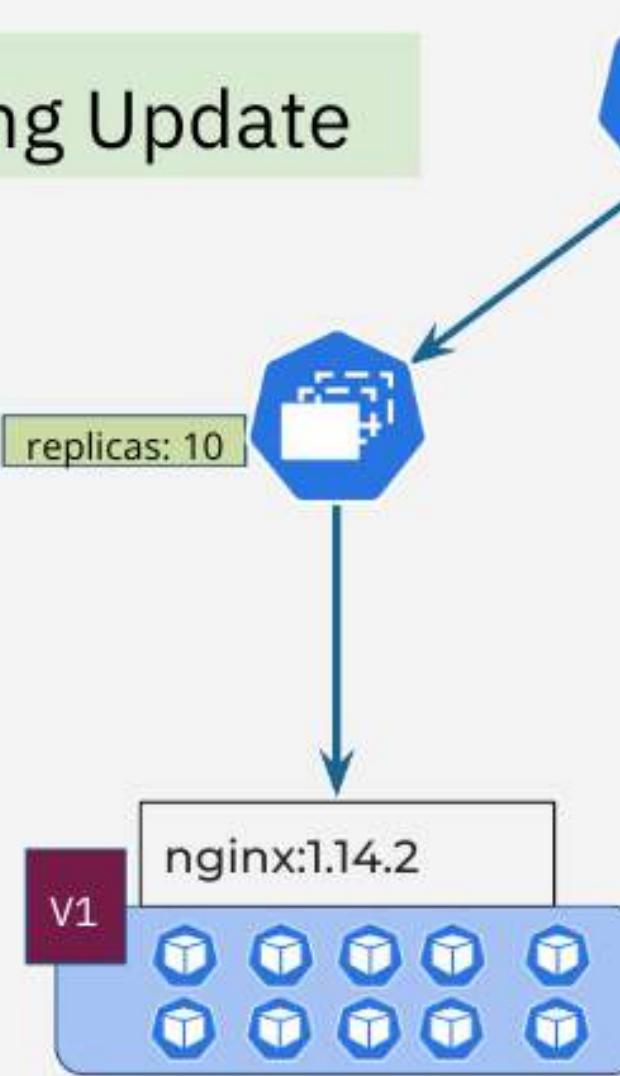
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  strategy:
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
```

Rolling Update

Deployment Strategies

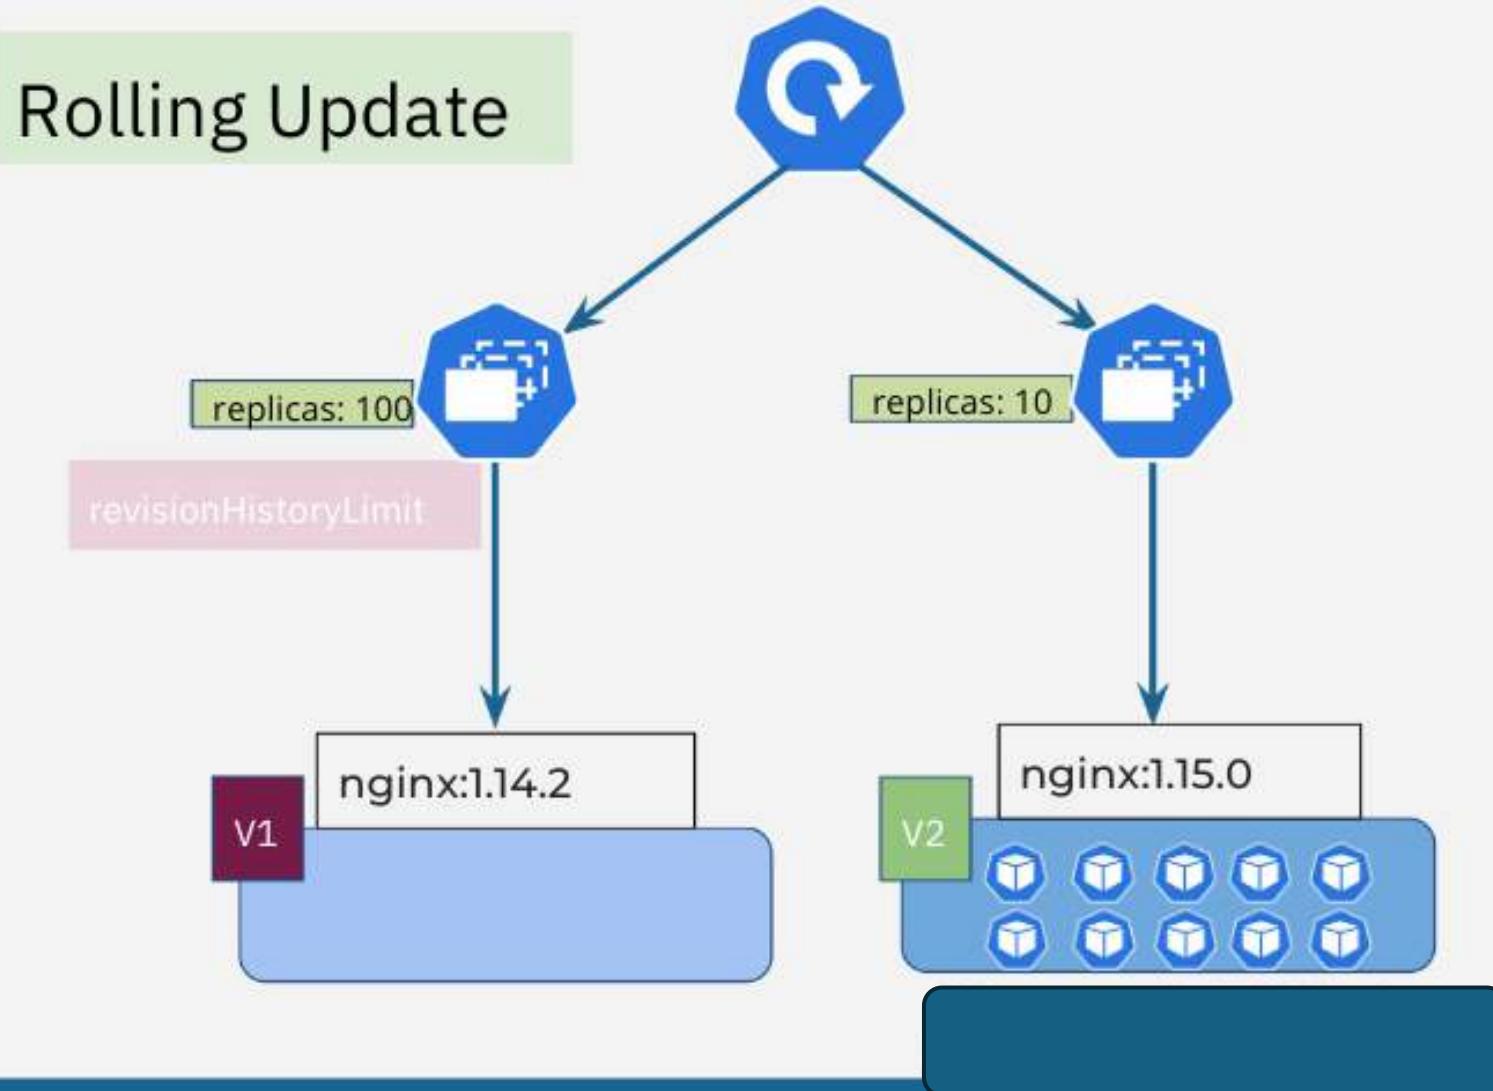
Rolling Update

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  strategy:
    rollingUpdate:
      maxSurge: 30%
      maxUnavailable: 30%
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```



Deployment Strategies

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 10
  selector:
    matchLabels:
      app: nginx
  strategy:
    rollingUpdate:
      maxSurge: 30%
      maxUnavailable: 30%
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.15.0
          ports:
            - containerPort: 80
```



Services

01

Service Discovery

A way to connect different services or components with a single stable endpoint

02

Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.

Services

01

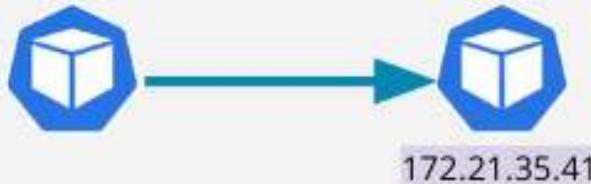
Service Discovery

A way to connect different services or components with a single stable endpoint

02

Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.



Services

01

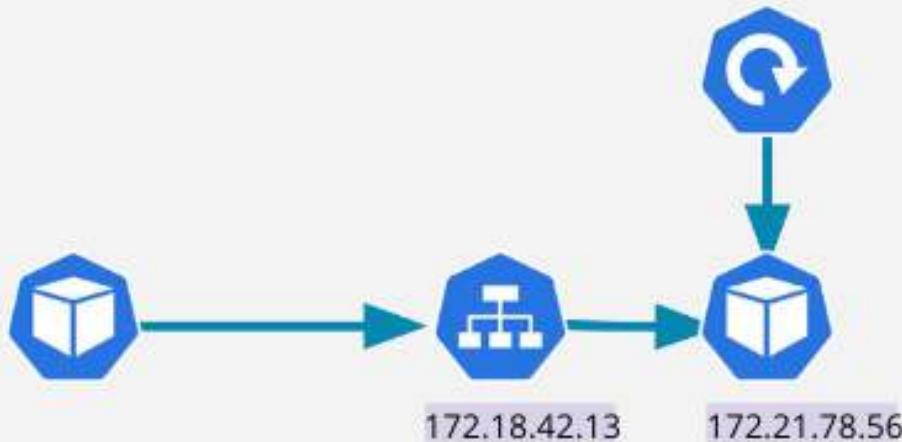
Service Discovery

A way to connect different services or components with a single stable endpoint

02

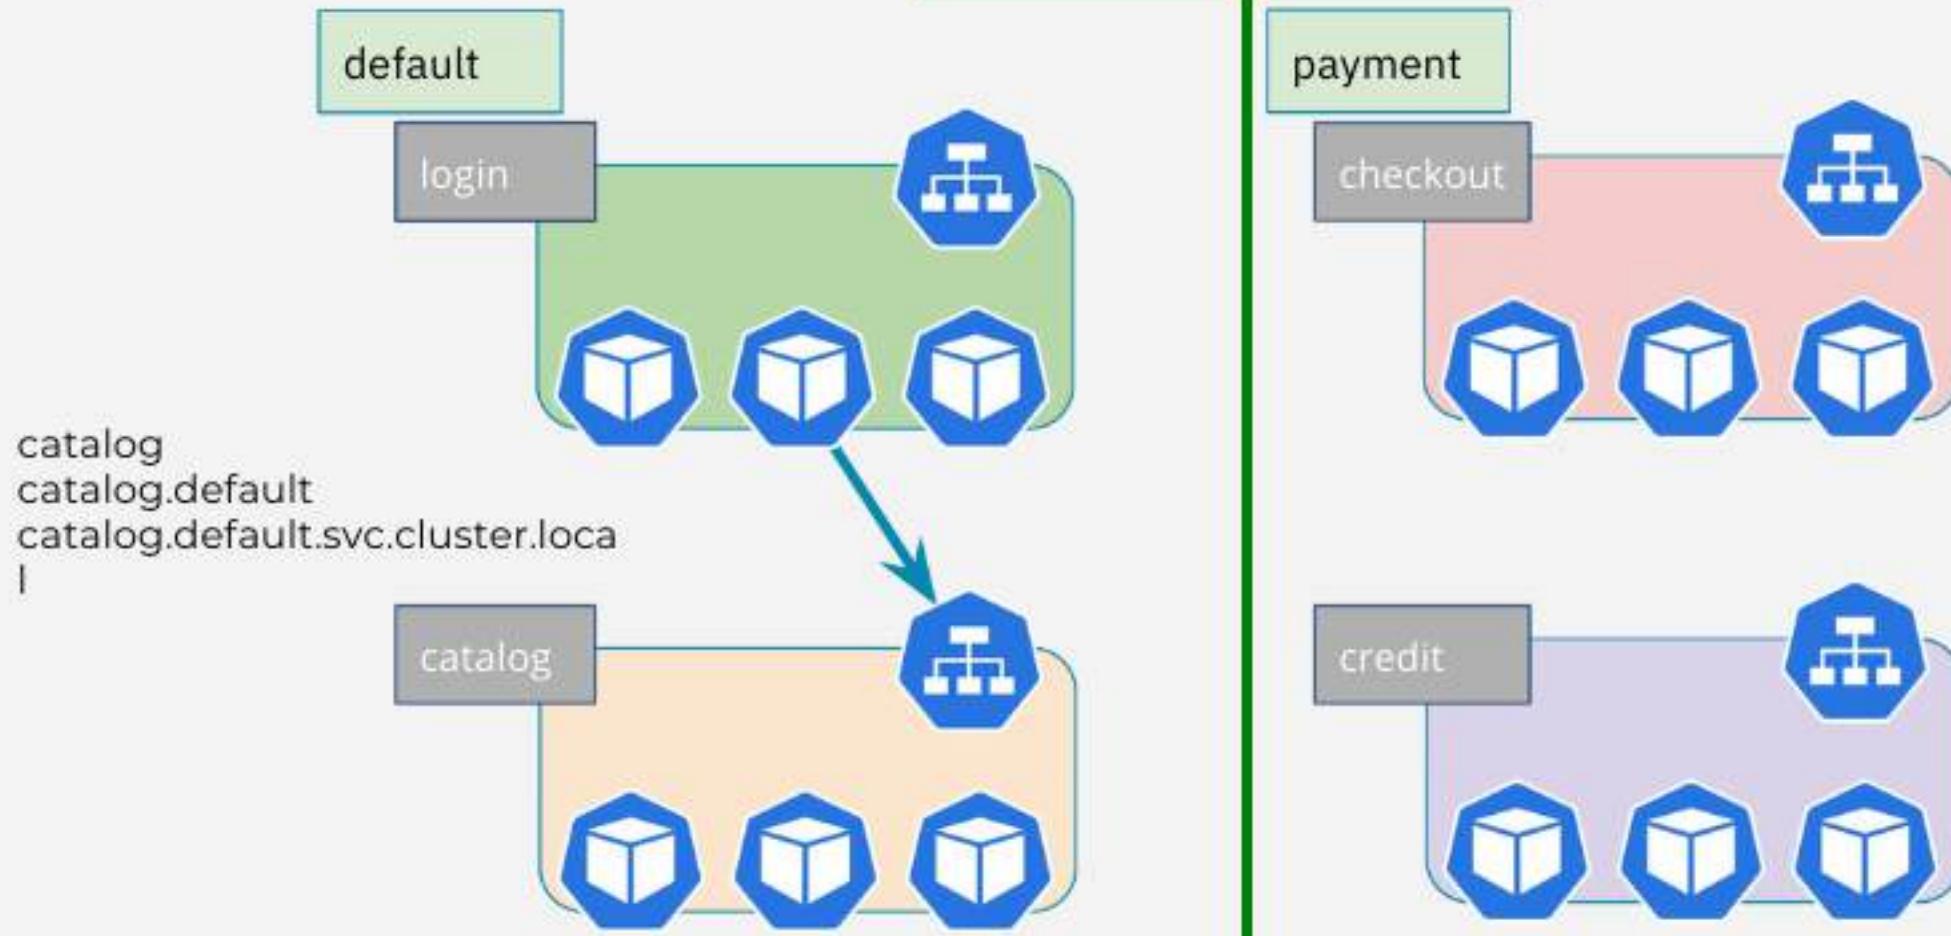
Load Balancing

Kubernetes automatically distributes traffic between different application backends to optimize performance and availability.



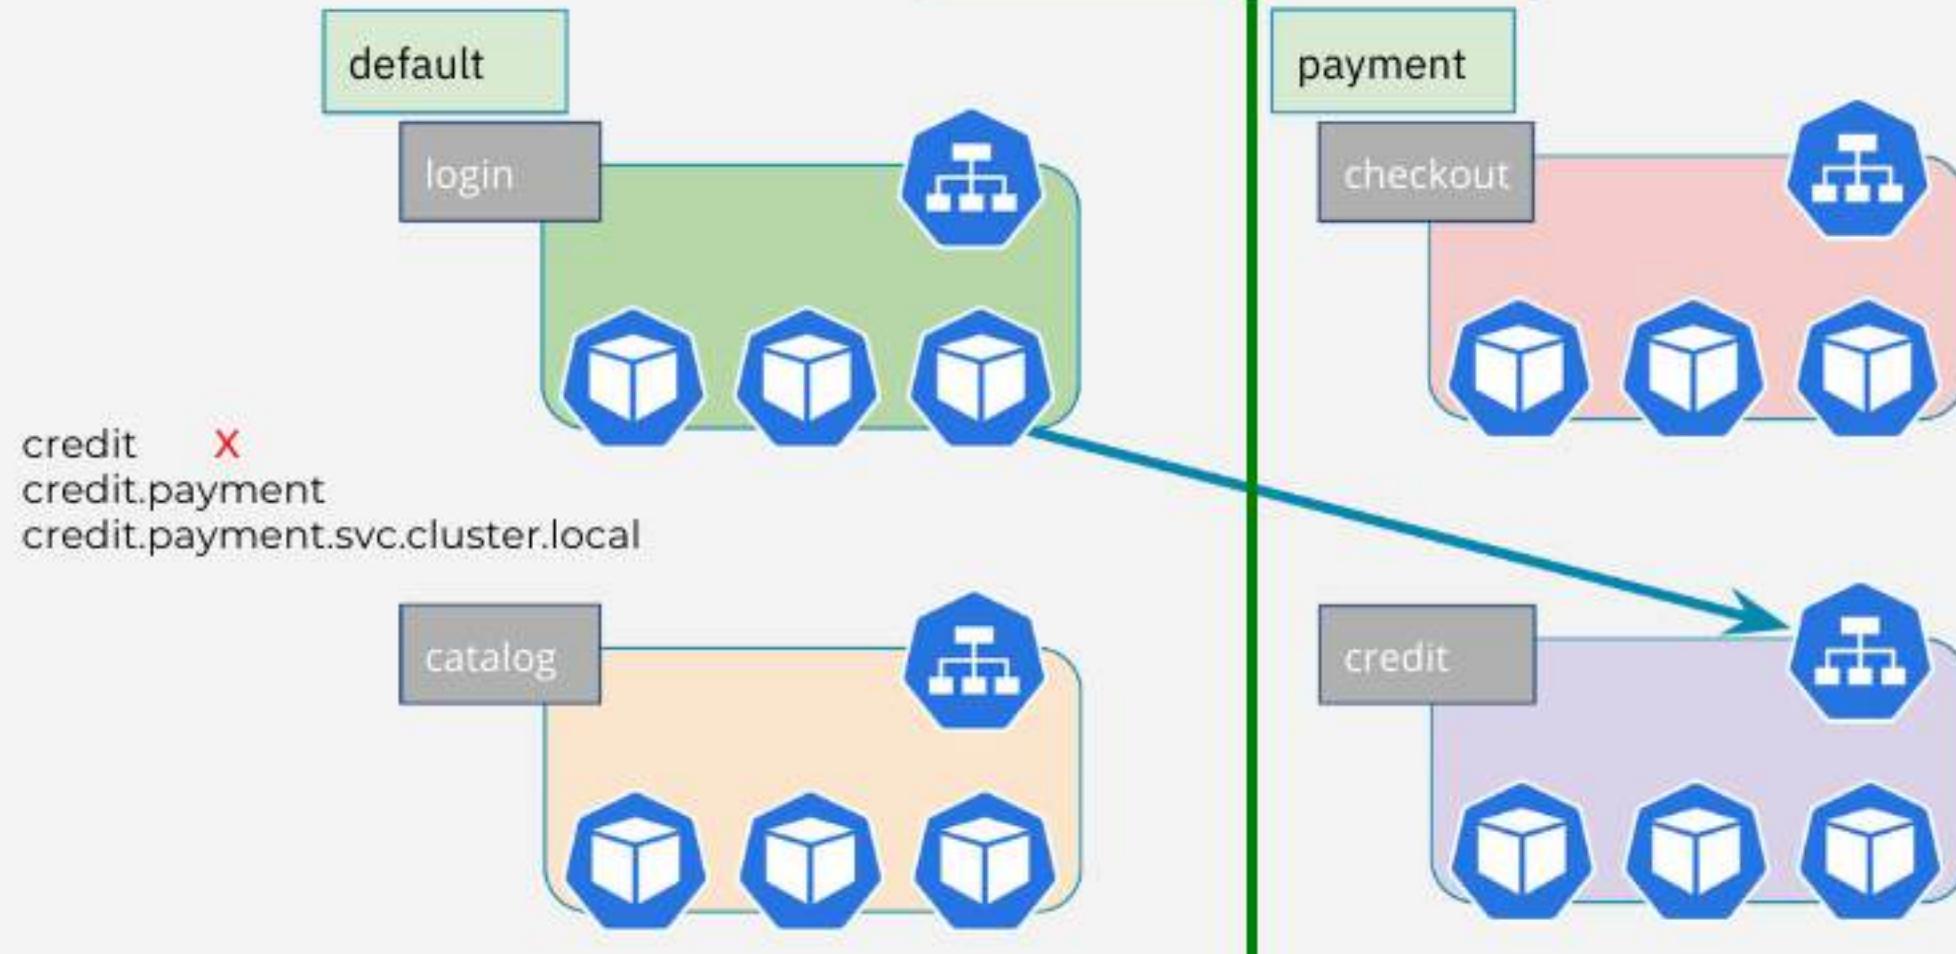
Services

Service Name



Services

Service Name



Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pod
  labels:
    app: nginx
spec:
  containers:
    - name: web
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
```

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pod
  labels:
    app: nginx
spec:
  containers:
    - name: web
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

Services

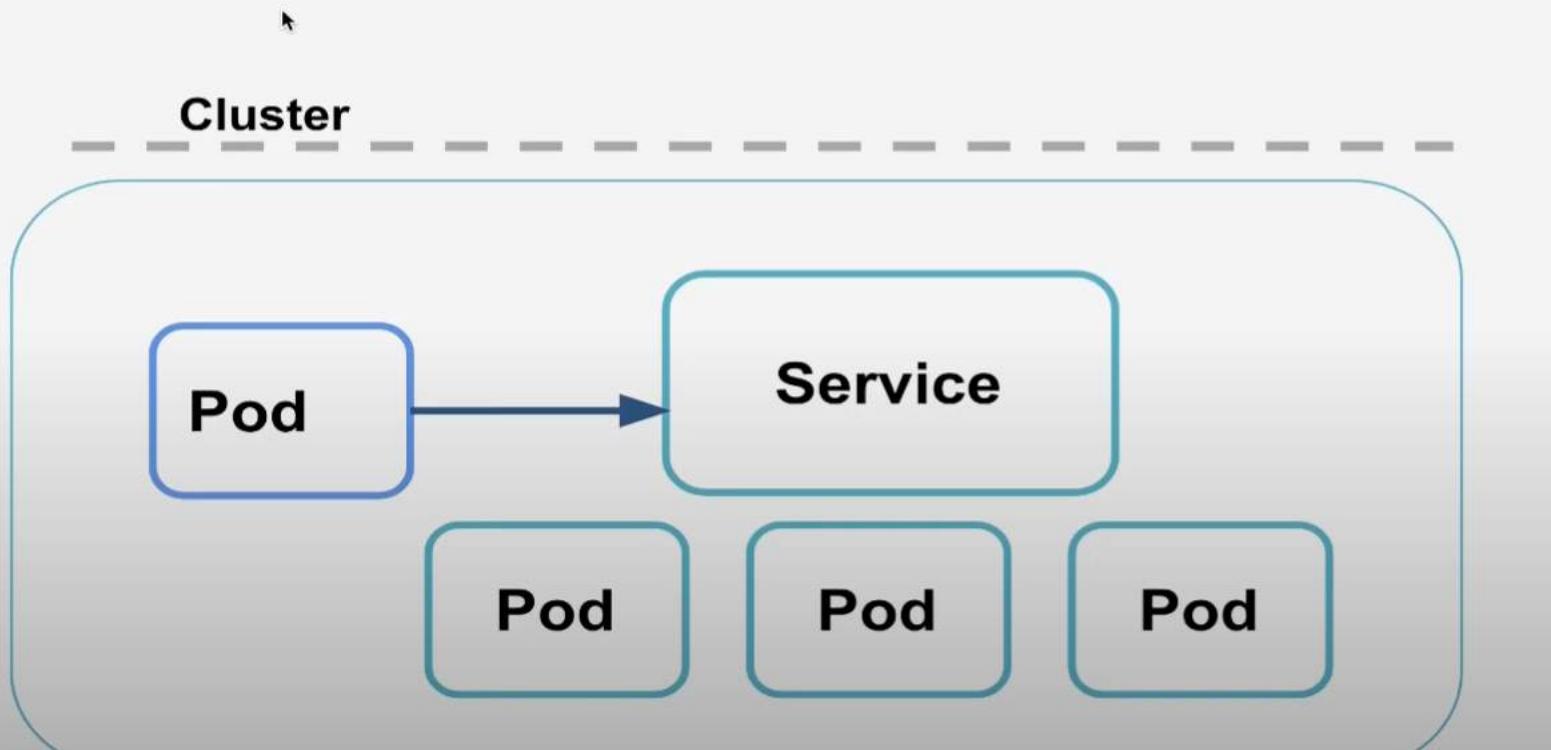
```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      name: http
      port: 80
      targetPort: 9376
    - name: https
      protocol: TCP
      port: 443
      targetPort: 9377
```

Service Types

- ClusterIP
- NodePort
- LoadBalancer

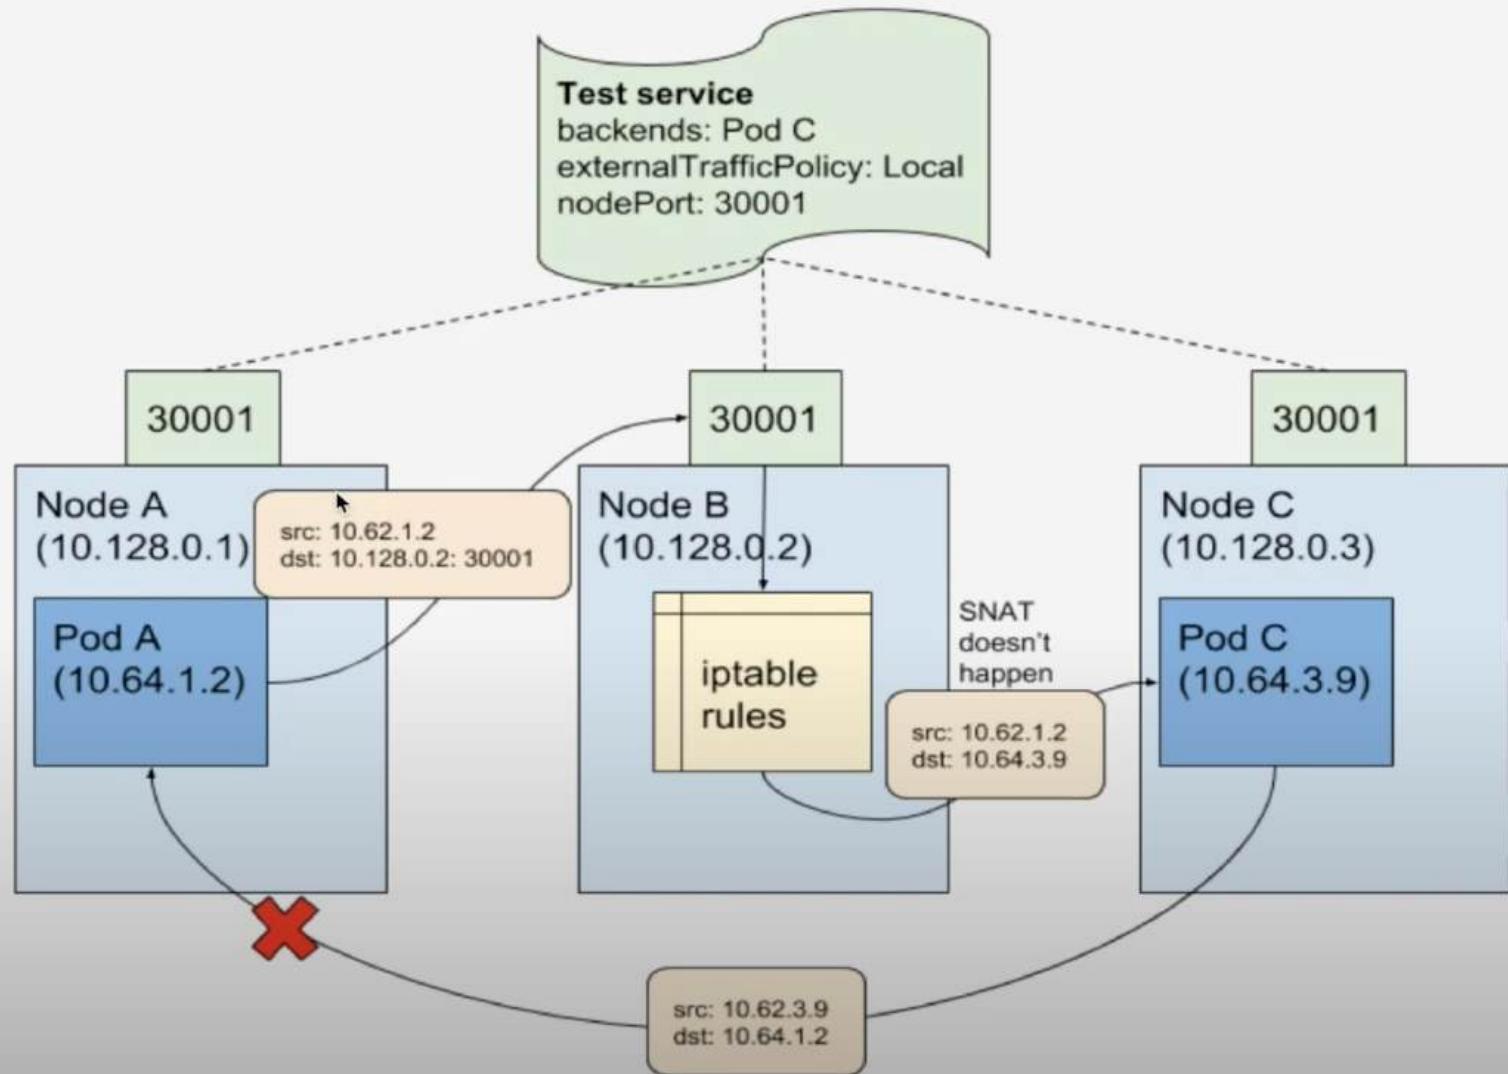
Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: ClusterIP
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```



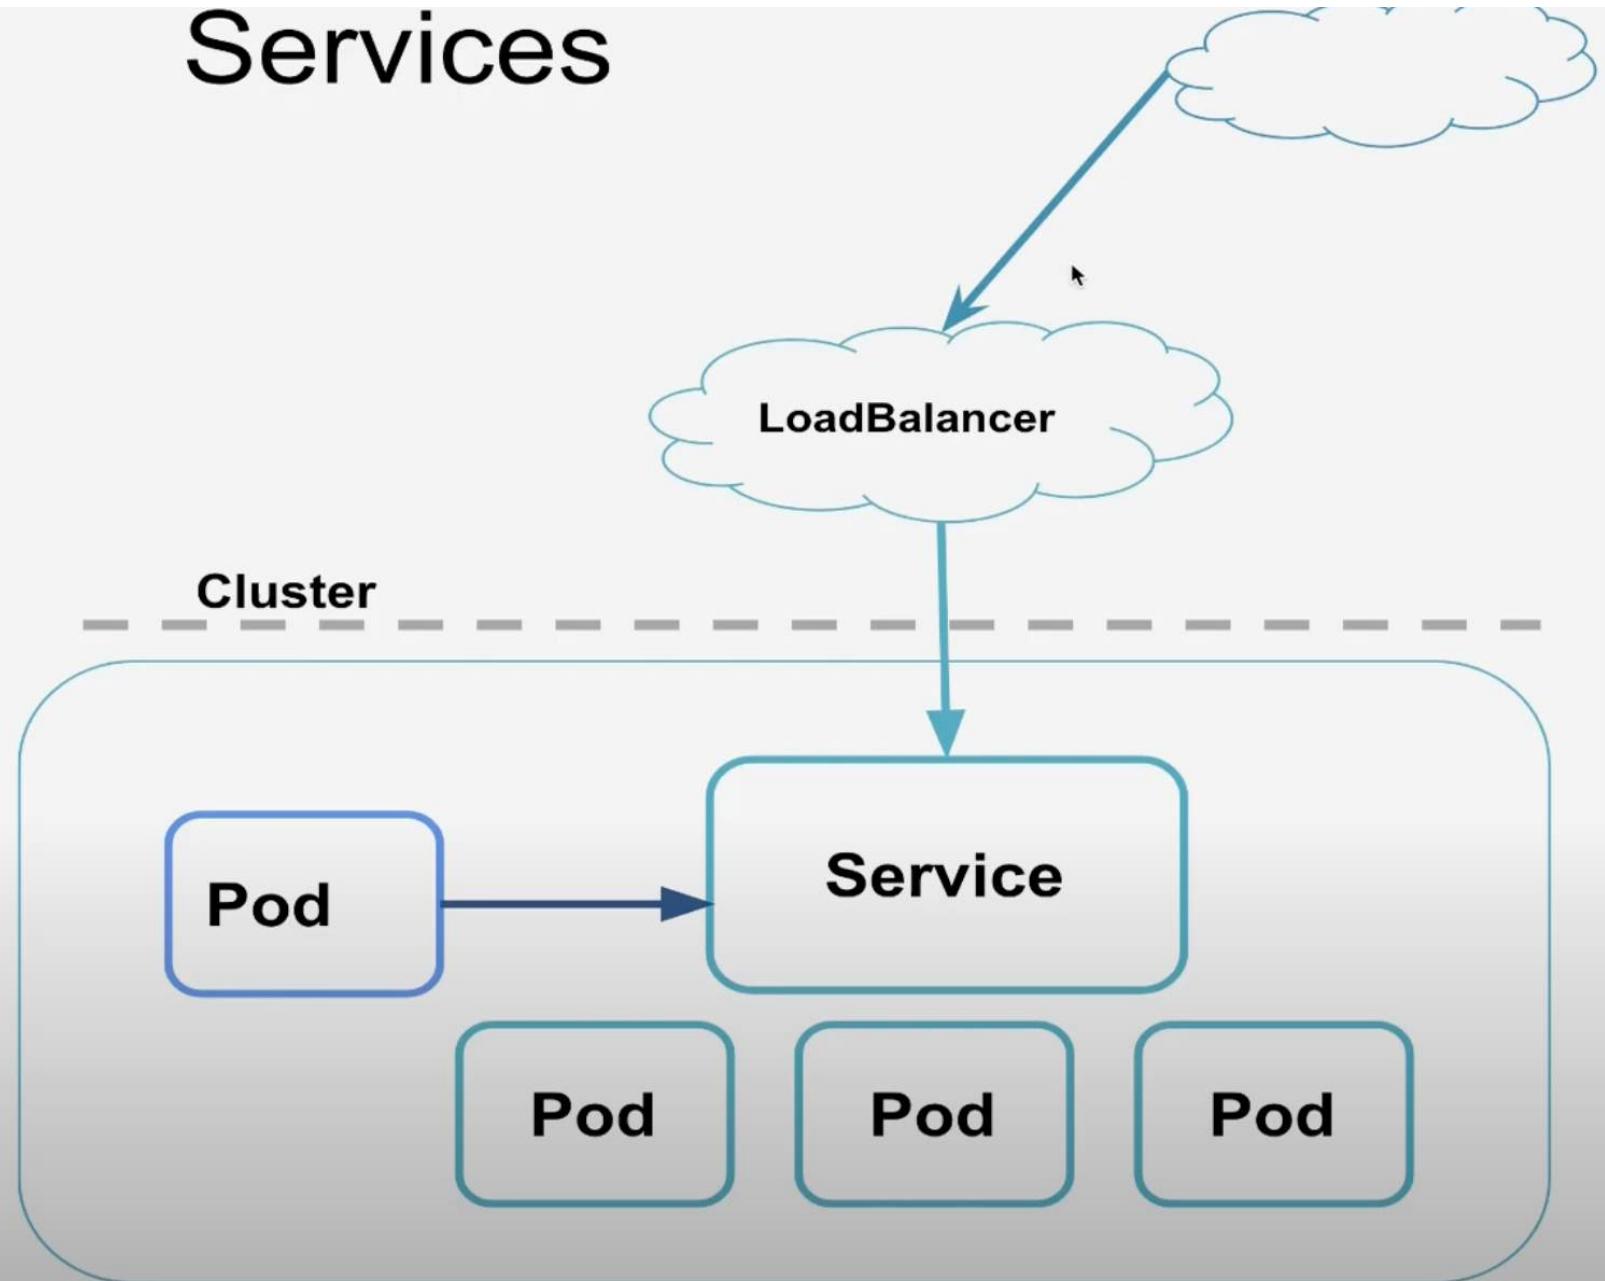
Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: NodePort
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
      nodePort: 30305
```

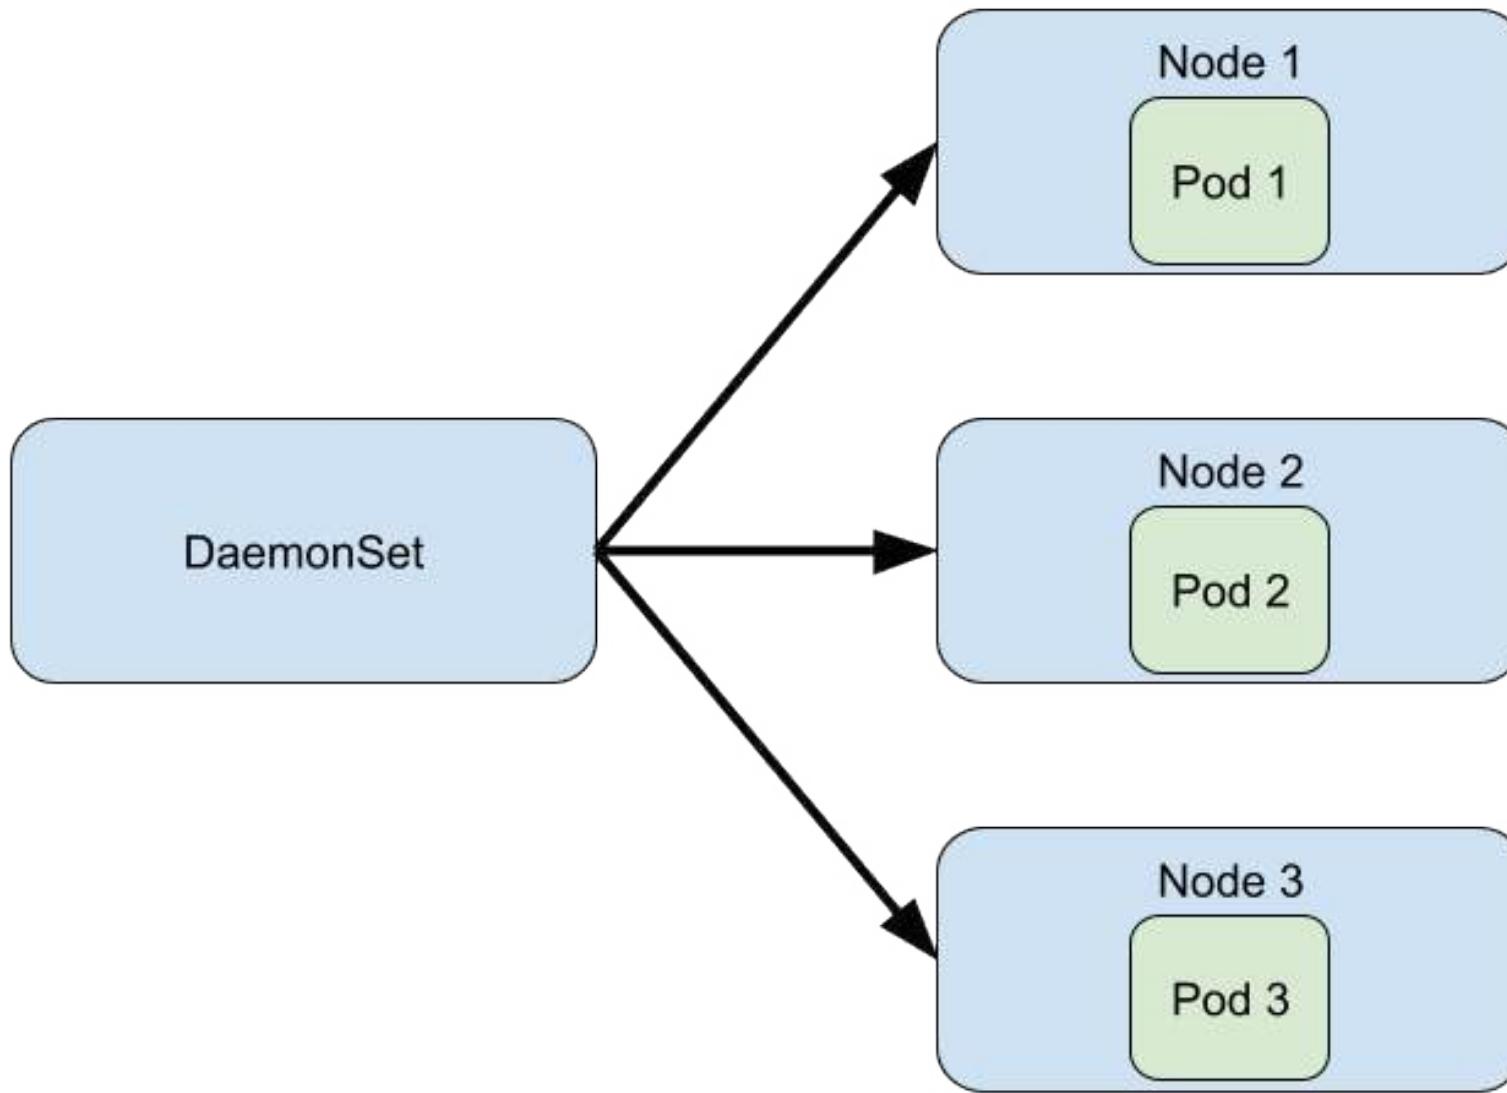


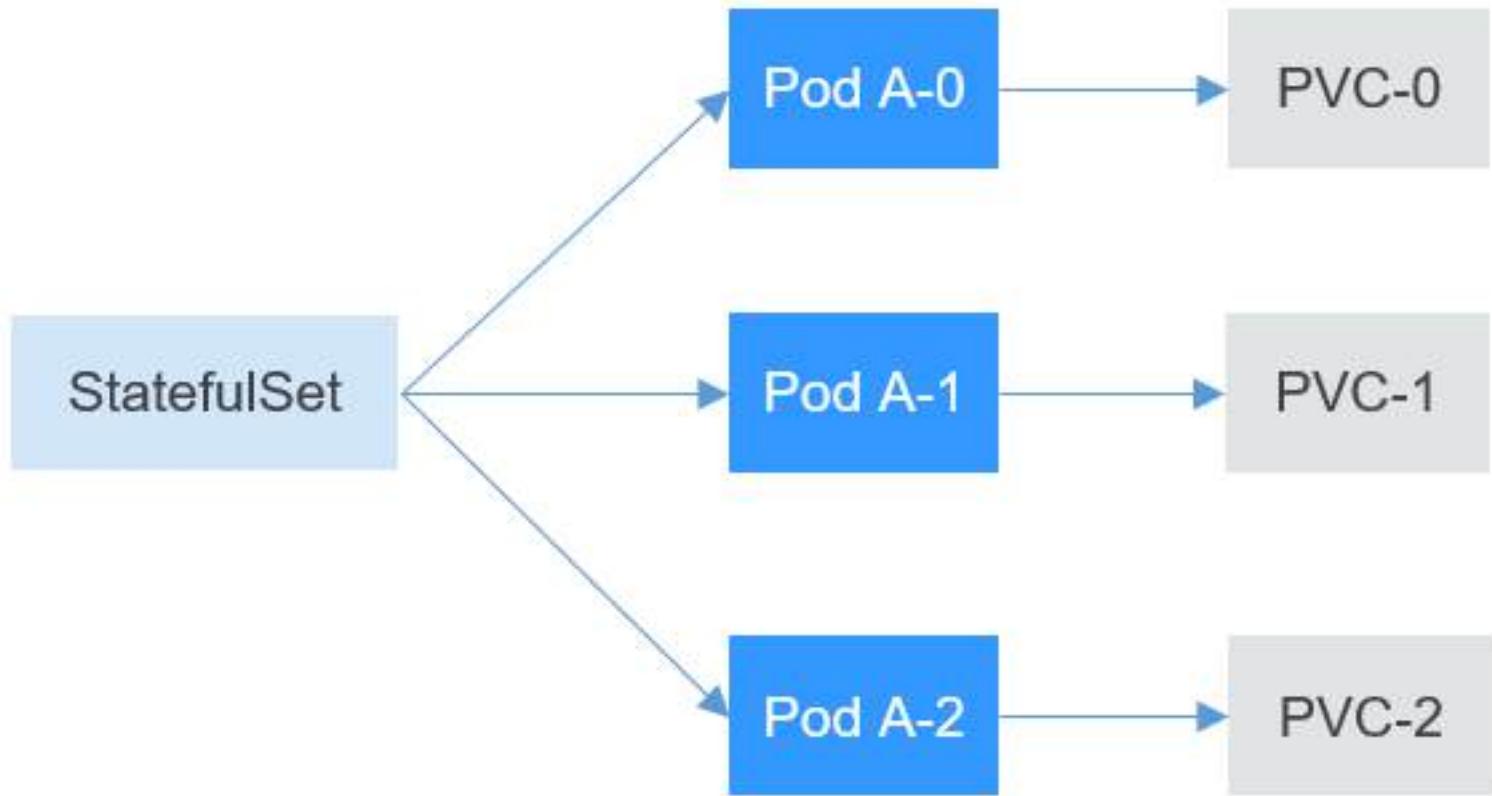
Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
Spec:
  type: LoadBalancer
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

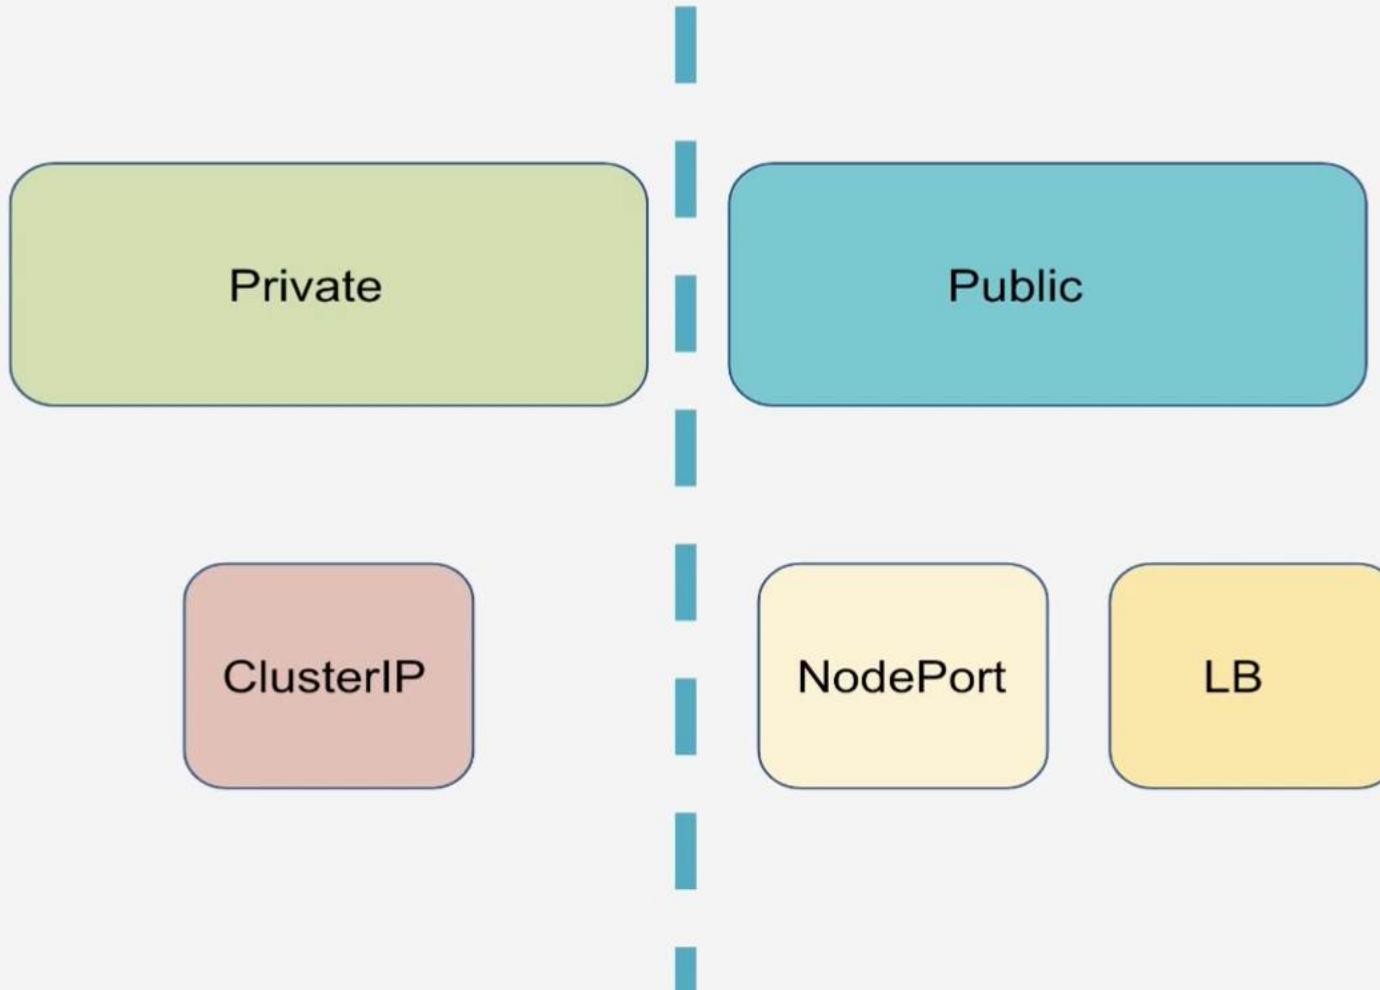


DaemonSet - Add Pods to Nodes

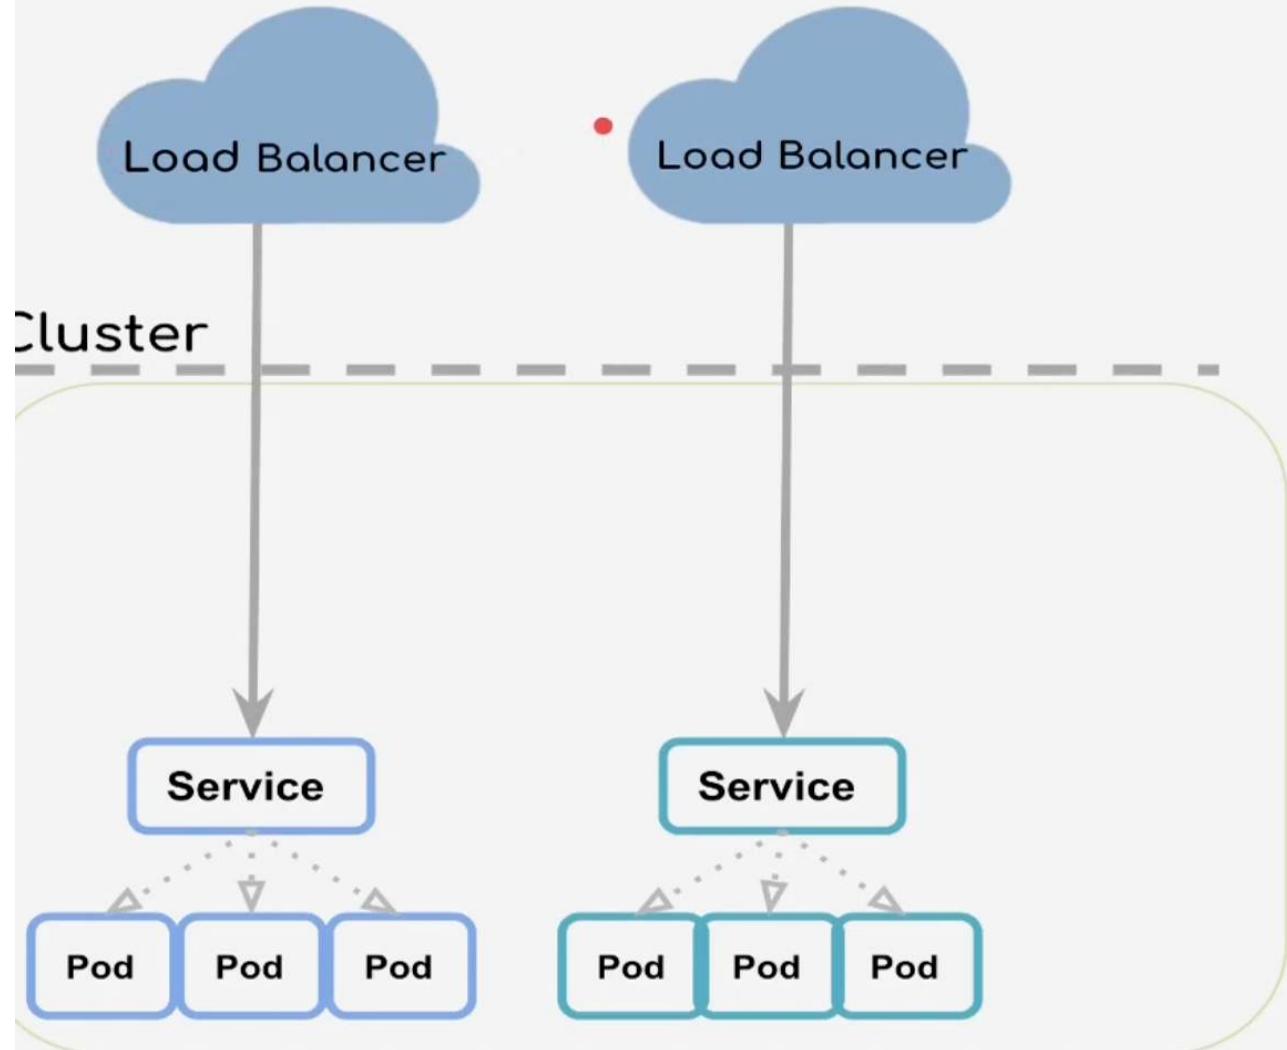




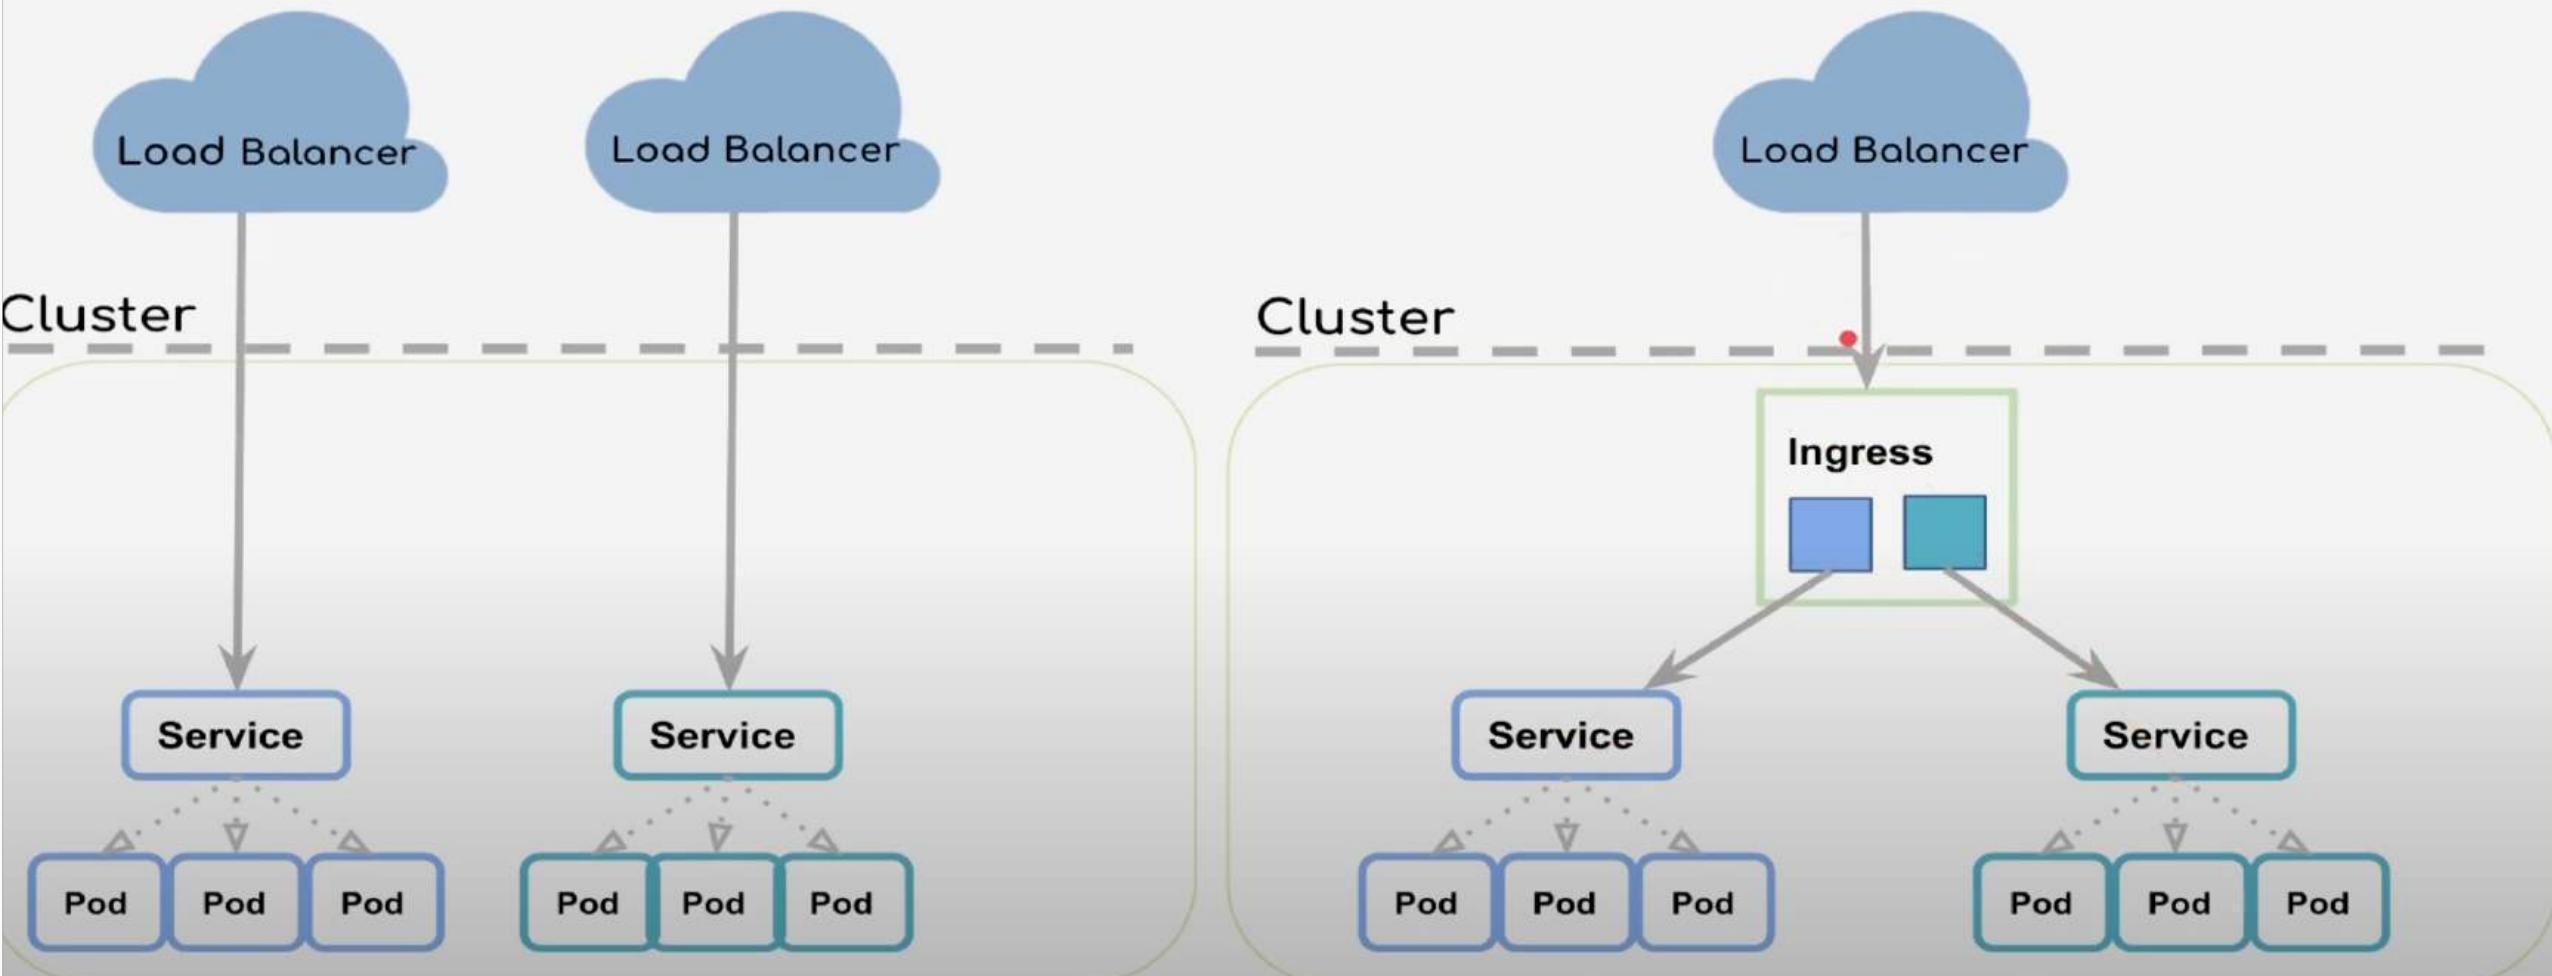
Service Types



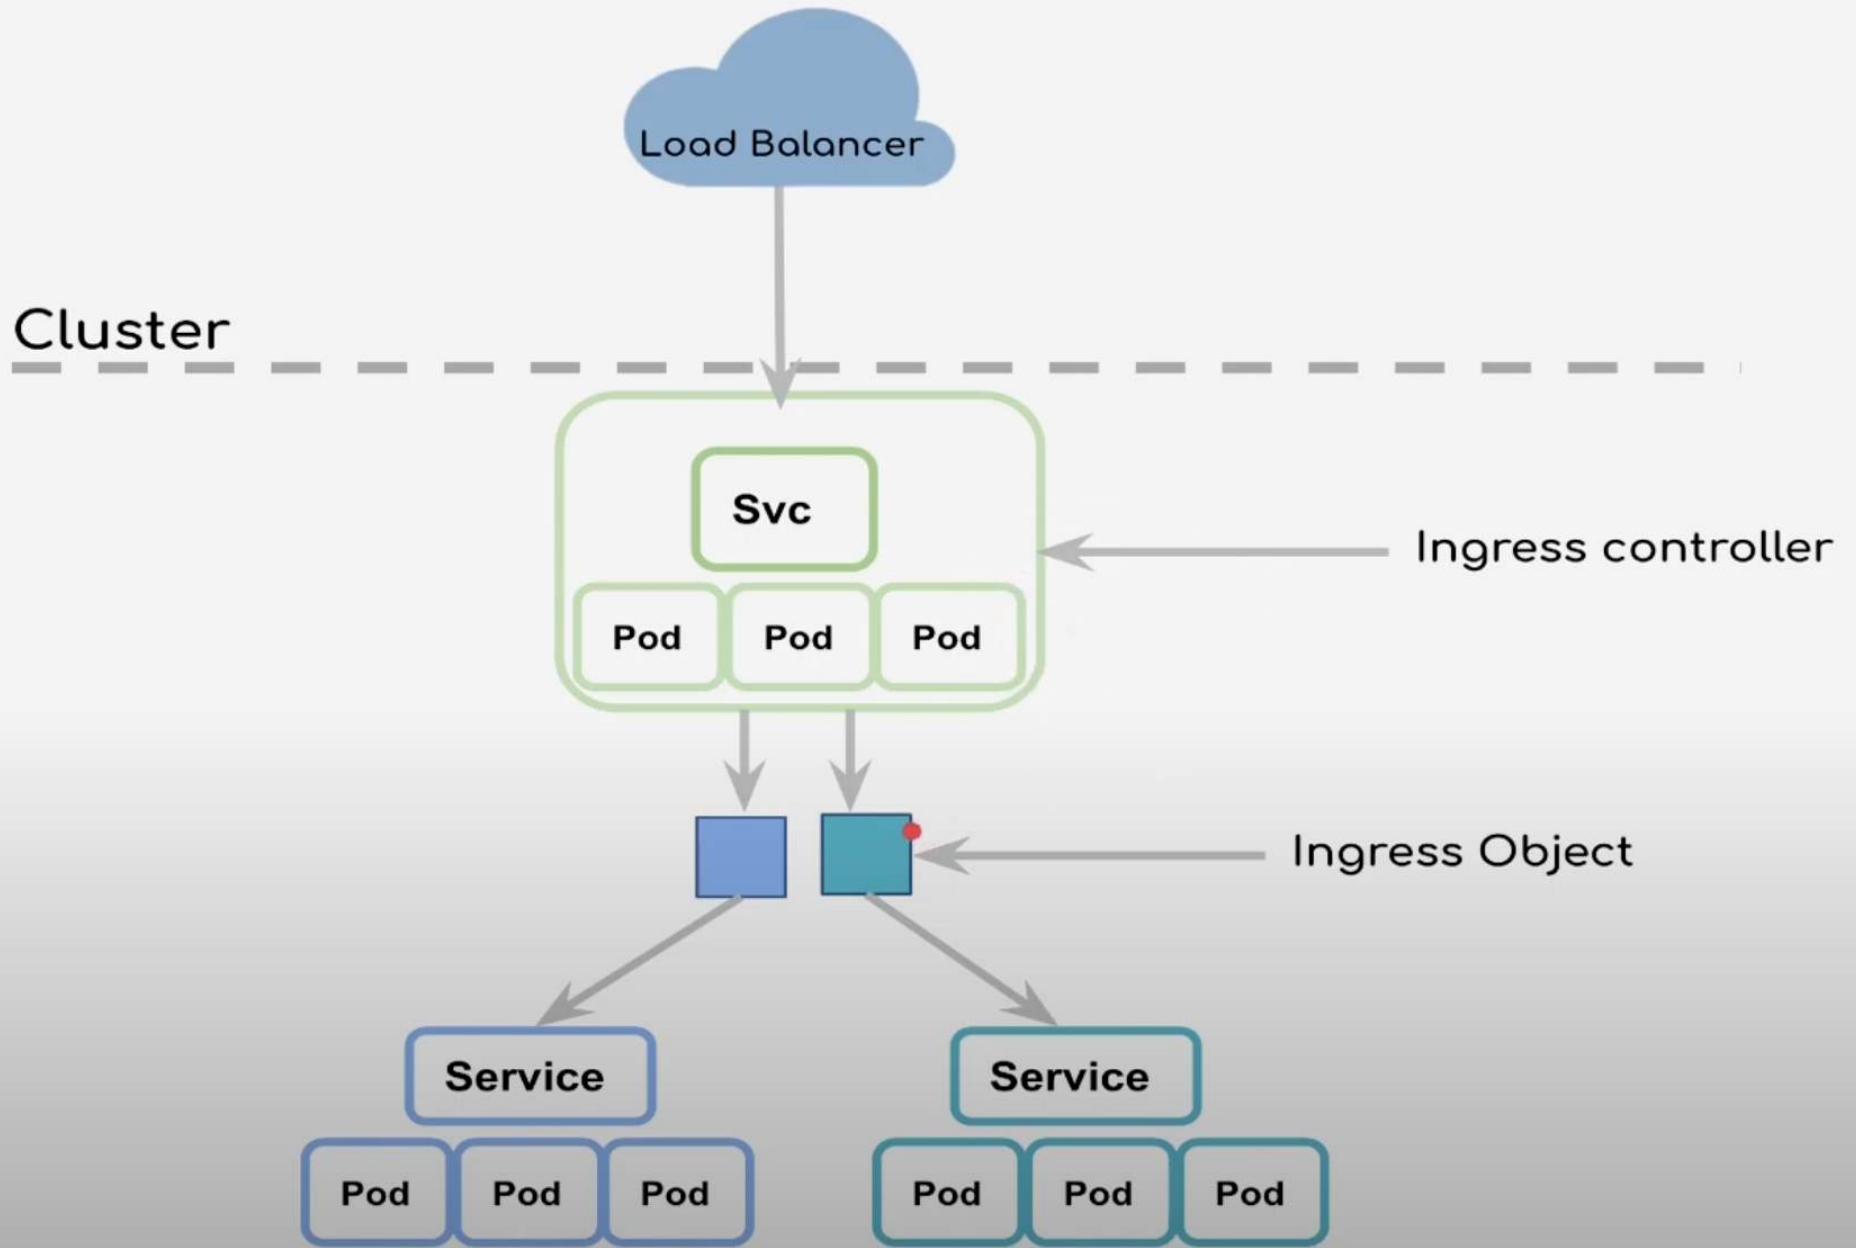
Ingress



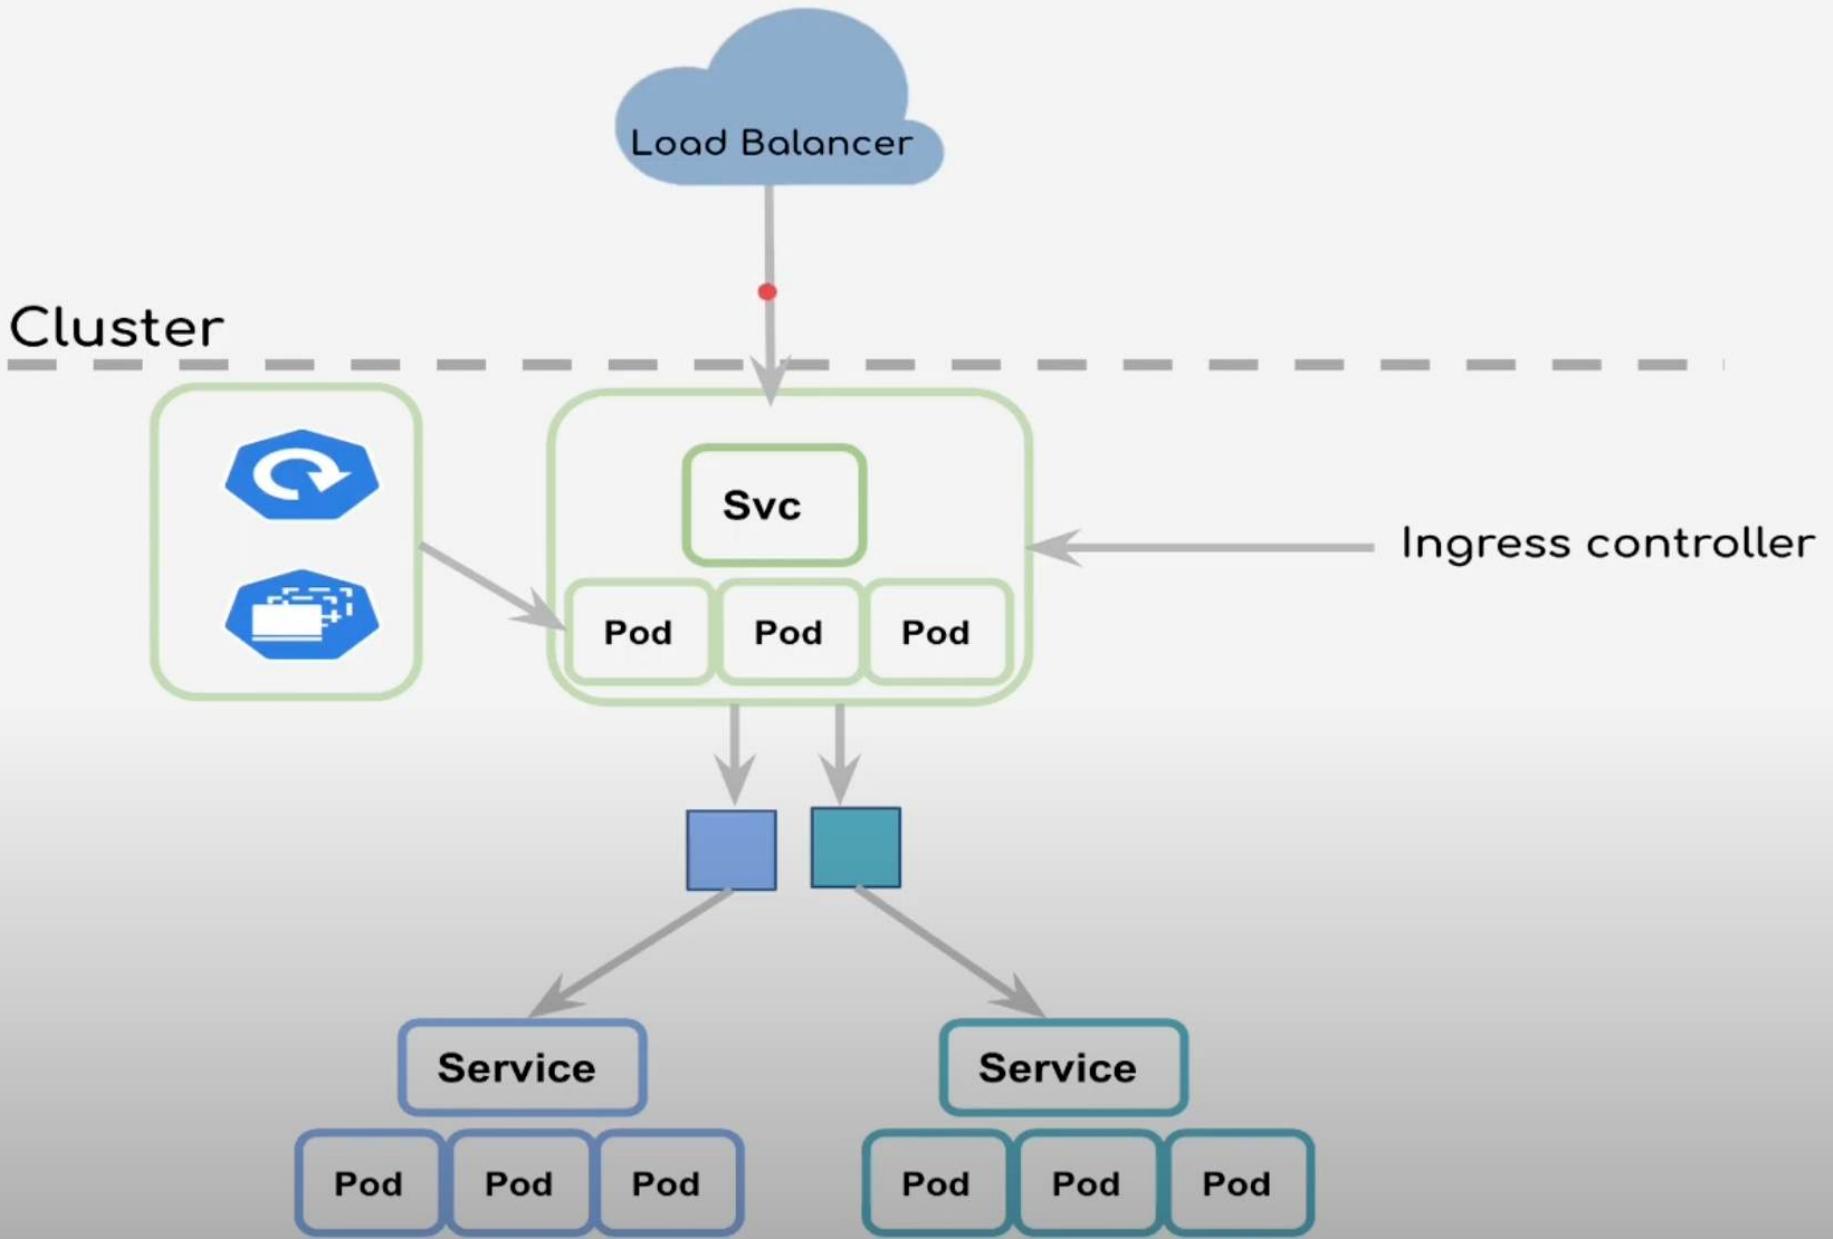
Ingress



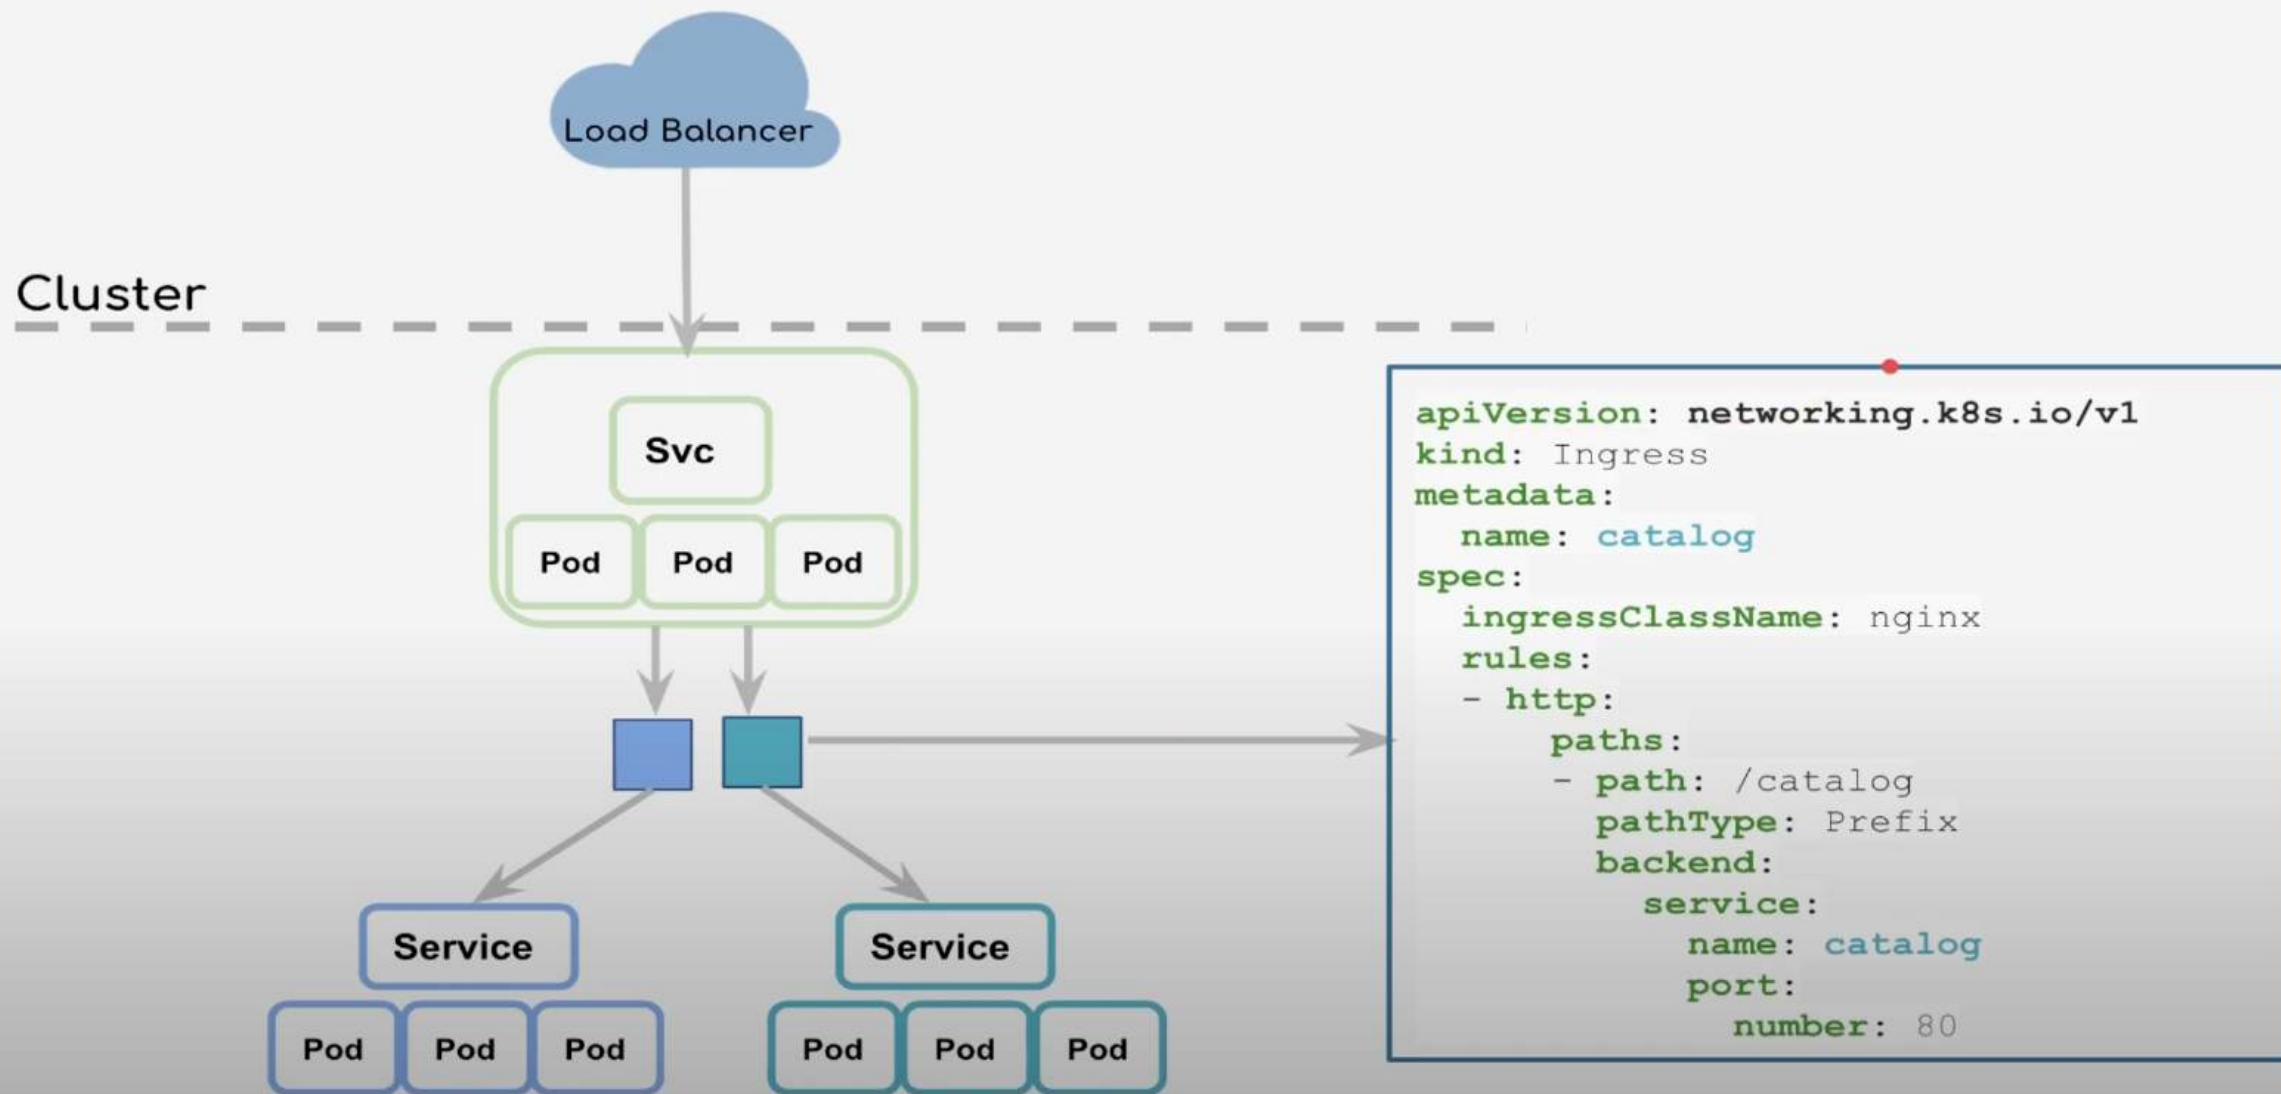
Ingress



Ingress

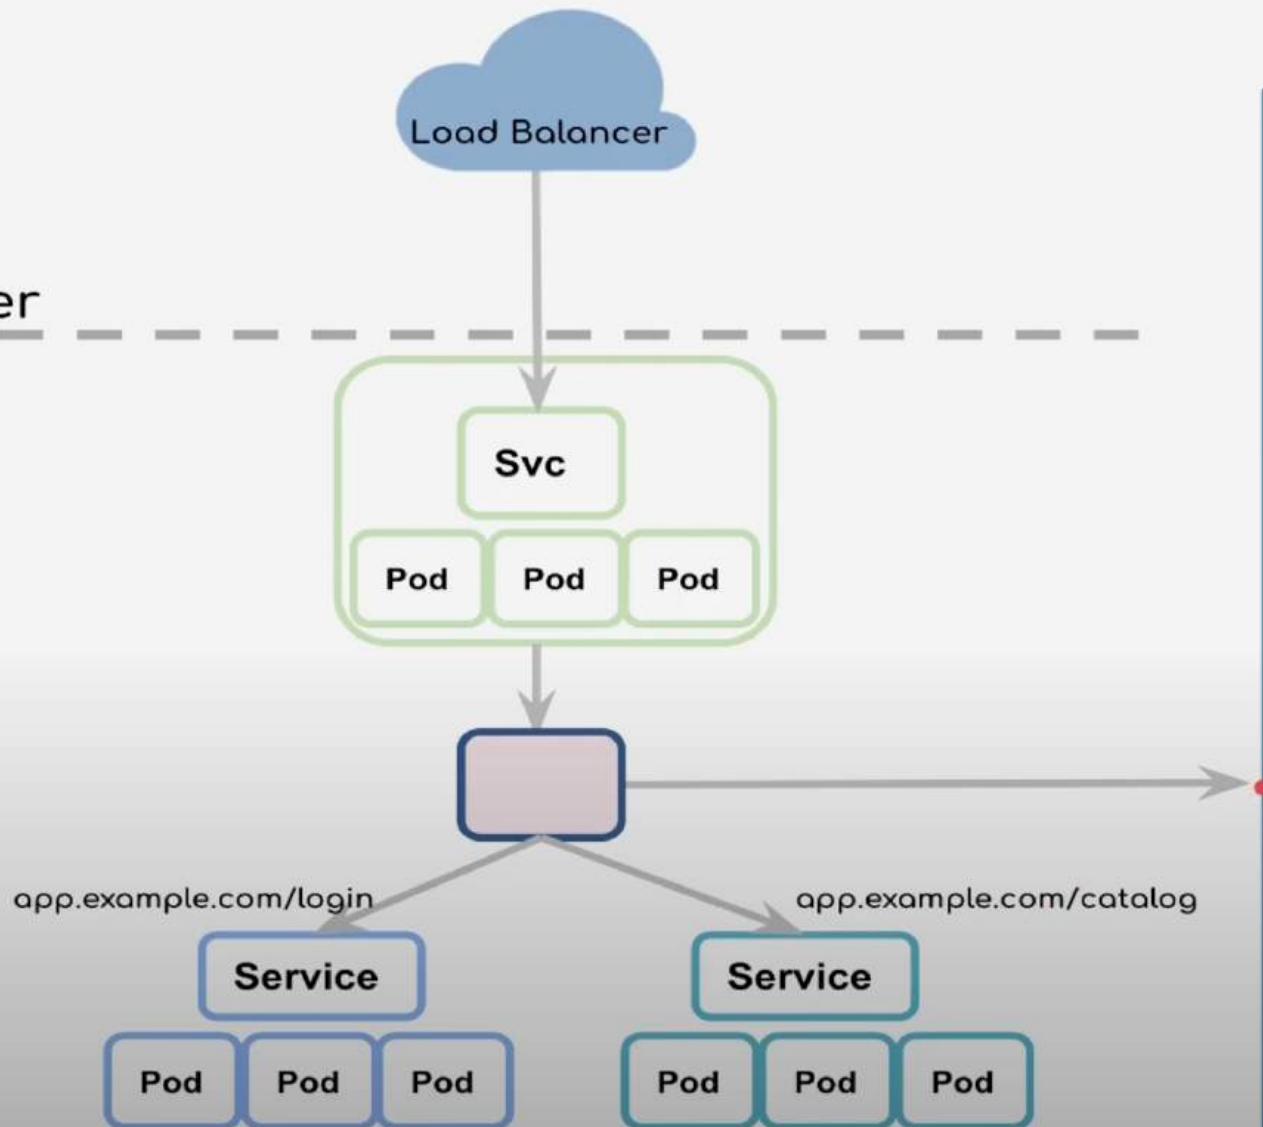


Ingress



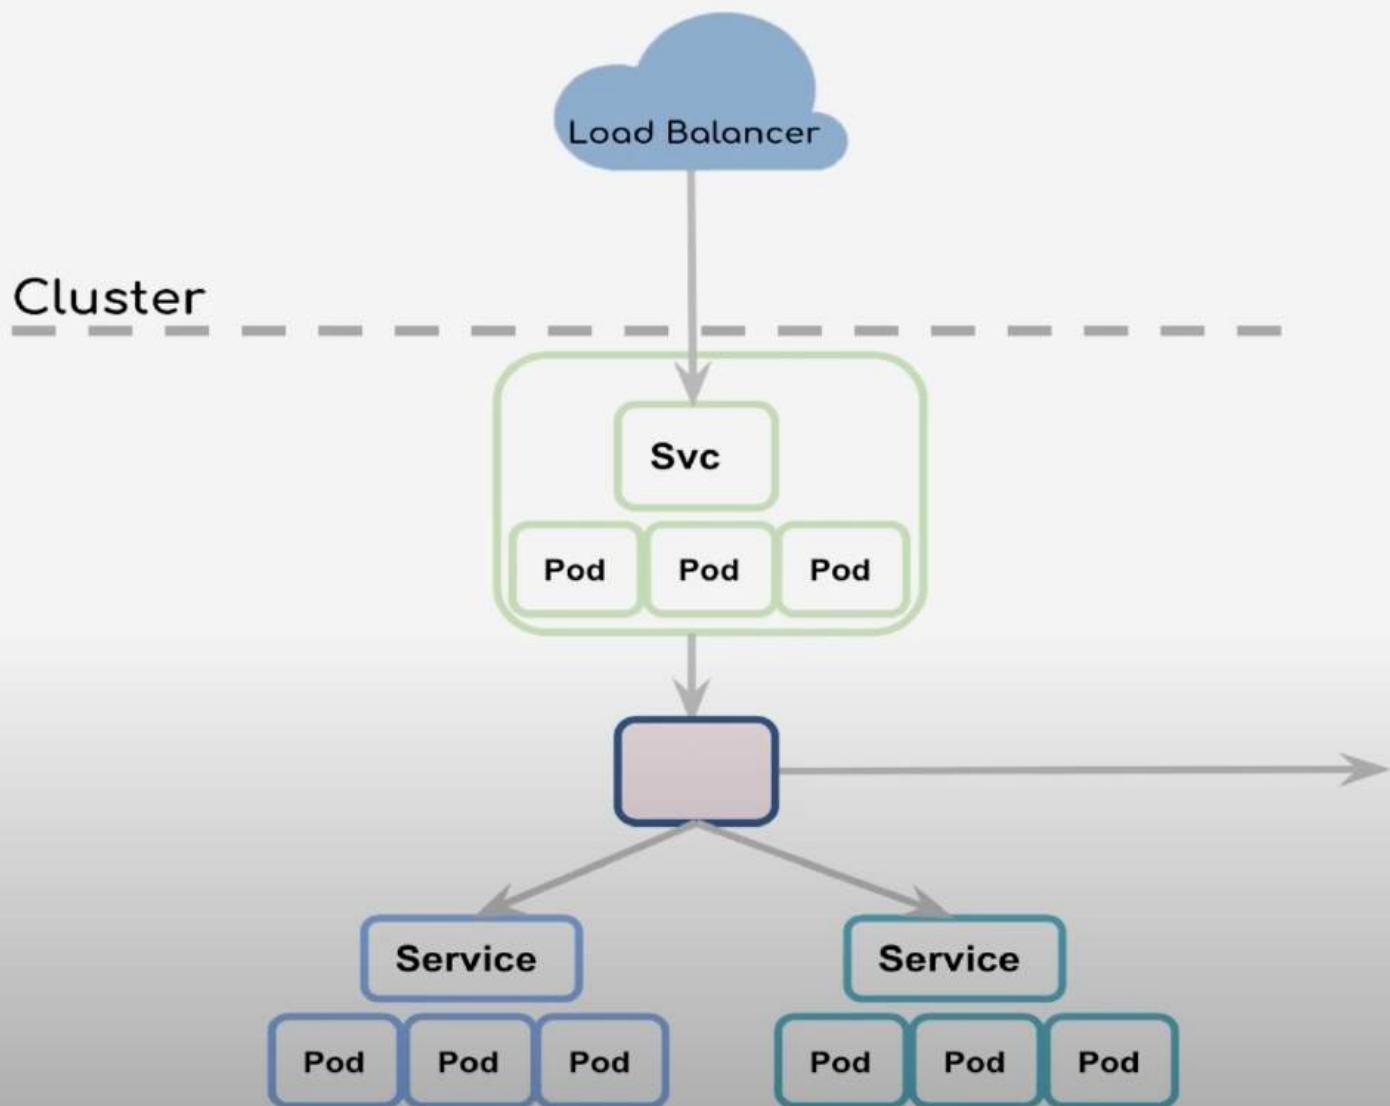
Ingress

Cluster



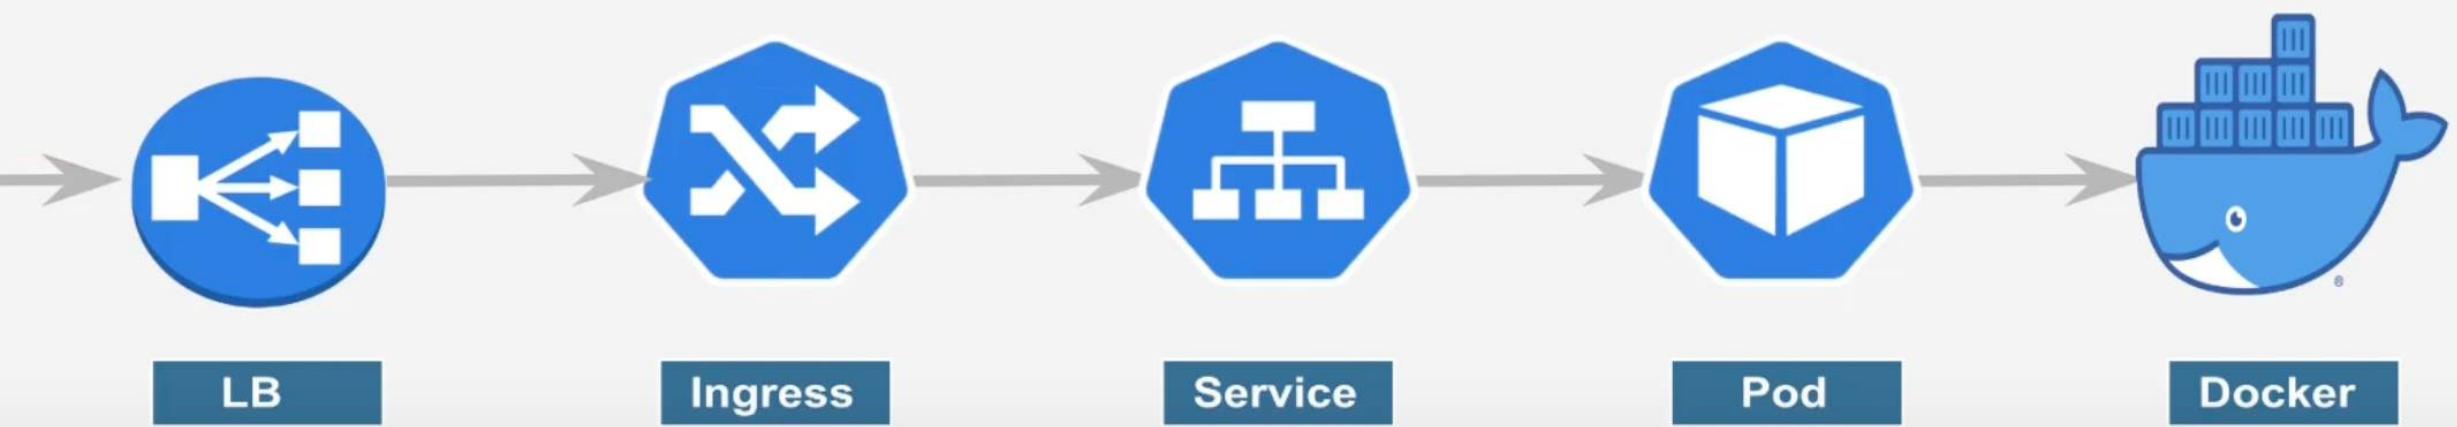
```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example
spec:
  rules:
    - host: app.example.com
      http:
        paths:
          - path: /login
            pathType: Prefix
            backend:
              service:
                name: login
                port:
                  number: 4200
          - path: /catalog
            pathType: Prefix
            backend:
              service:
                name: catalog
                port:
                  number: 8080
```

Ingress



```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example
spec:
  rules:
  - host: login.example.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            name: login
            port:
              number: 80
  - host: catalog.example.com
    http:
      paths:
      - pathType: Prefix
        path: "/"
        backend:
          service:
            name: catalog
            port:
              number: 80
```

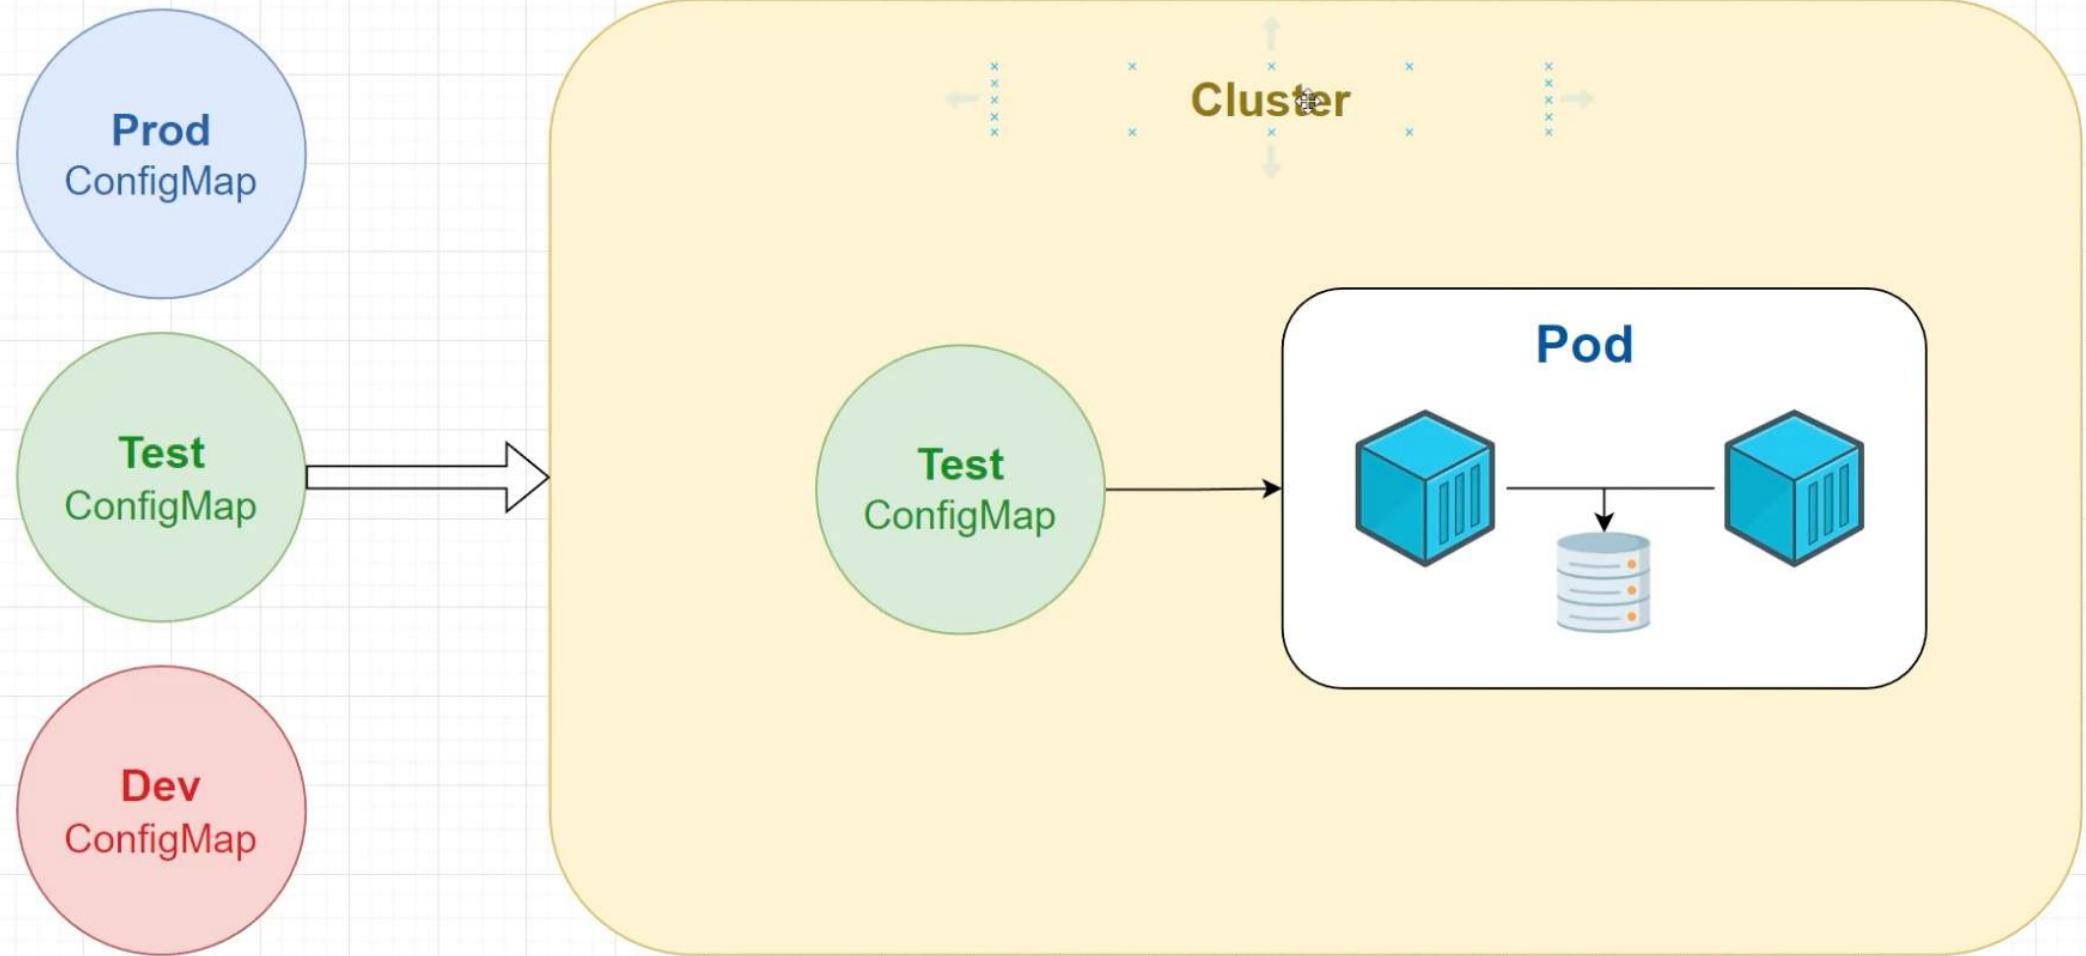
Ingress



Config Maps

A **ConfigMap** lets you **store configuration** for other **objects** to use. They accept **key-value pairs** as their values. **ConfigMaps** are designed to store **config parameters** and **inject** them into **running pods**.

ConfigMaps are ideal for most situations where you want to supply **environment-specific configuration values** to your **pods** such as IP address of your application's **database server** or the **URL** of a **proxy service**.



Secrets

A Kubernetes **Secret** is sensitive information – login **usernames** and **passwords**, **tokens**, **keys**, etc – that is used within a Kubernetes environment.

The primary purpose of **Secrets** is to reduce the risk of exposing sensitive data while deploying applications on Kubernetes.

Secrets have a **size limit of 1 MB**. When it comes to **implementation**, you can either **mount Secrets as volumes or expose them as environment variables** inside the Pod manifest files.